

Lecture 23: 3D Vision

Reminder: A5

Recurrent networks, attention, Transformers

Due on **Tuesday 4/12**, 11:59pm ET

A6

Will cover image generation and visualization:

Generative Models: GANs and VAEs

Network visualization: saliency maps, adversarial examples, class visualizations


Style Transfer

Should be released tonight; due 2 weeks after release

YOU CANNOT USE LATE DAYS ON A6!!!!

Last Time: Self-Supervised Learning


Context Prediction




Colorization



Contrastive Learning



Multimodal Pretraining



Previously: Predicting 2D Shapes of Objects


Classification



CAT

No spatial extent


Semantic Segmentation



GRASS, CAT, TREE,
SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Objects

Instance Segmentation




DOG, DOG, CAT

Today: Predicting 3D Shapes of Objects


Mask R-CNN:

2D Image -> 2D shapes



Mesh R-CNN:

2D Image -> 3D shapes



He, Gkioxari, Dollár, and
Girshick, "Mask R-CNN",
ICCV 2017

Gkioxari, Malik, and Johnson,
"Mesh R-CNN", ICCV 2019


Focus on Two Problems today

Predicting 3D Shapes
from single image


Processing 3D
input data



Input Image



3D Shape



3D Shape

→ Chair

Many more topics in 3D Vision!

Computing correspondences

Multi-view stereo

Structure from Motion

Simultaneous Localization and Mapping (SLAM)


Self-supervised learning

Differentiable graphics

3D Sensors


3D Shape Representations

8
8
2
2
2
2




Depth Map


Voxel Grid



Pointcloud




Mesh




Implicit Surface

3D Shape Representations


8
8
2
2
2
2




Depth Map



Pointcloud



Mesh




Implicit Surface

3D Shape Representations: Depth Map

For each pixel, **depth map** gives distance from the camera to the object in the world at that pixel

RGB image + Depth image
= RGB-D Image (2.5D)

This type of data can be recorded directly for some types of 3D sensors (e.g. Microsoft Kinect)




RGB Image: $3 \times H \times W$ Depth Map: $H \times W$

Eigen and Fergus, "Predicting Depth, Surface Normals and Semantic Labels with a Common Multi-Scale Convolutional Architecture", ICCV 2015

Predicting Depth Maps





RGB Input Image:
 $3 \times H \times W$



**Fully Convolutional
network**

Predicted Depth Image:
 $1 \times H \times W$




**Per-Pixel Loss
(L2 Distance)**



Predicted Depth Image:
 $1 \times H \times W$

Eigen, Puhrsh, and Fergus, "Depth Map Prediction from a Single Image using a Multi-Scale Deep Network", NeurIPS 2014
Eigen and Fergus, "Predicting Depth, Surface Normals and Semantic Labels with a Common Multi-Scale Convolutional Architecture", ICCV 2015

Problem: Scale / Depth Ambiguity



Predicting Depth Maps


Predicted Depth Image:
 $1 \times H \times W$

Scale invariant loss



$$\begin{aligned} D(y, y^*) &= \frac{1}{2n^2} \sum_{i,j} ((\log y_i - \log y_j) - (\log y_i^* - \log y_j^*))^2 \\ &= \frac{1}{n} \sum_i d_i^2 - \frac{1}{n^2} \sum_{i,j} d_i d_j = \frac{1}{n} \sum_i d_i^2 - \frac{1}{n^2} \left(\sum_i d_i \right)^2 \end{aligned}$$



RGB Input Image:
 $3 \times H \times W$




**Fully Convolutional
network**



Predicted Depth Image:
 $1 \times H \times W$

Per-Pixel Loss
(Scale invariant)




Eigen, Puhrsh, and Fergus, "Depth Map Prediction from a Single Image using a Multi-Scale Deep Network", NeurIPS 2014
Eigen and Fergus, "Predicting Depth, Surface Normals and Semantic Labels with a Common Multi-Scale Convolutional Architecture", ICCV 2015

3D Shape Representations: Surface Normals

For each pixel, **surface normals** give a vector giving the normal vector to the object in the world for that pixel



RGB Image: $3 \times H \times W$




Normals: $3 \times H \times W$

Eigen and Fergus, "Predicting Depth, Surface Normals and Semantic Labels with a Common Multi-Scale Convolutional Architecture", ICCV 2015

Predicting Normals





RGB Input Image:
 $3 \times H \times W$



**Fully Convolutional
network**

Ground-truth Normals:
 $3 \times H \times W$



Predicted Normals:
 $3 \times H \times W$


Per-Pixel Loss:
 $(x \cdot y) / (|x| |y|)$

Recall:

$$\begin{aligned} x \cdot y \\ = |x| |y| \cos \theta \end{aligned}$$


3D Shape Representations

8
8
2
2
2
2




Depth
Map


**Voxel
Grid**



Pointcloud




Mesh



Implicit
Surface


3D Shape Representations: Voxels

- Represent a shape with a $V \times V \times V$ grid of occupancies
- Just like segmentation masks in Mask R-CNN, but in 3D!
- (+) Conceptually simple: just a 3D grid!
- (-) Need high spatial resolution to capture fine structures
- (-) Scaling to high resolutions is nontrivial!



Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016


Processing Voxel Inputs: 3D Convolution



Train with classification loss

Wu et al, "3D ShapeNets: A Deep Representation for Volumetric Shapes", CVPR 2015



Generating Voxel Shapes: 3D Convolution



Train with per-voxel cross-entropy loss

Choy et al, "3D-R2N2: A Unified Approach for Single and Multi-view 3D Object Reconstruction", ECCV 2016

Generating Voxel Shapes: “Voxel Tubes”



Input image:
 $3 \times 112 \times 112$

2D Features:
 $C \times H \times W$


3D Features:
 $C' \times D' \times H' \times W'$

Voxels:
 $V \times V \times V$

Train with per-voxel cross-entropy loss


Voxel Problems: Memory Usage

Storing 1024^3 voxel grid takes 4GB of memory!



Scaling Voxels: Oct-Trees


Use voxel grids with heterogenous resolution!



Tatarchenko et al, "Octree Generating Networks: Efficient Convolutional Architectures for High-resolution 3D Outputs", ICCV 2017


3D Shape Representations

8
8
2
2
2
2




Depth
Map


Voxel
Grid



Pointcloud




Mesh



Implicit
Surface

3D Shape Representations: Point Cloud


- Represent shape as a set of P points in 3D space
- (+) Can represent fine structures without huge numbers of points
- (-) Requires new architecture, losses, etc
- (-) Doesn't explicitly represent the surface of the shape: extracting a mesh for rendering or other applications requires post-processing



Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Processing Pointcloud Inputs: PointNet


Want to process pointclouds as sets:
order should not matter



Qi et al, "PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation", CVPR 2017

Qi et al, "PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space", NeurIPS 2017

Generating Pointcloud Outputs



Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds **as sets!**

Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds as sets!

Chamfer distance is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$


Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds as sets!

Chamfer distance is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$




Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds as sets!

Chamfer distance is the sum of L2 distance to each point's nearest neighbor in the other set

$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$




Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017

Predicting Point Clouds: Loss Function

We need a (differentiable) way to compare pointclouds as sets!

Chamfer distance is the sum of L2 distance to each point's nearest neighbor in the other set


$$d_{CD}(S_1, S_2) = \sum_{x \in S_1} \min_{y \in S_2} \|x - y\|_2^2 + \sum_{y \in S_2} \min_{x \in S_1} \|x - y\|_2^2$$



Fan et al, "A Point Set Generation Network for 3D Object Reconstruction from a Single Image", CVPR 2017


3D Shape Representations

8
8
2
2
2
2




Depth
Map


Voxel
Grid



Pointcloud



Mesh



Implicit
Surface

3D Shape Representations: Triangle Mesh


Represent a 3D shape as a set of triangles

Vertices: Set of V points in 3D space

Faces: Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes



3D Shape Representations: Triangle Mesh

Represent a 3D shape as a set of triangles


Vertices: Set of V points in 3D space

Faces: Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes

(+) Adaptive: Can represent flat surfaces very efficiently, can allocate more faces to areas with fine detail



[Dolphin image](#) is in the public domain

3D Shape Representations: Triangle Mesh

Represent a 3D shape as a set of triangles

Vertices: Set of V points in 3D space


Faces: Set of triangles over the vertices

(+) Standard representation for graphics

(+) Explicitly represents 3D shapes

(+) Adaptive: Can represent flat surfaces very efficiently, can allocate more faces to areas with fine detail

(+) Can attach data on verts and interpolate over the whole surface: RGB colors, texture coordinates, normal vectors, etc.




UV mapping figure is licensed under CC BY-SA 3.0. Figure slightly reorganized.

Predicting Meshes: Pixel2Mesh

Input: Single RGB
Image of an object

Output: Triangle
mesh for the object




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

Predicting Meshes: Pixel2Mesh

Input: Single RGB
Image of an object

Key ideas:
Iterative Refinement
Graph Convolution
Vertex Aligned-Features
Chamfer Loss Function

Output: Triangle
mesh for the object



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018


Predicting Triangle Meshes: Iterative Refinement

Idea #1: Iterative mesh refinement

Start from initial ellipsoid mesh

Network predicts offsets for each vertex

Repeat.



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018


Predicting Triangle Meshes: Graph Convolution

$$f'_i = W_0 f_i + \sum_{j \in N(i)} W_1 f_j$$

Vertex v_i has feature f_i

New feature f'_i for vertex v_i depends on feature of neighboring vertices $N(i)$

Use same weights W_0 and W_1 to compute all outputs




Input: Graph with a feature vector at each vertex

Output: New feature vector for each vertex

Predicting Triangle Meshes: Graph Convolution

Each of these blocks consists of a stack of **graph convolution layers** operating on edges of the mesh




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

Predicting Triangle Meshes: Graph Convolution

Each of these blocks consists of a stack of **graph convolution layers** operating on edges of the mesh

Problem: How to incorporate image features?




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

Predicting Triangle Meshes: Vertex-Aligned Features

Idea #2: Aligned vertex features

For each vertex of the mesh:

- Use camera information to project onto image plane
- Use bilinear interpolation to sample a CNN feature



Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018


Predicting Triangle Meshes: Vertex-Aligned Features

Idea #2: Aligned vertex features

For each vertex of the mesh:


- Use camera information to project onto image plane
- Use bilinear interpolation to sample a CNN feature

Similar to RoI-Align operation from detection: maintains alignment between input image and feature vectors



Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018


Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Idea: Convert meshes to pointclouds, then compute loss



vs




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018


Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Idea: Convert meshes to pointclouds, then compute loss



vs




Sample points from the surface of the ground-truth mesh (offline)

Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Loss = Chamfer distance between **predicted verts** and **ground-truth samples**




Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Loss = Chamfer distance between **predicted verts** and **ground-truth samples**

Problem: Doesn't take the interior of predicted faces into account!




Sample points from the surface of the ground-truth mesh (offline)

Predicting Meshes: Loss Function

The same shape can be represented with different meshes – how can we define a loss between predicted and ground-truth mesh?

Loss = Chamfer distance between **predicted samples** and **ground-truth samples**

Sample points
from the surface
of the predicted
mesh (online!)



Sample points from the
surface of the ground-
truth mesh (offline)

Smith et al, "GEOMetrics: Exploiting Geometric Structure for Graph-Encoded Objects", ICML 2019


Predicting Meshes: Loss Function

Problem: Need to sample online! Must be efficient!

Problem: Need to backprop through sampling!

Loss = Chamfer distance between **predicted samples** and **ground-truth samples**

Sample points
from the surface
of the predicted
mesh (online!)




Sample points from the
surface of the ground-
truth mesh (offline)

Predicting Meshes: Pixel2Mesh

Input: Single RGB
Image of an object

Key ideas:
Iterative Refinement
Graph Convolution
Vertex Aligned-Features
Chamfer Loss Function

Output: Triangle
mesh for the object





Wang et al, "Pixel2Mesh: Generating 3D Mesh Models from Single RGB Images", ECCV 2018

3D Shape Prediction: Mesh R-CNN



Mask R-CNN:

2D Image -> 2D shapes



Mesh R-CNN:

2D Image -> Triangle Meshes



He, Gkioxari, Dollár, and
Girshick, "Mask R-CNN",
ICCV 2017

Gkioxari, Malik, and Johnson,
"Mesh R-CNN", ICCV 2019

Mesh R-CNN: Task

Input: Single RGB image

Output:


A set of detected objects

For each object:

- Bounding box
- Category label
- Instance segmentation
- 3D triangle mesh


Mask R-CNN

Mesh head




Mesh R-CNN: Hybrid 3D shape representation

Mesh deformation gives good results, but the topology (verts, faces, genus, connected components) fixed by the initial mesh




Mesh R-CNN: Hybrid 3D shape representation

Mesh deformation gives good results, but the topology (verts, faces, genus, connected components) fixed by the initial mesh



Our approach: Use voxel predictions to create initial mesh prediction!



Mesh R-CNN Pipeline

Input image



Mesh R-CNN Pipeline

Input image



2D object recognition



Mesh R-CNN Pipeline

Input image



2D object recognition



3D object voxels

Mesh R-CNN Pipeline

Input image



2D object recognition



3D object meshes

3D object voxels

Mesh R-CNN: ShapeNet Results




Mesh R-CNN: Pix3D Results




Mesh R-CNN: Pix3D Results

Predicting many objects per scene




Box & Mask Predictions




Mesh Predictions

Mesh R-CNN: Pix3D Results

Amodal completion: predict occluded parts of objects



Box & Mask Predictions



Mesh Predictions

Mesh R-CNN: Pix3D Results

Segmentation failures propagate to meshes




Box & Mask Predictions



Mesh Predictions


3D Shape Representations

8
8
2
2
2
2




Depth Map


Voxel Grid



Pointcloud



Mesh



Implicit Surface


3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape


$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set

$$\{x : o(x) = \frac{1}{2}\}$$



Implicit function



Explicit Shape


3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape


$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set

$$\{x : o(x) = \frac{1}{2}\}$$



Implicit function



Explicit Shape

Same idea: **signed distance function (SDF)** gives the Euclidean distance to the surface of the shape; sign gives inside / outside


3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set

$$\{x : o(x) = \frac{1}{2}\}$$



Allows for multiscale outputs like Oct-Trees

Mescheder et al, “Occupancy Networks: Learning 3D Reconstruction in Function Space”, CVPR 2019


3D Shape Representations: Implicit Functions

Learn a function to classify arbitrary 3D points as inside / outside the shape

$$o : \mathbb{R}^3 \rightarrow \{0, 1\}$$

The surface of the 3D object is the level set

$$\{x : o(x) = \frac{1}{2}\}$$



Mescheder et al, "Occupancy Networks: Learning 3D Reconstruction in Function Space", CVPR 2019

Neural Radiance Fields (NeRF) for View Synthesis

View Synthesis

Input: Many images of the same scene
(with known camera parameters)



Image source: Mildenhall et al, "Representing Scenes as Neural Radiance Fields for View Synthesis", ECCV 2020

View Synthesis

Input: Many images of the same scene
(with known camera parameters)



Output: Images showing the
scene from novel viewpoints





Image source: Mildenhall et al, "Representing Scenes as Neural Radiance Fields for View Synthesis", ECCV 2020

Stepping Back: Pinhole Camera Model


Camera




Stepping Back: Pinhole Camera Model




Stepping Back: Pinhole Camera Model



Stepping Back: Pinhole Camera Model



Stepping Back: Pinhole Camera Model




Stepping Back: Pinhole Camera Model

Reflected light
captured by
the camera

Camera


Sensors:
Records
light



Light sources
emit light



Objects in the
world reflect light




Stepping Back: Pinhole Camera Model

Reflected light
captured by
the camera

Camera

Sensors:
Records
light



Focal Length

Pinhole:
Lets light in


Light sources
emit light



Opaque objects
block light



Objects in the
world reflect light




Stepping Back: Pinhole Camera Model

Reflected light
captured by
the camera

Camera

Sensors:
Records
light




Focal Length

Light sources
emit light



Transparent objects
attenuate light





Objects in the
world reflect light

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$




Volume Rendering

Abstract away light sources, objects.


For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Point on car:

- (1) Emits red light in hemisphere
- (2) Complete opaque
 $\sigma = 1$




Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Point in empty space:
(1) Emits no light (black)
(2) Completely transparent $\sigma = 0$

Point on car:
(1) Emits red light in hemisphere
(2) Complete opaque
 $\sigma = 1$




Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$


Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:


$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t)\sigma(\mathbf{r}(t))c(\mathbf{r}(t), \mathbf{d})dt$$

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t)\sigma(\mathbf{r}(t))\mathbf{c}(\mathbf{r}(t), \mathbf{d})dt$$

Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(p) \in [0,1]$


Color that a point \mathbf{p} emits in direction \mathbf{d} is $\mathbf{c}(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$




Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$




Transmittance: How much light from the current point will reach the camera?

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$




Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$



Transmittance: How much light from the current point will reach the camera?


Opacity: How opaque is the current point?

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$




Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(p) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$




- Transmittance**: How much light from the current point will reach the camera?
- Opacity**: How opaque is the current point?
- Color**: What color does the current point emit along the direction toward the camera?

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(p) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$

$$T(t) = \exp \left(- \int_{t_n}^t \sigma(\mathbf{r}(s)) ds \right)$$

A diagram showing a blue ray starting from a camera sensor at the left edge and extending to the right. Along the ray, there are two white circles representing "Current point: t " and "Far point: t_f ". The region between these points is shaded grey, representing a volume element. The ray is shown as a thin blue line with arrows at both ends.

Transmittance: How much light from the current point will reach the camera?


Compute transmittance by accumulating volume density up to current point

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$


Volume Density is $\sigma(p) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$

$$T(t) = \exp \left(- \int_{t_n}^t \sigma(\mathbf{r}(s)) ds \right)$$




Approximate integrals with a set of samples:

Volume Rendering

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Color observed by the camera given by **volume rendering equation**:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) c(\mathbf{r}(t), \mathbf{d}) dt$$

$$T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Approximate integrals with a set of samples:

$$C(\mathbf{r}) \approx \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i$$

$$T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Mildenhall et al, “Representing Scenes as Neural Radiance Fields for View Synthesis”, ECCV 2020


Neural Radiance Fields (NeRF)

Abstract away light sources, objects.

For each point in space, need to know:

- (1) How much light does it emit?
- (2) How opaque is it? $\sigma \in [0,1]$


Train a neural network to input position \mathbf{p} and direction \mathbf{d} , output $\sigma(\mathbf{p})$ and $c(\mathbf{p}, \mathbf{d})$



Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$



Approximate integrals with a set of samples:

$$C(\mathbf{r}) \approx \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i$$

$$T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Mildenhall et al, "Representing Scenes as Neural Radiance Fields for View Synthesis", ECCV 2020


Neural Radiance Fields (NeRF)

Abstract away light sources, objects.

For each point in space, need to know:

(1) How much light does it emit?

(2) How opaque is it? $\sigma \in [0,1]$




Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Train a neural network to input position \mathbf{p} and direction \mathbf{d} , output $\sigma(\mathbf{p})$ and $c(\mathbf{p}, \mathbf{d})$

Training loss: Estimated pixel colors $C(\mathbf{r})$ should match actual pixel colors from images



Approximate integrals with a set of samples:

$$C(\mathbf{r}) \approx \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i$$

$$T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Mildenhall et al, “Representing Scenes as Neural Radiance Fields for View Synthesis”, ECCV 2020


Neural Radiance Fields (NeRF)

Abstract away light sources, objects.

For each point in space, need to know:

(1) How much light does it emit?

(2) How opaque is it? $\sigma \in [0,1]$




Parameterize each ray as origin plus direction: $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$

Volume Density is $\sigma(\mathbf{p}) \in [0,1]$

Color that a point \mathbf{p} emits in direction \mathbf{d} is $c(\mathbf{p}, \mathbf{d}) \in [0,1]^3$

Train a neural network to input position \mathbf{p} and direction \mathbf{d} , output $\sigma(\mathbf{p})$ and $c(\mathbf{p}, \mathbf{d})$

Training loss: Estimated pixel colors $C(\mathbf{r})$ should match actual pixel colors from images




After training, can generate novel views of the scene by integrating along rays corresponding to new pixels

Mildenhall et al, “Representing Scenes as Neural Radiance Fields for View Synthesis”, ECCV 2020

Neural Radiance Fields (NeRF): Network Architecture


Fully-connected network: Input position x and direction d , and output volume density and RGB color



Mildenhall et al, “Representing Scenes as Neural Radiance Fields for View Synthesis”, ECCV 2020

Neural Radiance Fields (NeRF): Network Architecture

Fully-connected network: Input position x and direction d , and output volume density and RGB color




Rather than pass raw xyz values to network,
instead use **positional encodings**:

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

Mildenhall et al, "Representing Scenes as Neural Radiance Fields for View Synthesis", ECCV 2020

Neural Radiance Fields: Very Strong Results!



Neural Radiance Fields: Very Strong Results!



Neural Radiance Fields: Very Strong Results!



Neural Radiance Fields: Can extract 3D geometry!



Neural Radiance Fields


Main Problem: Very slow!

Training: 1-2 days on a V100 GPU, for just a single scene!

Inference: Sampling an image from a trained model:
 $(256 \times 256 \text{ pixels}) \times (224 \text{ samples per pixel})$
= 14.6M forward passes through MLP


Tons of follow-up work!

Mip-NeRF: Model cones rather than rays




Barron et al, "Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields", ICCV 2021

Dynamic NeRF: Deformable Scenes



(a) Capture Process



(b) Input




(c) Nerfie



(d) Nerfie Depth


Park et al, "Nerfies: Deformable Neural Radiance Fields", ICCV 2021

RawNeRF: High-Dynamic Range Imagery



Mildenhall et al, "NeRF in the Dark: High Dynamic Range View Synthesis from Noisy Raw Images", CVPR 2022


BlockNeRF: A Neighborhood of San Francisco




Tancik et al, "Block-NeRF: Scalable Large Scene Neural View Synthesis", arXiv 2022

Training NeRF models in minutes!


TensoRF



Plenoxels



Instant-NGP




Yu et al, “Plenoxels: Radiance Fields without Neural Networks”, CVPR 2022

Muller et al, “Instant Neural Graphics Primitives with a Multiresolution Hash Encoding”, arXiv 2022

Chen et al, “TensoRF: Tensorial Radiance Fields”, arXiv 2022

Generalizable NeRF: Same model for many scenes




Yu et al, “pixelNeRF: Neural Radiance Fields from One or Few Images”, CVPR 2021

Wang et al, “IBRNet: Learning Multi-View Image-Based Rendering”, CVPR 2021


Summary: 3D Shape Representations

8
8
2
2
2
2




Depth Map


Voxel Grid



Pointcloud



Mesh



Implicit Surface

Summary: Neural Radiance Fields

Represent neural radiance fields with neural networks

Train using posed RGB images of a scene

Render novel views, extract 3D scene representations

One of the hottest topics in computer vision for past few years

Next Time: Videos