

Lecture 24: Videos

Reminder: A5

Recurrent networks, attention, Transformers

Due on **Tuesday 4/12**, 11:59pm ET

A6

Covers image generation and generative models:

Generative Models: GANs and VAEs

Network visualization: saliency maps, adversarial examples, class visualizations

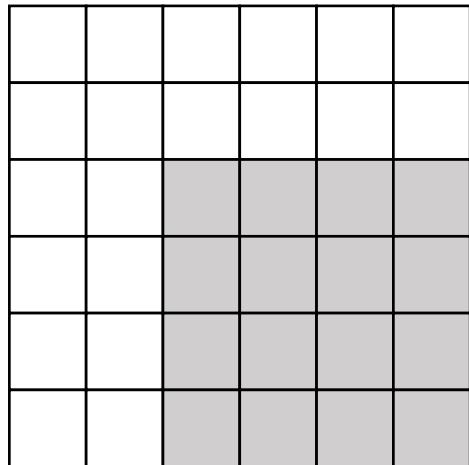
Style Transfer

Due Tuesday 4/26, 11:59pm ET

YOU CANNOT USE LATE DAYS ON A6!!!!

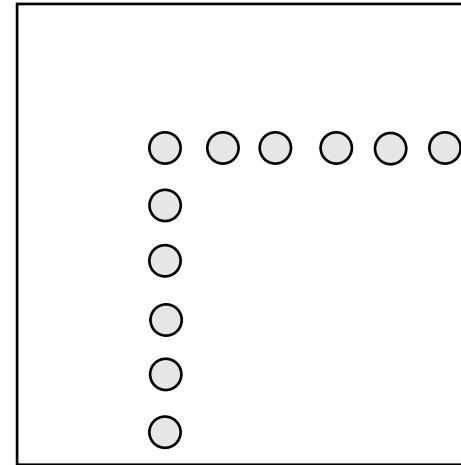
Last Time: 3D Shape Representations

∞
∞
2
2
2
2

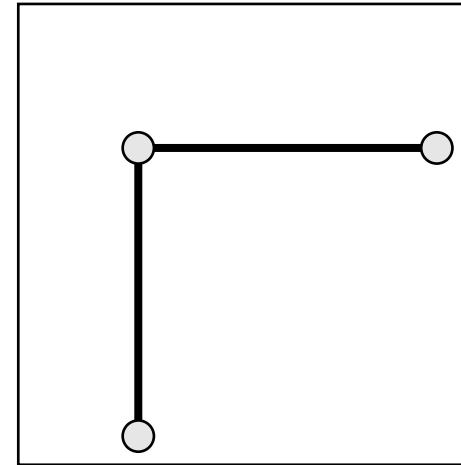


Depth
Map

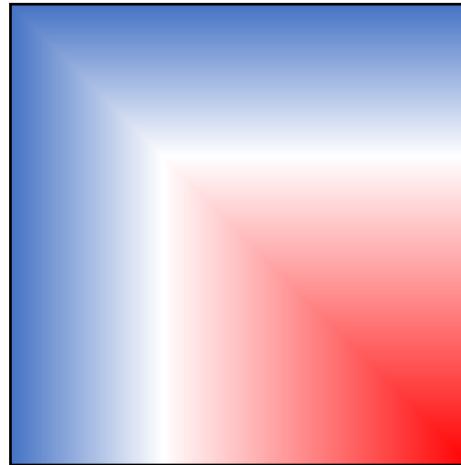
Voxel
Grid



Pointcloud



Mesh



Implicit
Surface

Last Time: Neural Radiance Fields (NeRF)



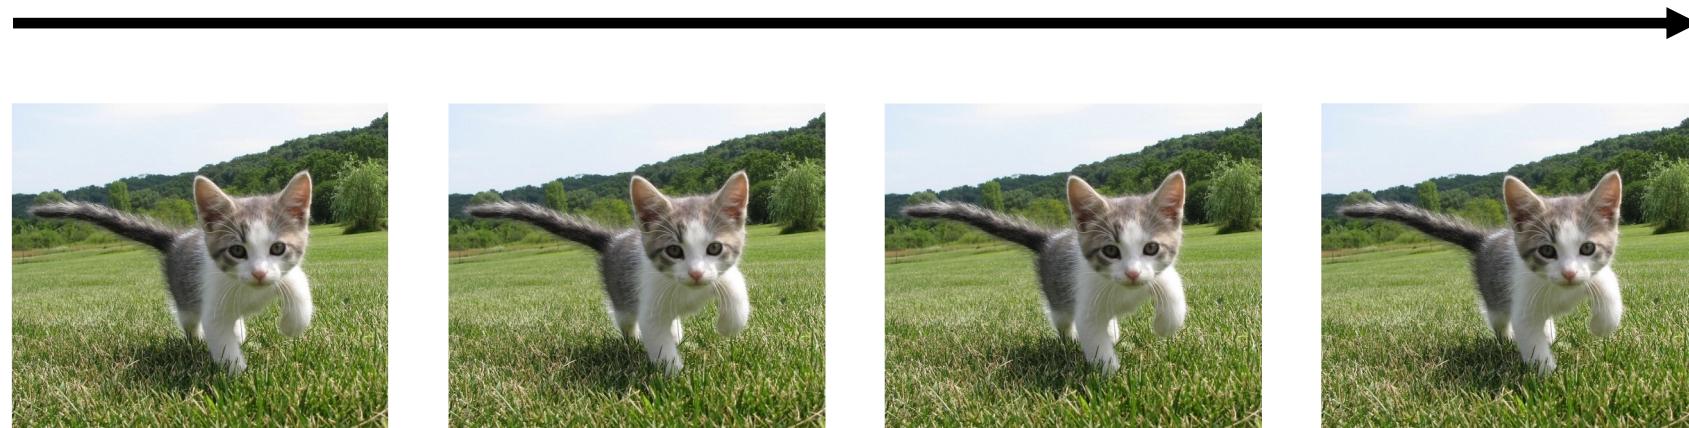
Mildenhall et al, "Representing Scenes as Neural Radiance Fields for View Synthesis", ECCV 2020

Today:
Videos

Today: Video = 2D + Time

A video is a **sequence** of images

4D tensor: $T \times 3 \times H \times W$
(or $3 \times T \times H \times W$)



This image is [CC0 public domain](#)

Example task: Video Classification



Input video:

$T \times 3 \times H \times W$

Swimming
Running
Jumping
Eating
Standing

[Running video](#) is in the [public domain](#)

Example task: Video Classification



Images: Recognize **objects**



Dog
Cat
Fish
Truck



Videos: Recognize **actions**



Swimming
Running
Jumping
Eating
Standing

[Running video](#) is in the [public domain](#)

Problem: Videos are big!

Videos are ~30 frames per second (fps)



Size of uncompressed video
(3 bytes per pixel):

SD (640 x 480): **~1.5 GB per minute**

HD (1920 x 1080): **~10 GB per minute**

Input video:

$T \times 3 \times H \times W$

Problem: Videos are big!



Input video:

$T \times 3 \times H \times W$

Videos are ~30 frames per second (fps)

Size of uncompressed video
(3 bytes per pixel):

SD (640 x 480): **~1.5 GB per minute**

HD (1920 x 1080): **~10 GB per minute**

Solution: Train on short **clips**: low
fps and low spatial resolution
e.g. $T = 16$, $H=W=112$
(3.2 seconds at 5 fps, 588 KB)

Training on Clips

Raw video: Long, high FPS



Training on Clips

Raw video: Long, high FPS



Training: Train model to classify short **clips** with low FPS



Training on Clips

Raw video: Long, high FPS



Training: Train model to classify short **clips** with low FPS



Testing: Run model on different clips, average predictions

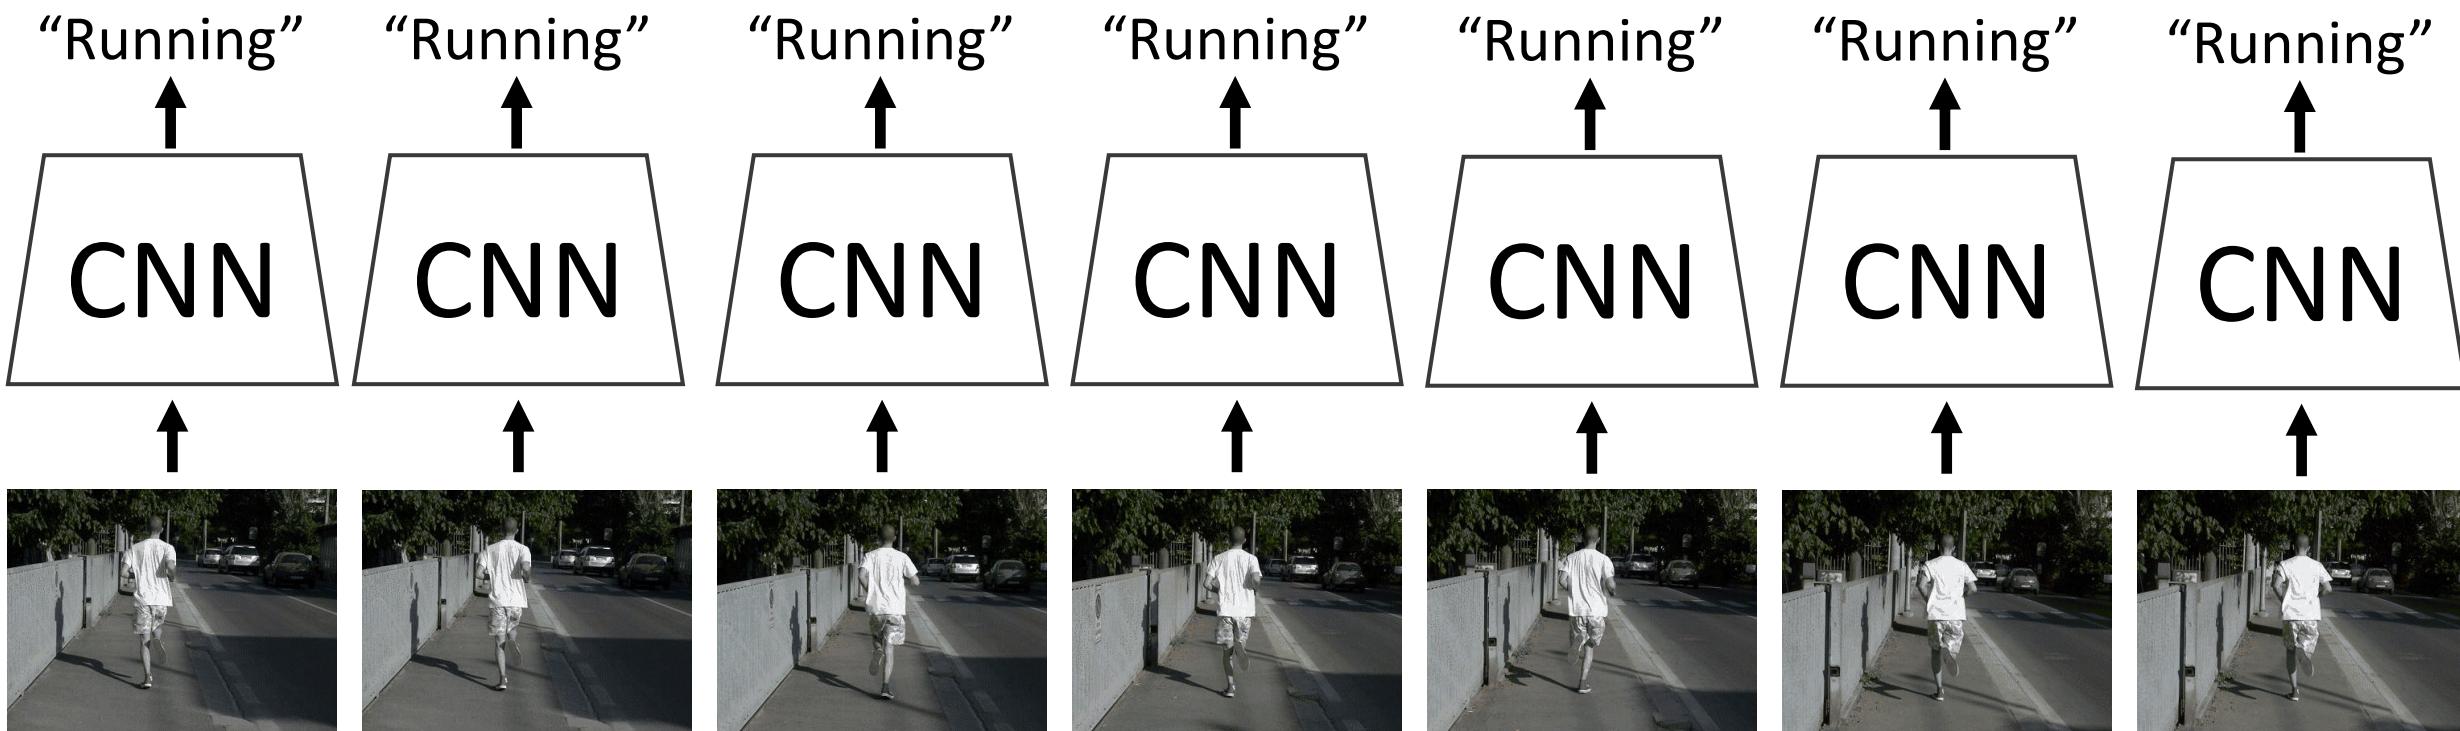


Video Classification: Single-Frame CNN

Simple idea: train normal 2D CNN to classify video frames independently!

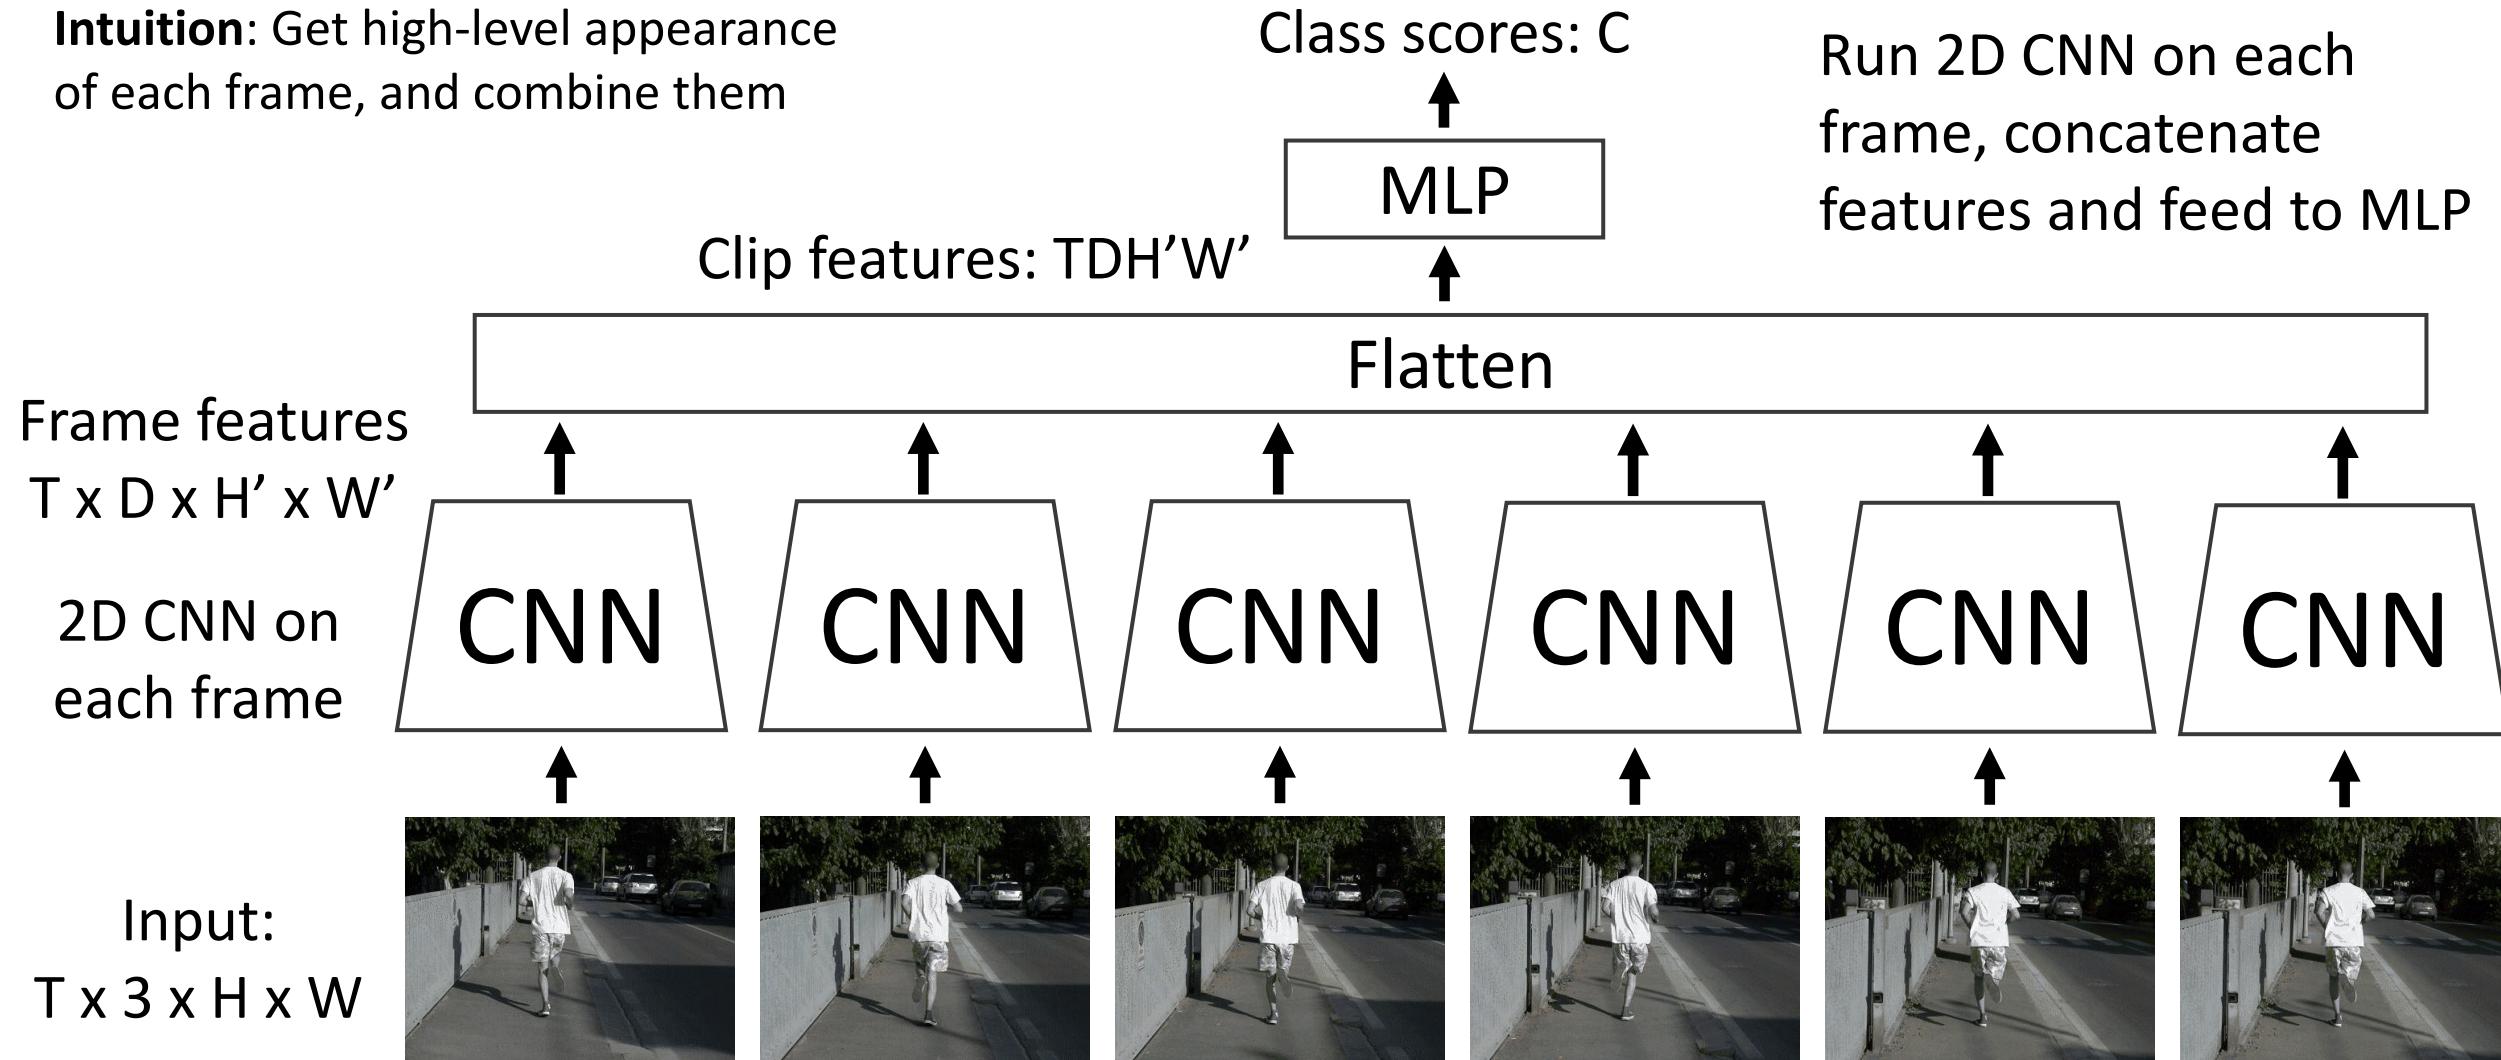
(Average predicted probs at test-time)

Often a **very** strong baseline for video classification



Video Classification: Late Fusion (with FC layers)

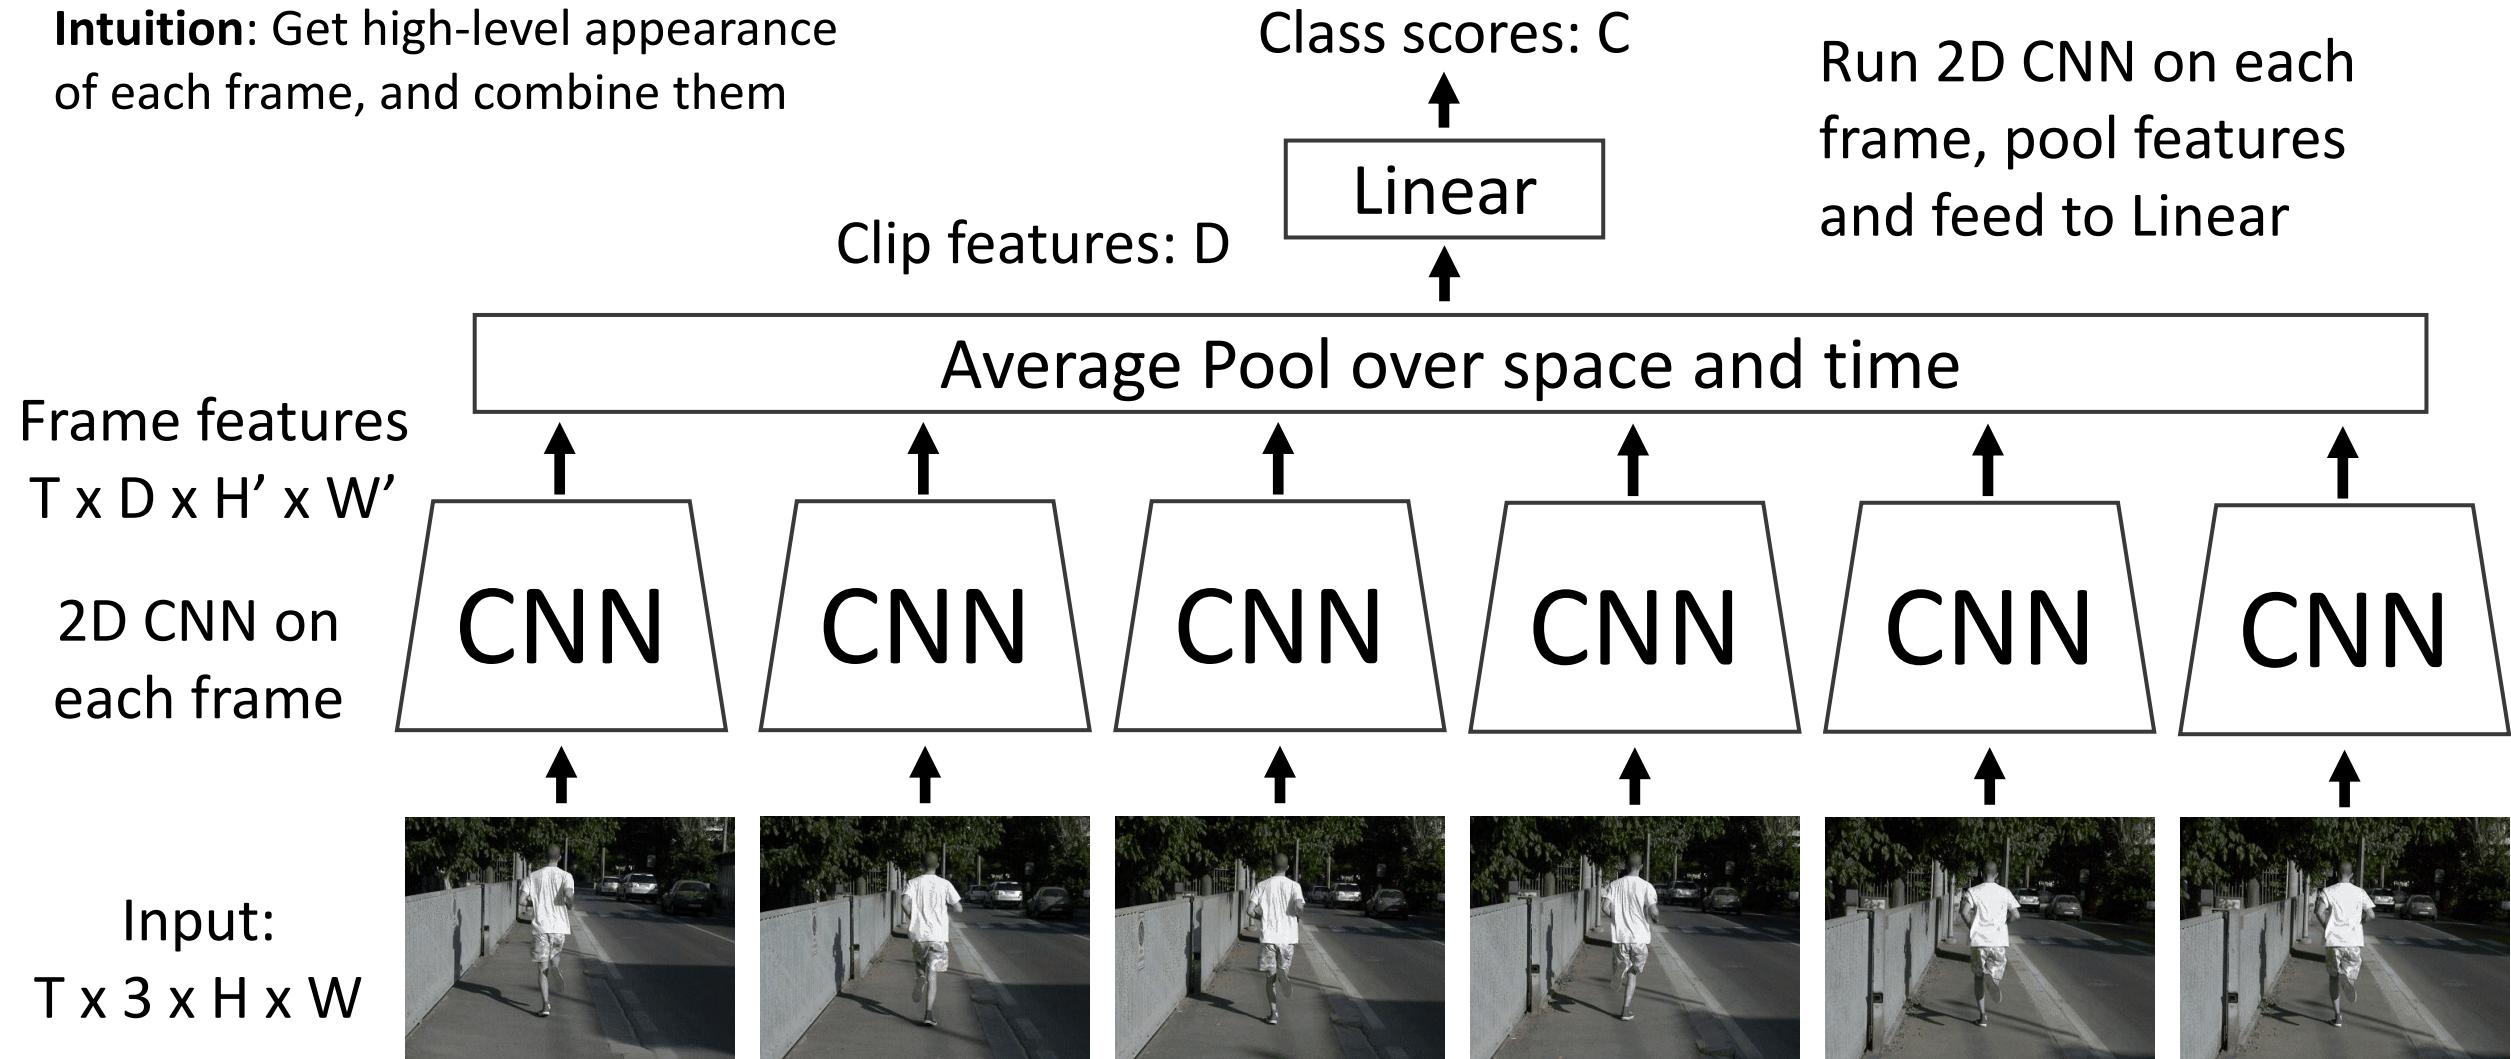
Intuition: Get high-level appearance of each frame, and combine them



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

Video Classification: Late Fusion (with pooling)

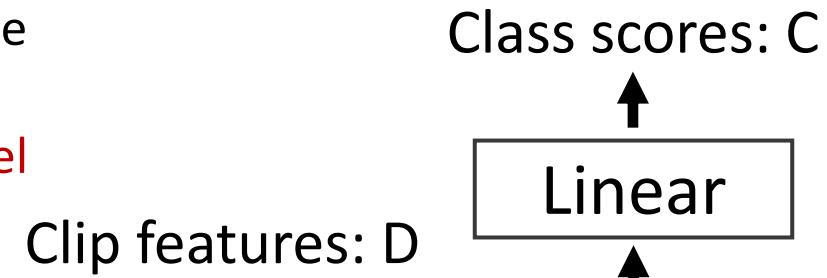
Intuition: Get high-level appearance of each frame, and combine them



Video Classification: Late Fusion (with pooling)

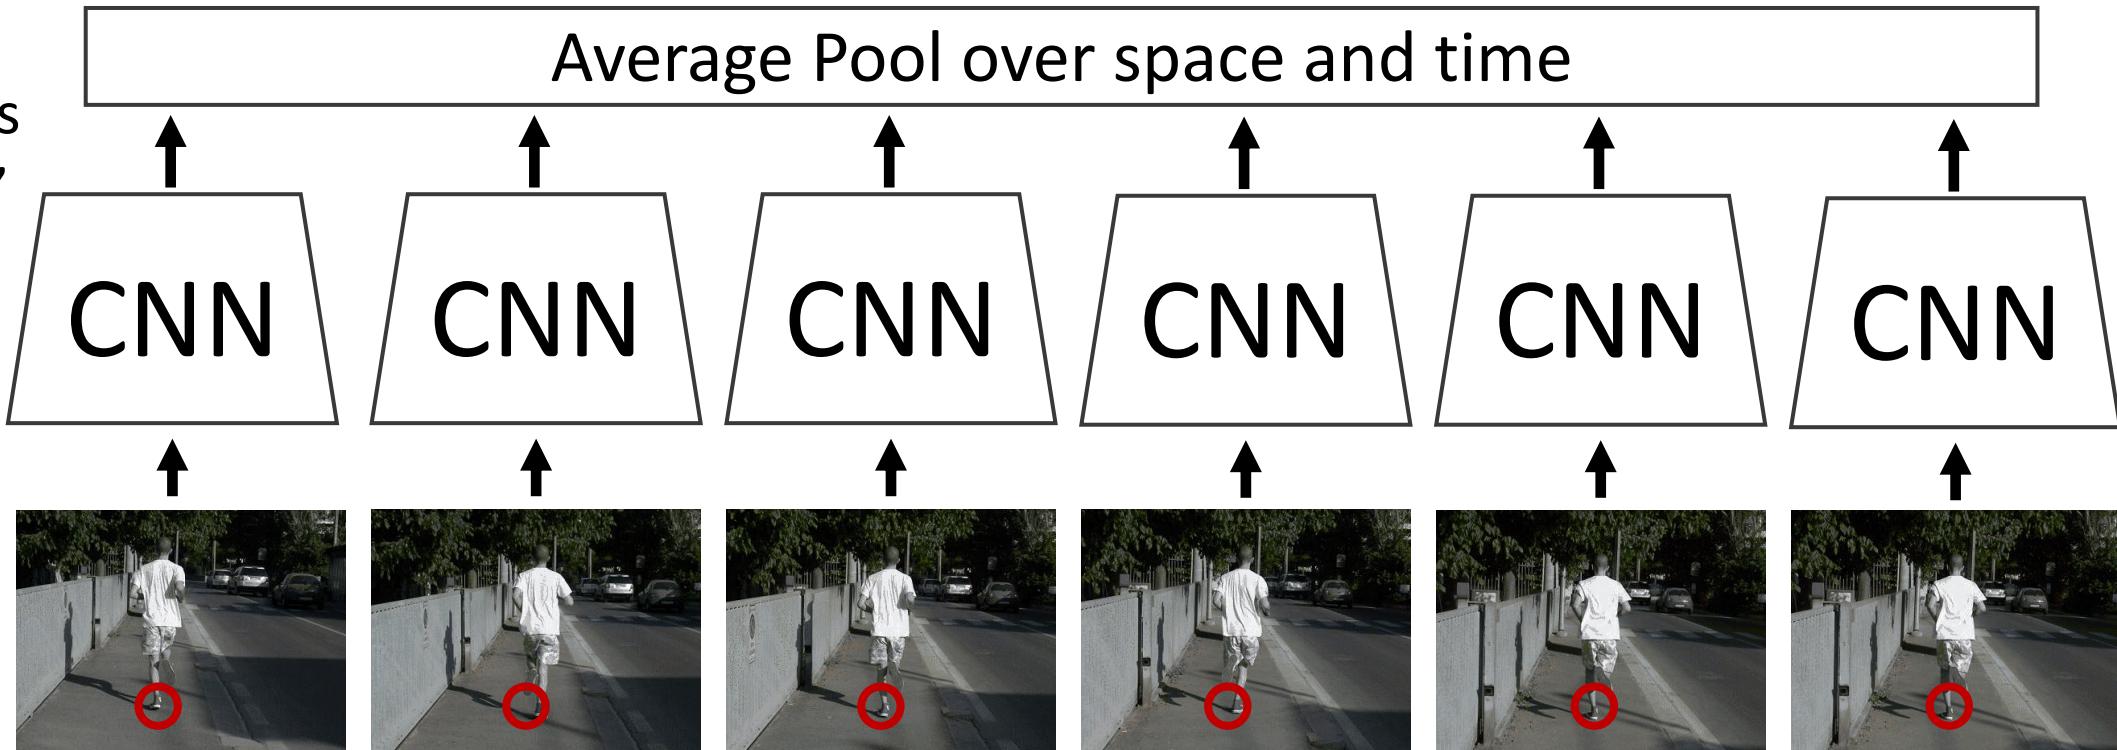
Intuition: Get high-level appearance of each frame, and combine them

Problem: Hard to compare low-level motion between frames



Run 2D CNN on each frame, pool features and feed to Linear

Frame features
 $T \times D \times H' \times W'$
2D CNN on each frame



Input:
 $T \times 3 \times H \times W$

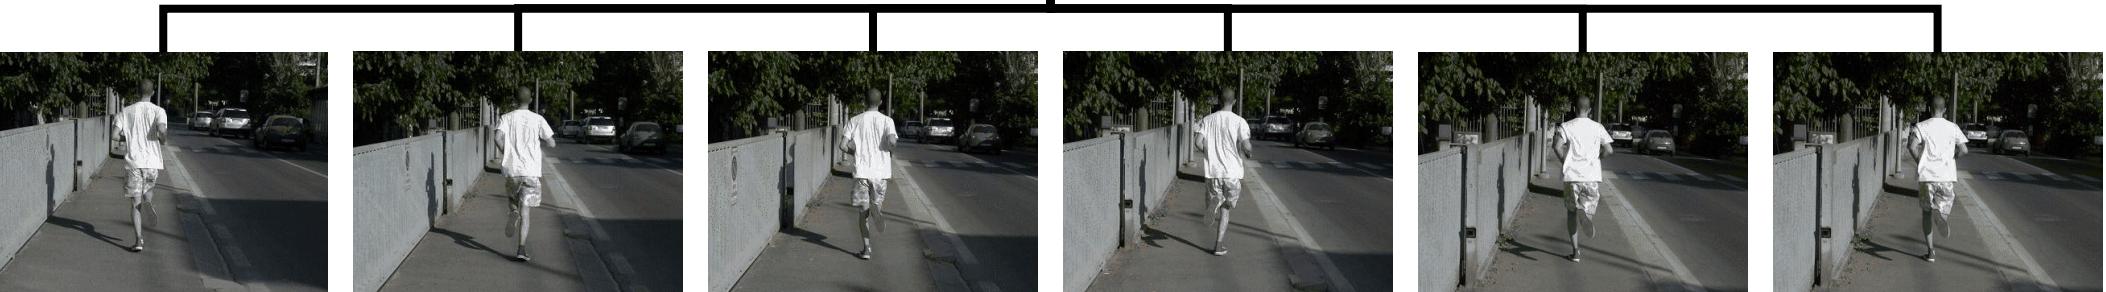
Video Classification: Early Fusion

Intuition: Compare frames with very first conv layer, after that normal 2D CNN

First 2D convolution collapses all temporal information:
Input: $3T \times H \times W$
Output: $D \times H \times W$

Reshape:
 $3T \times H \times W$

Input:
 $T \times 3 \times H \times W$



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

Video Classification: Early Fusion

Intuition: Compare frames with very first conv layer, after that normal 2D CNN

Problem: One layer of temporal processing may not be enough!

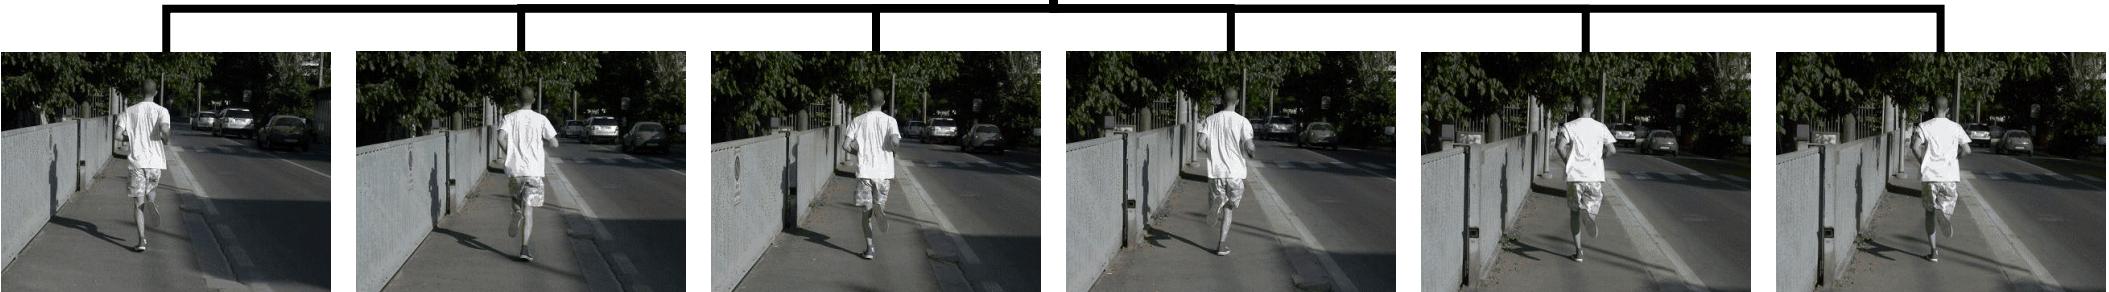
First 2D convolution collapses all temporal information:

Input: $3T \times H \times W$

Output: $D \times H \times W$

Reshape:
 $3T \times H \times W$

Input:
 $T \times 3 \times H \times W$

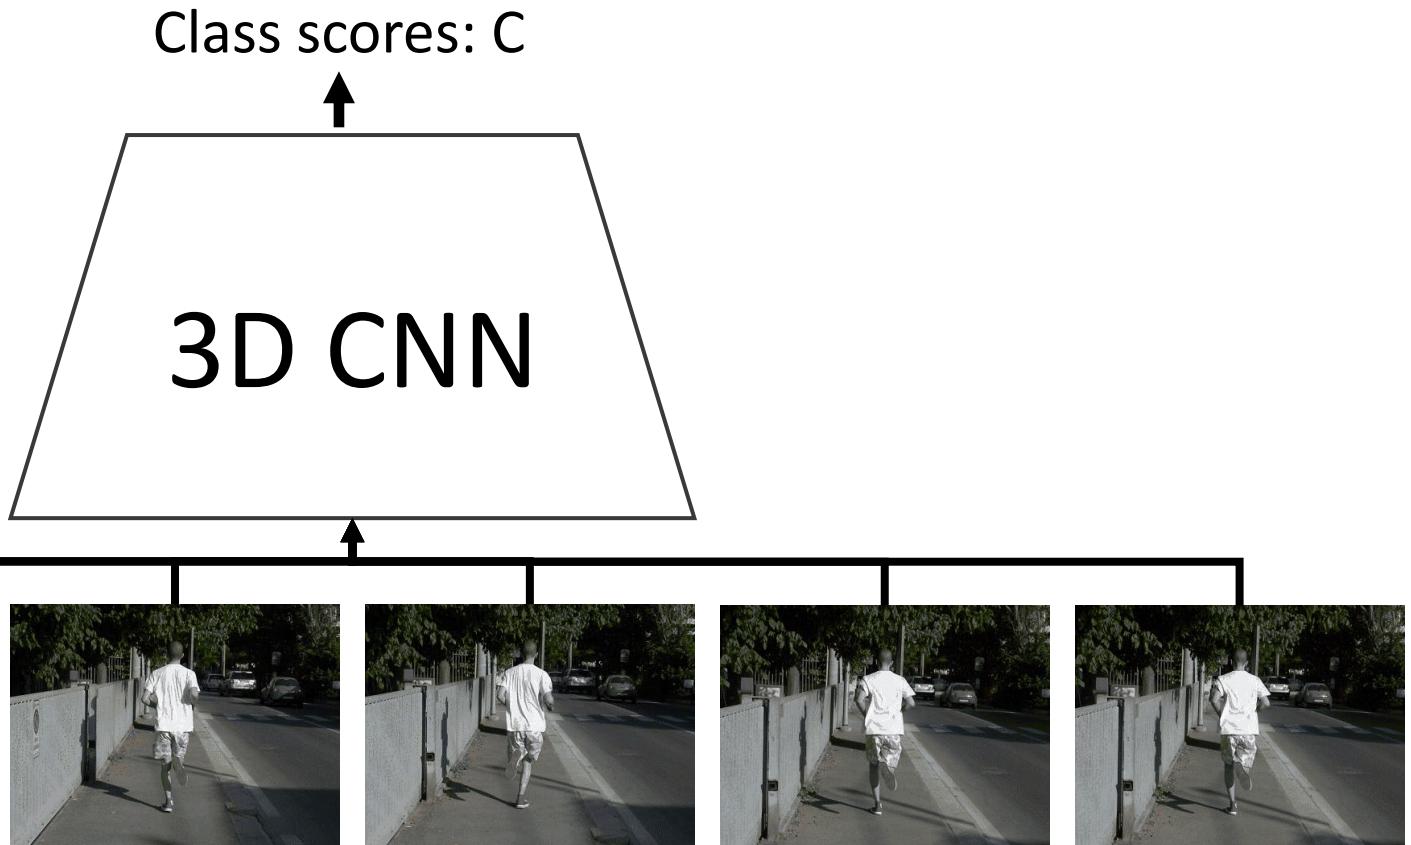


Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

Video Classification: 3D CNN

Intuition: Use 3D versions of convolution and pooling to slowly fuse temporal information over the course of the network

Each layer in the network is a 4D tensor: $D \times T \times H \times W$
Use 3D conv and 3D pooling operations



Ji et al, "3D Convolutional Neural Networks for Human Action Recognition", TPAMI 2010 ; Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

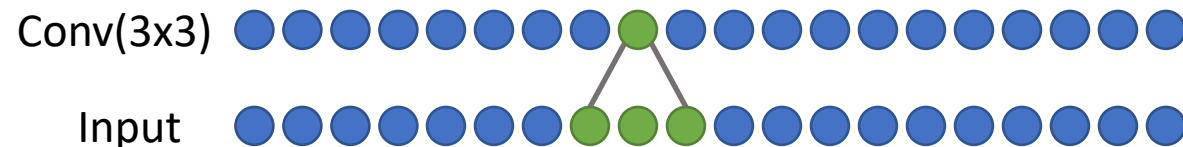
Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3

(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3

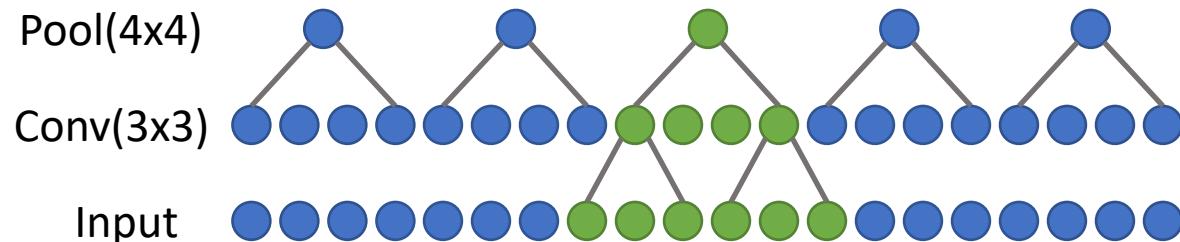


(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$1 \times 3 \times 3$
Pool2D(4x4)	$12 \times 20 \times 16 \times 16$	$1 \times 6 \times 6$



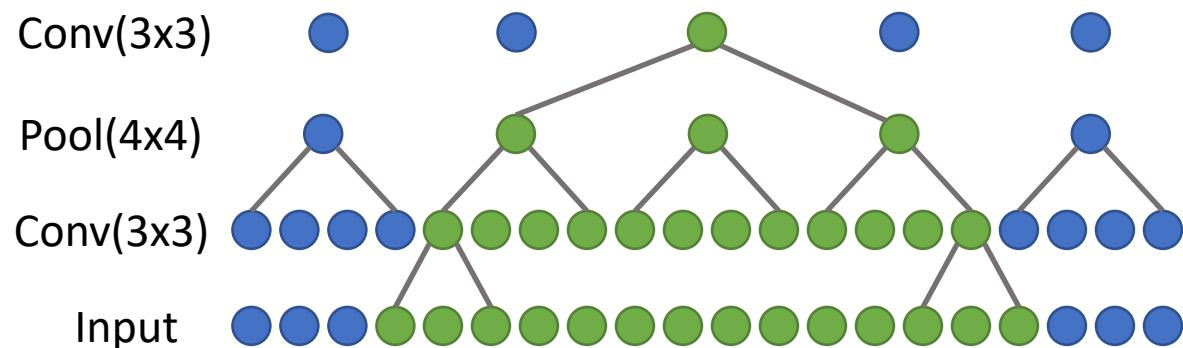
(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$1 \times 3 \times 3$
Pool2D(4x4)	$12 \times 20 \times 16 \times 16$	$1 \times 6 \times 6$
Conv2D(3x3, 12->24)	$24 \times 20 \times 16 \times 16$	$1 \times 14 \times 14$

Build slowly in space



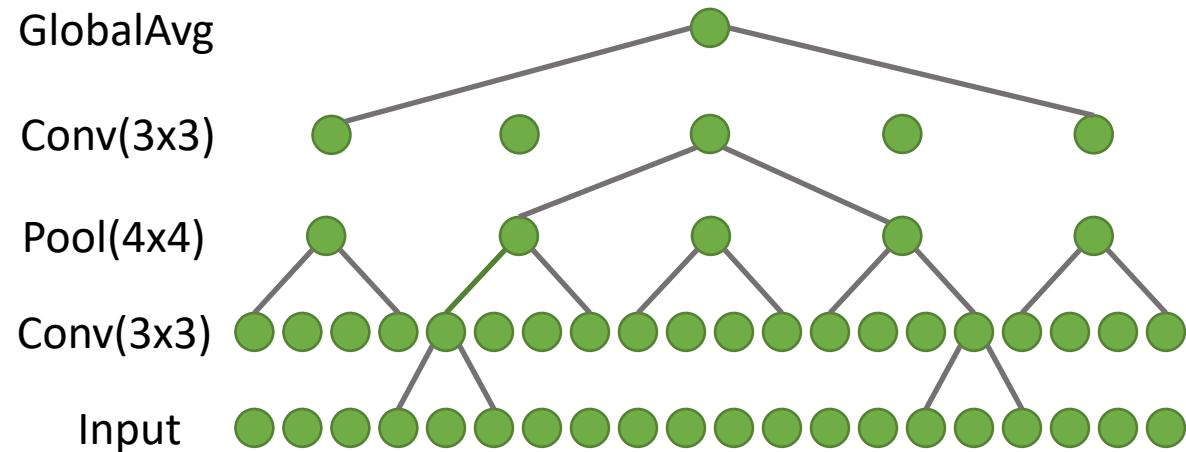
(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$1 \times 3 \times 3$
Pool2D(4x4)	$12 \times 20 \times 16 \times 16$	$1 \times 6 \times 6$
Conv2D(3x3, 12->24)	$24 \times 20 \times 16 \times 16$	$1 \times 14 \times 14$
GlobalAvgPool	$24 \times 1 \times 1 \times 1$	$20 \times 64 \times 64$

Build slowly in space,
All-at-once in time at end



(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3*10->12)	12 x 64 x 64	20 x 3 x 3
Pool2D(4x4)	12 x 16 x 16	20 x 6 x 6
Conv2D(3x3, 12->24)	24 x 16 x 16	20 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

Build slowly in space,
All-at-once in time at end

Early
Fusion

Build slowly in space,
All-at-once in time at start

(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

Late
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$1 \times 3 \times 3$
Pool2D(4x4)	$12 \times 20 \times 16 \times 16$	$1 \times 6 \times 6$
Conv2D(3x3, 12->24)	$24 \times 20 \times 16 \times 16$	$1 \times 14 \times 14$
GlobalAvgPool	$24 \times 1 \times 1 \times 1$	$20 \times 64 \times 64$
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3*10->12)	$12 \times 64 \times 64$	$20 \times 3 \times 3$
Pool2D(4x4)	$12 \times 16 \times 16$	$20 \times 6 \times 6$
Conv2D(3x3, 12->24)	$24 \times 16 \times 16$	$20 \times 14 \times 14$
GlobalAvgPool	$24 \times 1 \times 1$	$20 \times 64 \times 64$
Input	$3 \times 20 \times 64 \times 64$	
Conv3D(3x3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$3 \times 3 \times 3$
Pool3D(4x4x4)	$12 \times 5 \times 16 \times 16$	$6 \times 6 \times 6$
Conv3D(3x3x3, 12->24)	$24 \times 5 \times 16 \times 16$	$14 \times 14 \times 14$
GlobalAvgPool	$24 \times 1 \times 1$	$20 \times 64 \times 64$

Build slowly in space,
All-at-once in time at end

Early
Fusion

Build slowly in space,
All-at-once in time at start

3D CNN

Build slowly in space,
Build slowly in time
"Slow Fusion"

(Small example
architectures,
in practice
much bigger)

Early Fusion vs Late Fusion vs 3D CNN

What is the difference?

Late Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64

Build slowly in space,
All-at-once in time at end

Early Fusion

Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3*10->12)	12 x 64 x 64	20 x 3 x 3
Pool2D(4x4)	12 x 16 x 16	20 x 6 x 6
Conv2D(3x3, 12->24)	24 x 16 x 16	20 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

Build slowly in space,
All-at-once in time at start

3D CNN

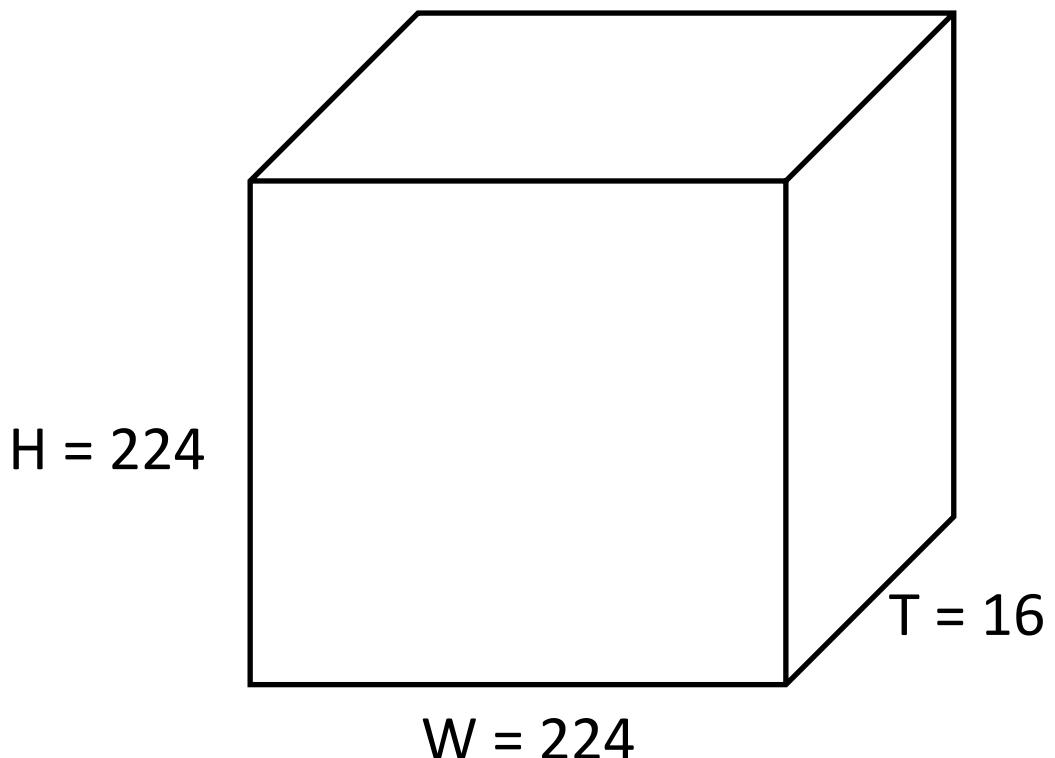
Input	3 x 20 x 64 x 64	
Conv3D(3x3x3, 3->12)	12 x 20 x 64 x 64	3 x 3 x 3
Pool3D(4x4x4)	12 x 5 x 16 x 16	6 x 6 x 6
Conv3D(3x3x3, 12->24)	24 x 5 x 16 x 16	14 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

Build slowly in space,
Build slowly in time
"Slow Fusion"

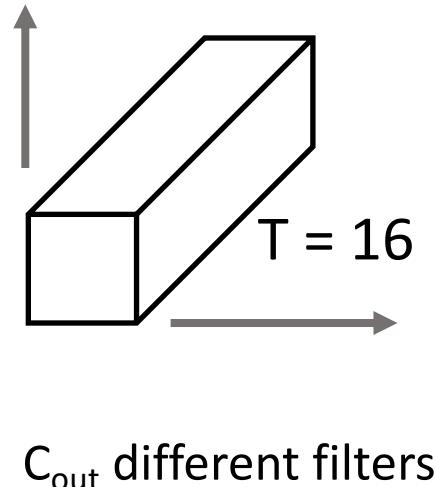
(Small example architectures,
in practice much bigger)

2D Conv (Early Fusion) vs 3D Conv (3D CNN)

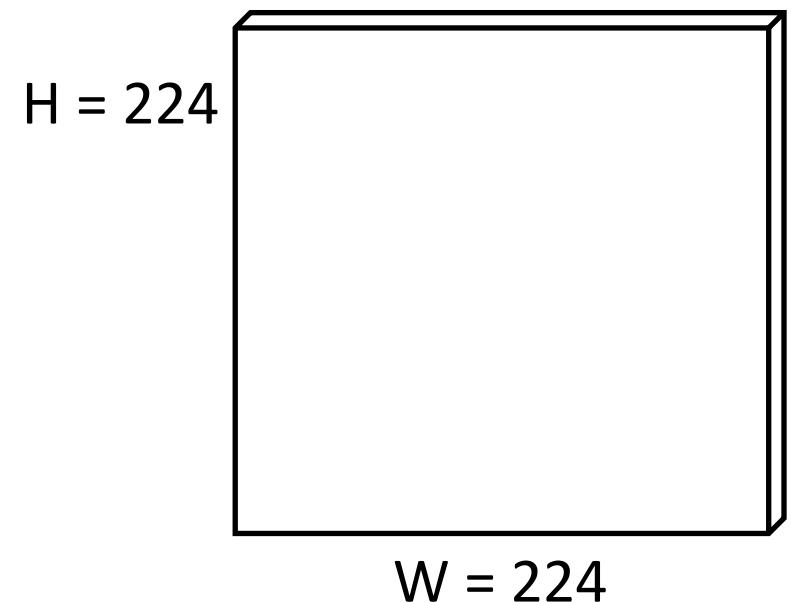
Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)



Weight:
 $C_{out} \times C_{in} \times T \times 3 \times 3$
Slide over x and y

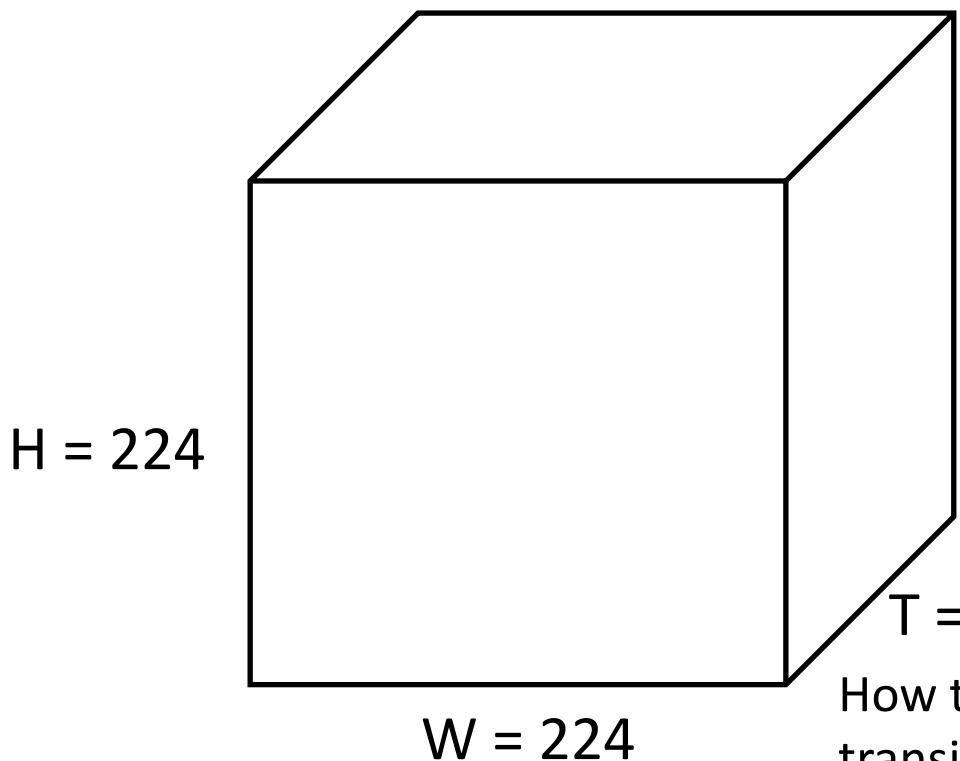


Output:
 $C_{out} \times H \times W$
2D grid with C_{out} -dim
feat at each point

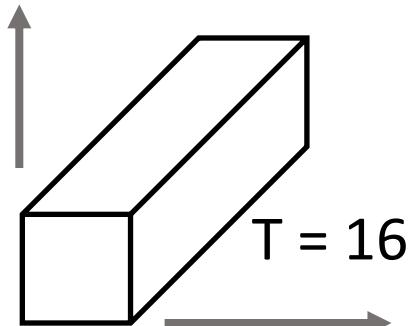


2D Conv (Early Fusion) vs 3D Conv (3D CNN)

Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)

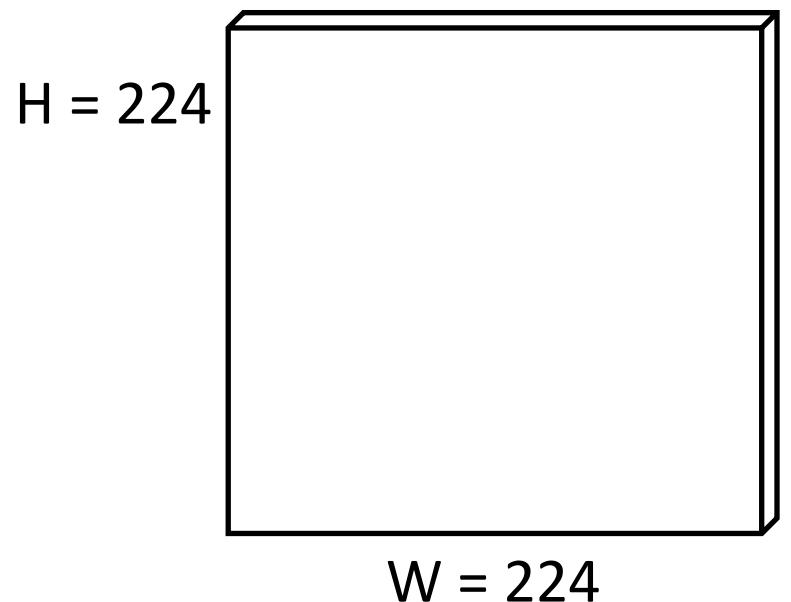


Weight:
 $C_{out} \times C_{in} \times T \times 3 \times 3$
Slide over x and y



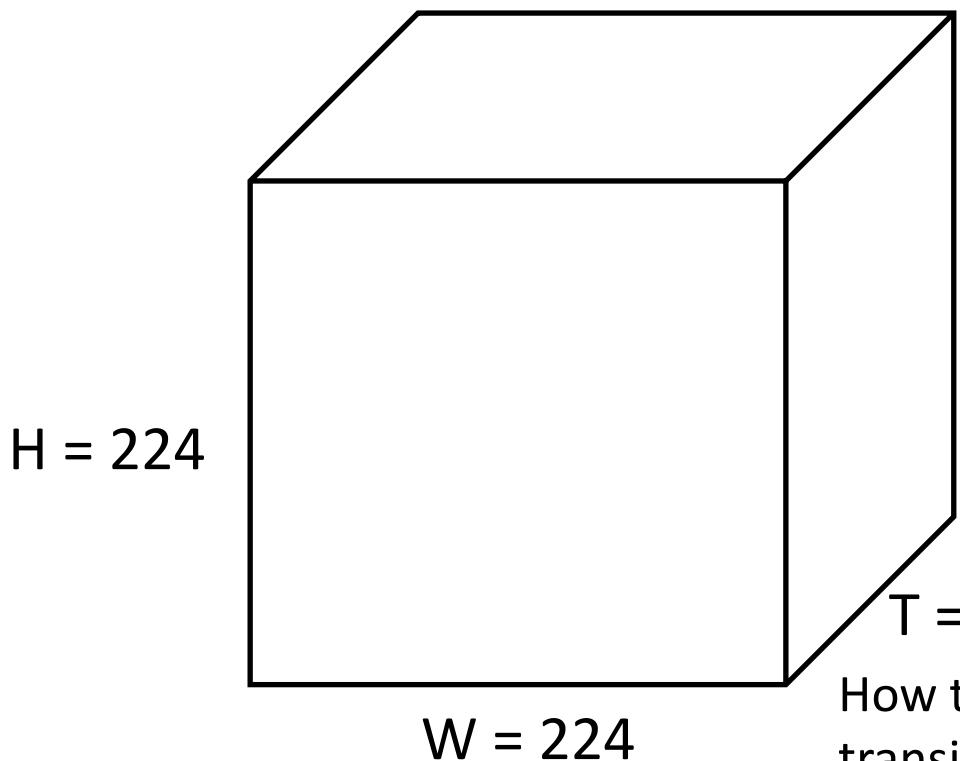
C_{out} different filters
How to recognize **blue** to **orange**
transitions anywhere in space and time?

Output:
 $C_{out} \times H \times W$
2D grid with C_{out} -dim
feat at each point



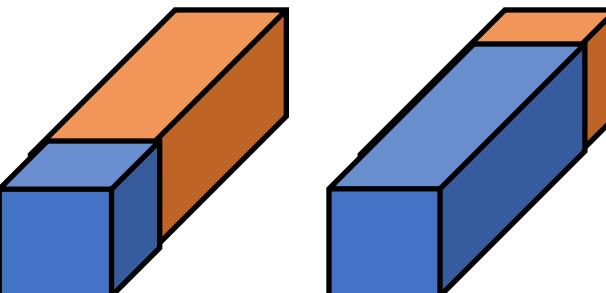
2D Conv (Early Fusion) vs 3D Conv (3D CNN)

Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)



Weight:
 $C_{out} \times C_{in} \times T \times 3 \times 3$
Slide over x and y

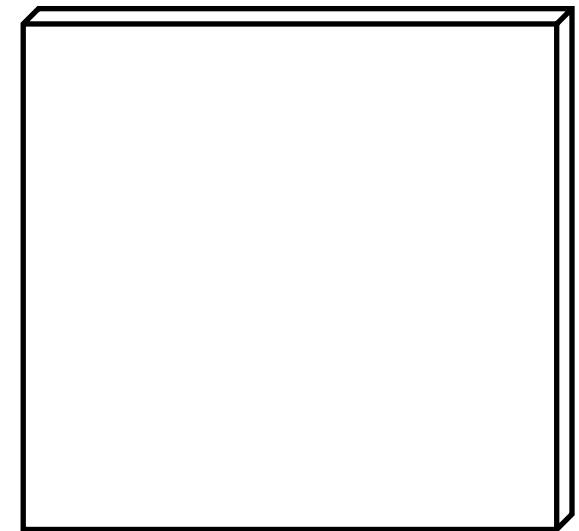
No temporal shift-invariance! Needs
to learn separate filters for the same
motion at different times in the clip



C_{out} different filters

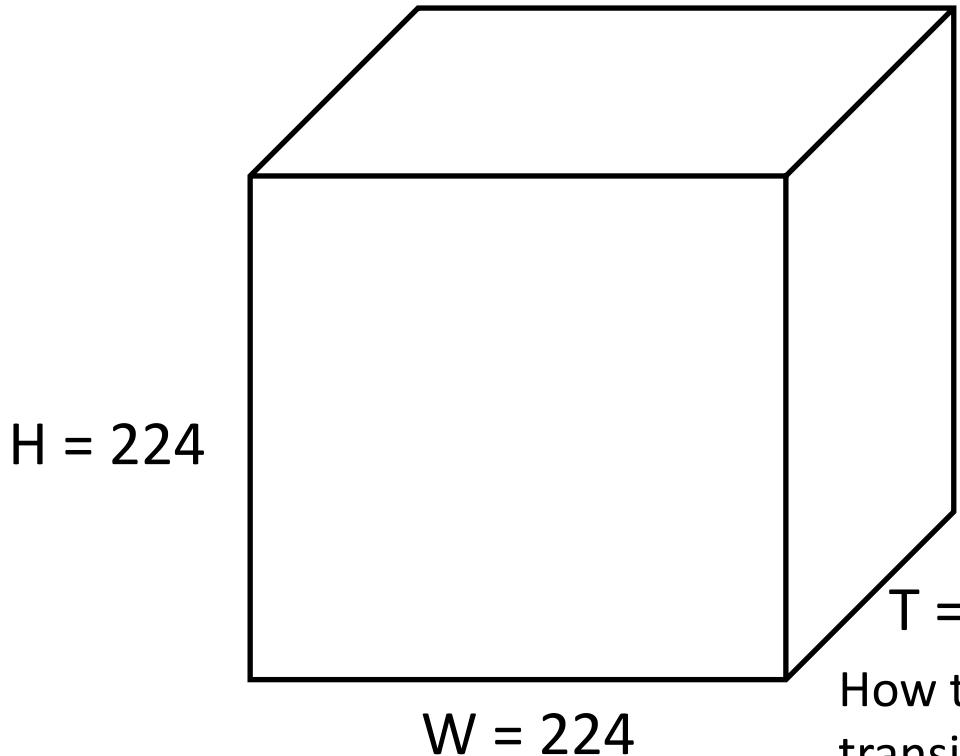
How to recognize **blue** to **orange**
transitions anywhere in space and time?

Output:
 $C_{out} \times H \times W$
2D grid with C_{out} -dim
feat at each point



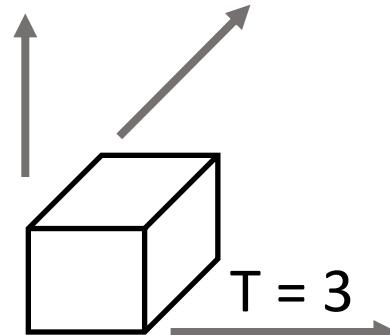
2D Conv (Early Fusion) vs 3D Conv (3D CNN)

Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)



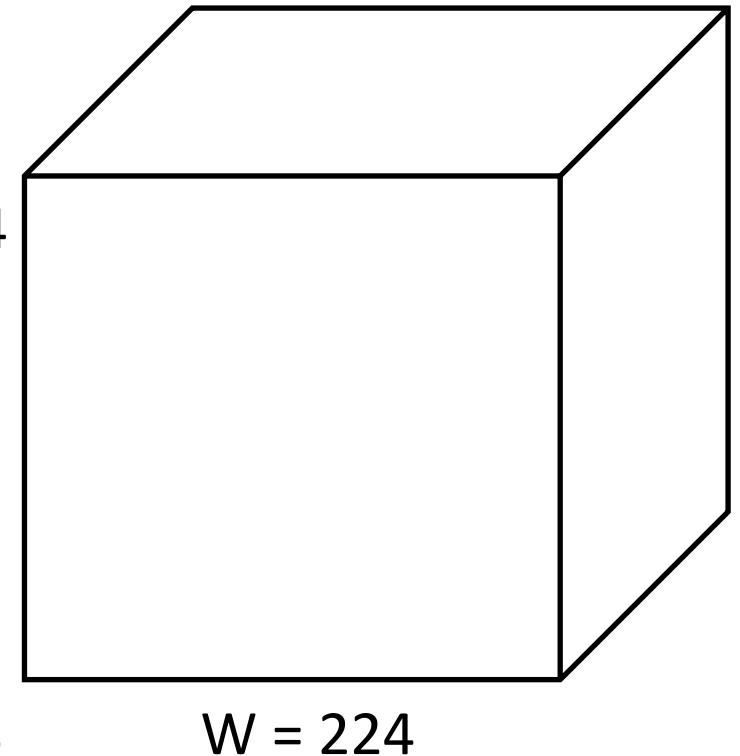
How to recognize **blue** to **orange**
transitions anywhere in space and time?

Weight:
 $C_{out} \times C_{in} \times 3 \times 3 \times 3$
Slide over x and y



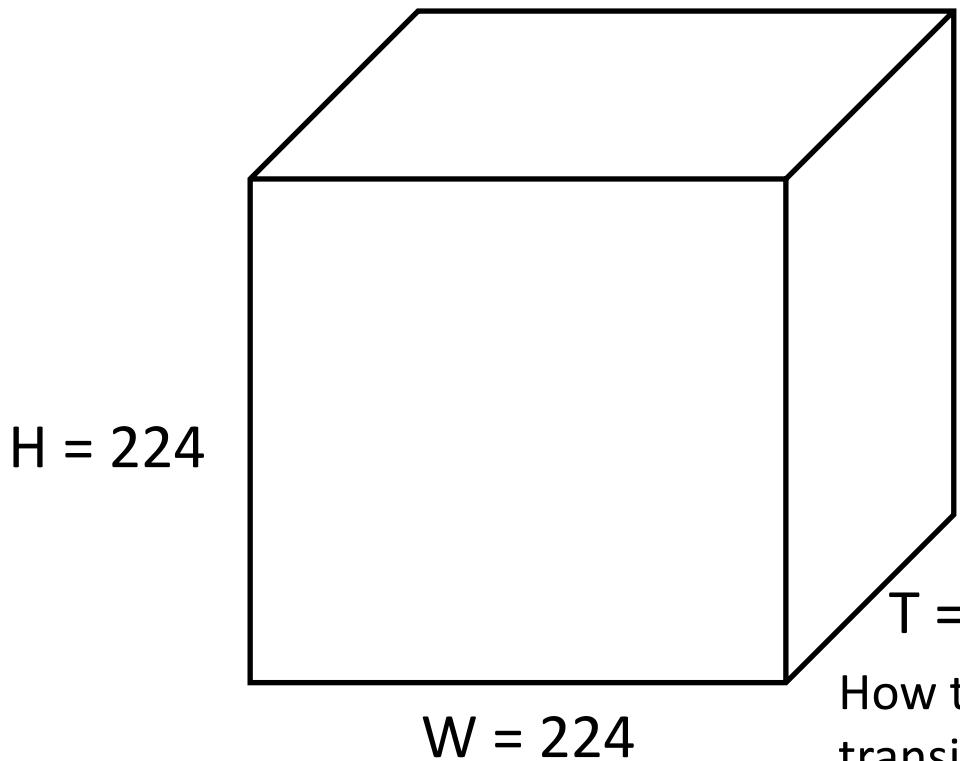
C_{out} different filters

Output:
 $C_{out} \times T \times H \times W$
3D grid with C_{out} -dim
feat at each point



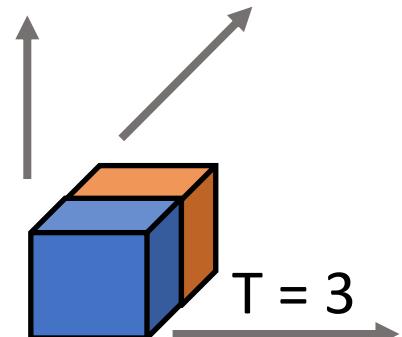
2D Conv (Early Fusion) vs 3D Conv (3D CNN)

Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)



How to recognize **blue** to **orange**
transitions anywhere in space and time?

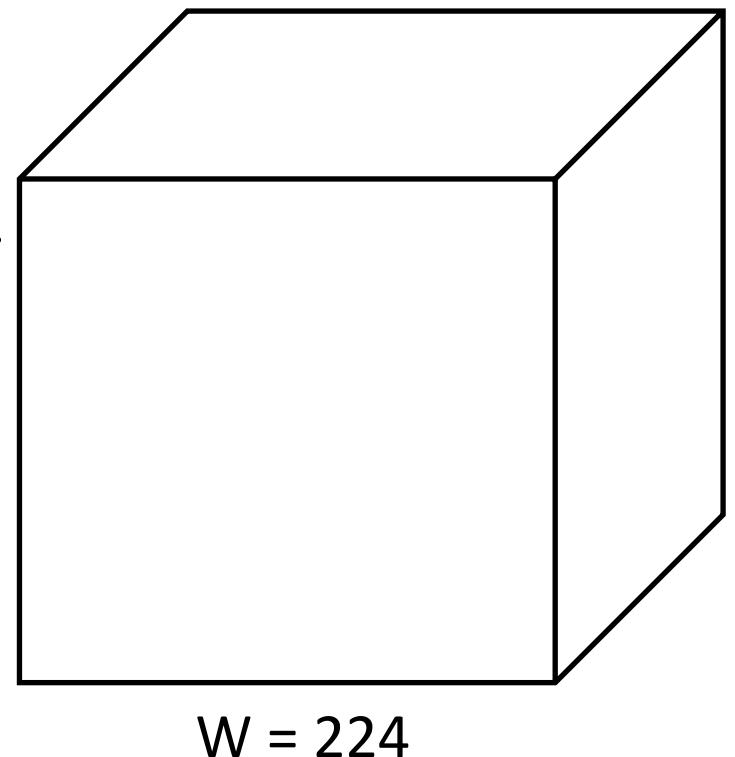
Weight:
 $C_{out} \times C_{in} \times 3 \times 3 \times 3$
Slide over x and y



Temporal shift-invariant since
each filter slides over time!

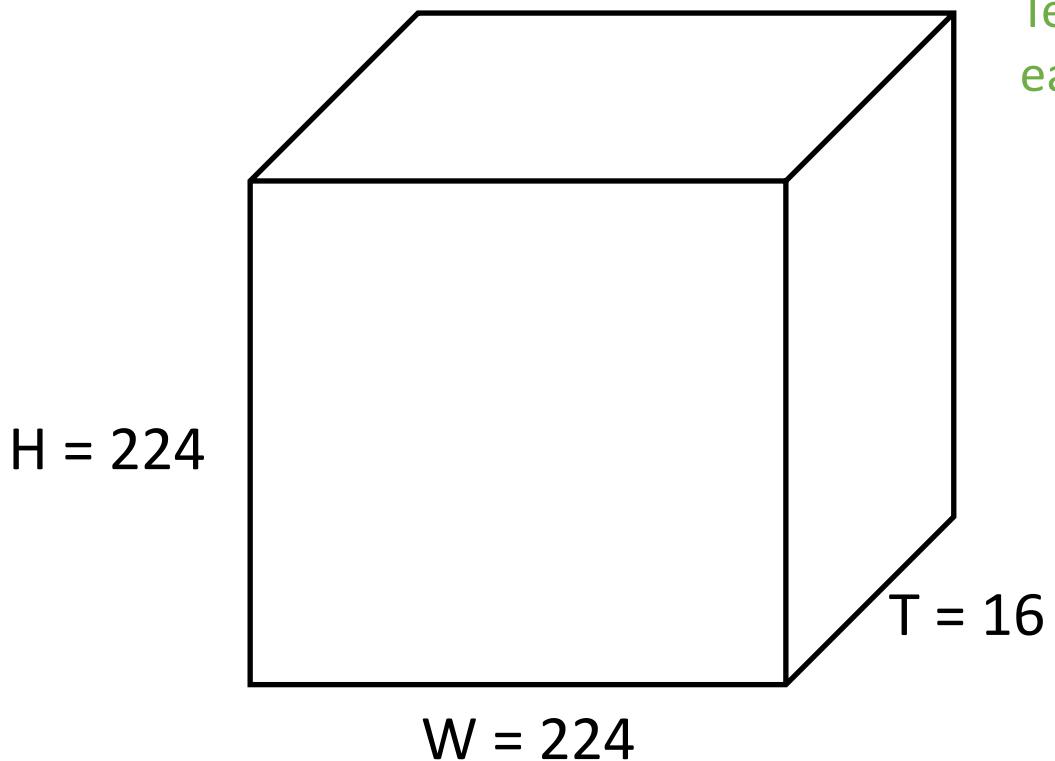
C_{out} different filters

Output:
 $C_{out} \times T \times H \times W$
3D grid with C_{out} -dim
feat at each point



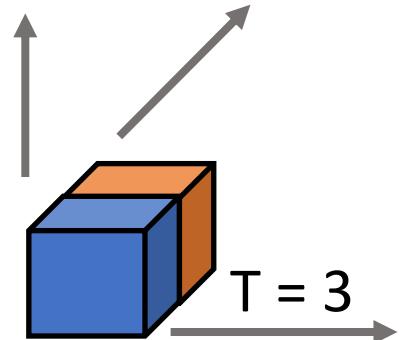
2D Conv (Early Fusion) vs 3D Conv (3D CNN)

Input: $C_{in} \times T \times H \times W$
(3D grid with C_{in} -dim
feat at each point)



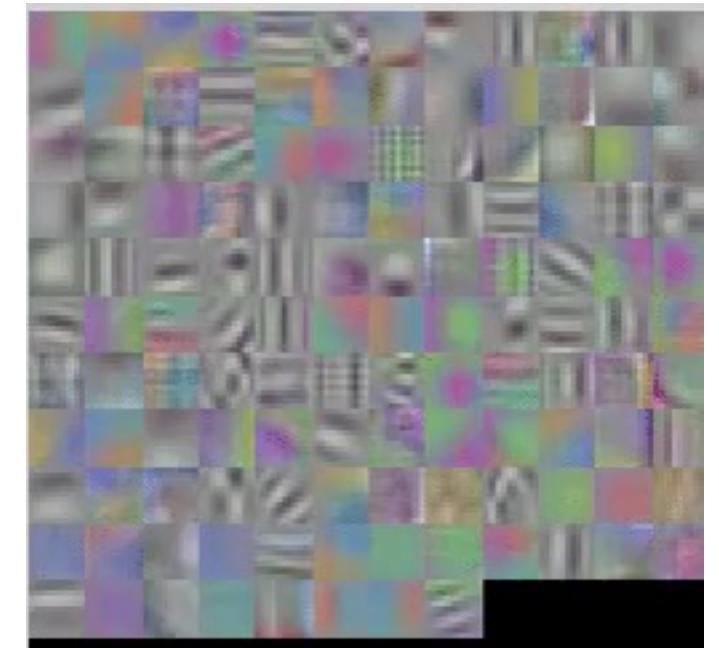
Weight:
 $C_{out} \times C_{in} \times 3 \times 3 \times 3$
Slide over x and y

Temporal shift-invariant since
each filter slides over time!



C_{out} different filters

First-layer filters have shape
3 (RGB) \times 4 (frames) \times 5 \times 5 (space)
Can visualize as video clips!



Karpathy et al, "Large-scale Video Classification
with Convolutional Neural Networks", CVPR 2014

Example Video Dataset: Sports-1M

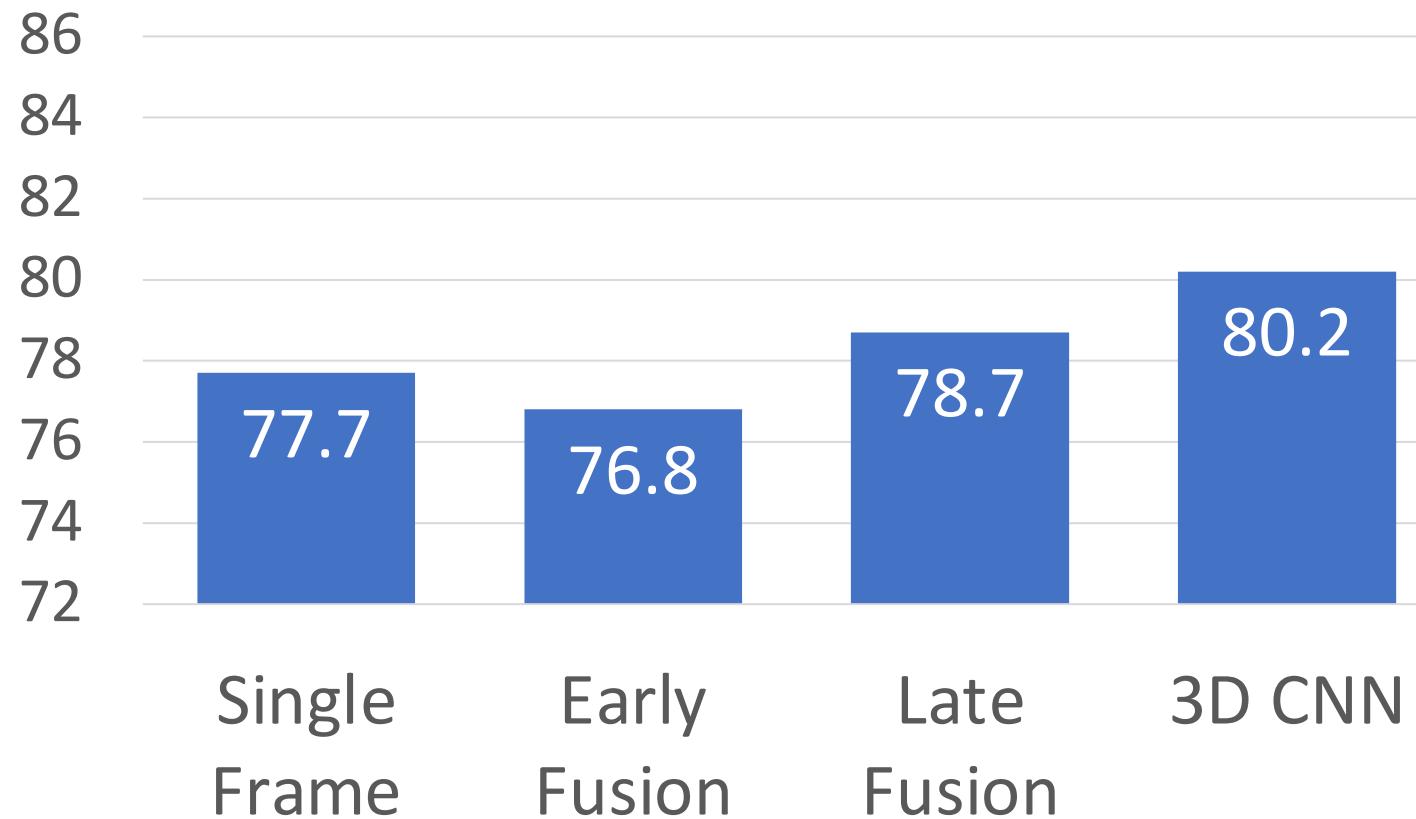


1 million YouTube videos
annotated with labels for
487 different types of sports

Ground Truth
Correct prediction
Incorrect prediction

Early Fusion vs Late Fusion vs 3D CNN

Sports-1M Top-5 Accuracy



Single Frame model works well – always try this first!

3D CNNs have improved a lot since 2014!

C3D: The VGG of 3D CNNs

3D CNN that uses all $3 \times 3 \times 3$ conv and
 $2 \times 2 \times 2$ pooling
(except Pool1 which is $1 \times 2 \times 2$)

Released model pretrained on Sports-
1M: Many people used this as a video
feature extractor

Layer	Size
Input	$3 \times 16 \times 112 \times 112$
Conv1 ($3 \times 3 \times 3$)	$64 \times 16 \times 112 \times 112$
Pool1 ($1 \times 2 \times 2$)	$64 \times 16 \times 56 \times 56$
Conv2 ($3 \times 3 \times 3$)	$128 \times 16 \times 56 \times 56$
Pool2 ($2 \times 2 \times 2$)	$128 \times 8 \times 28 \times 28$
Conv3a ($3 \times 3 \times 3$)	$256 \times 8 \times 28 \times 28$
Conv3b ($3 \times 3 \times 3$)	$256 \times 8 \times 28 \times 28$
Pool3 ($2 \times 2 \times 2$)	$256 \times 4 \times 14 \times 14$
Conv4a ($3 \times 3 \times 3$)	$512 \times 4 \times 14 \times 14$
Conv4b ($3 \times 3 \times 3$)	$512 \times 4 \times 14 \times 14$
Pool4 ($2 \times 2 \times 2$)	$512 \times 2 \times 7 \times 7$
Conv5a ($3 \times 3 \times 3$)	$512 \times 2 \times 7 \times 7$
Conv5b ($3 \times 3 \times 3$)	$512 \times 2 \times 7 \times 7$
Pool5	$512 \times 1 \times 3 \times 3$
FC6	4096
FC7	4096
FC8	C

C3D: The VGG of 3D CNNs

3D CNN that uses all $3 \times 3 \times 3$ conv and
 $2 \times 2 \times 2$ pooling
(except Pool1 which is $1 \times 2 \times 2$)

Released model pretrained on Sports-1M: Many people used this as a video feature extractor

Problem: $3 \times 3 \times 3$ conv is very expensive!

AlexNet: 0.7 GFLOP

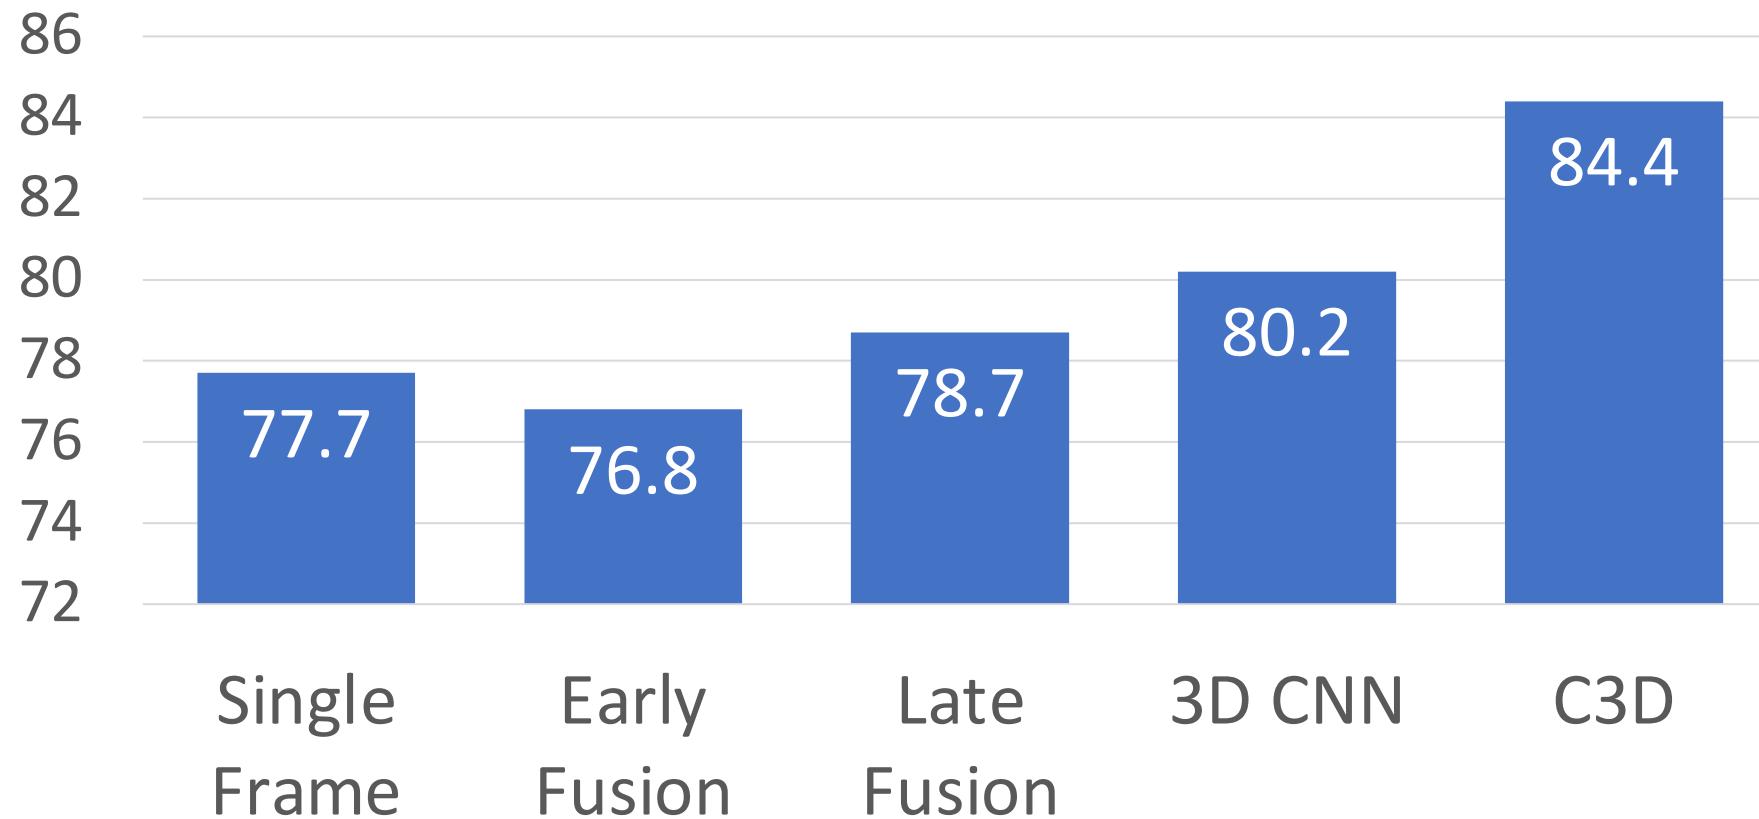
VGG-16: 13.6 GFLOP

C3D: **39.5 GFLOP (2.9x VGG!)**

Layer	Size	MFLOPs
Input	$3 \times 16 \times 112 \times 112$	
Conv1 ($3 \times 3 \times 3$)	$64 \times 16 \times 112 \times 112$	1.04
Pool1 ($1 \times 2 \times 2$)	$64 \times 16 \times 56 \times 56$	
Conv2 ($3 \times 3 \times 3$)	$128 \times 16 \times 56 \times 56$	11.10
Pool2 ($2 \times 2 \times 2$)	$128 \times 8 \times 28 \times 28$	
Conv3a ($3 \times 3 \times 3$)	$256 \times 8 \times 28 \times 28$	5.55
Conv3b ($3 \times 3 \times 3$)	$256 \times 8 \times 28 \times 28$	11.10
Pool3 ($2 \times 2 \times 2$)	$256 \times 4 \times 14 \times 14$	
Conv4a ($3 \times 3 \times 3$)	$512 \times 4 \times 14 \times 14$	2.77
Conv4b ($3 \times 3 \times 3$)	$512 \times 4 \times 14 \times 14$	5.55
Pool4 ($2 \times 2 \times 2$)	$512 \times 2 \times 7 \times 7$	
Conv5a ($3 \times 3 \times 3$)	$512 \times 2 \times 7 \times 7$	0.69
Conv5b ($3 \times 3 \times 3$)	$512 \times 2 \times 7 \times 7$	0.69
Pool5	$512 \times 1 \times 3 \times 3$	
FC6	4096	0.51
FC7	4096	0.45
FC8	C	0.05

Early Fusion vs Late Fusion vs 3D CNN

Sports-1M Top-5 Accuracy



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014
Tran et al, "Learning Spatiotemporal Features with 3D Convolutional Networks", ICCV 2015

Recognizing Actions from Motion

We can easily recognize actions using only **motion information**



Johansson, "Visual perception of biological motion and a model for its analysis." Perception & Psychophysics. 14(2):201-211. 1973.

Measuring Motion: Optical Flow

Image at frame t



Image at frame t+1

Simonyan and Zisserman, "Two-stream convolutional networks for action recognition in videos", NeurIPS 2014

Measuring Motion: Optical Flow

Image at frame t

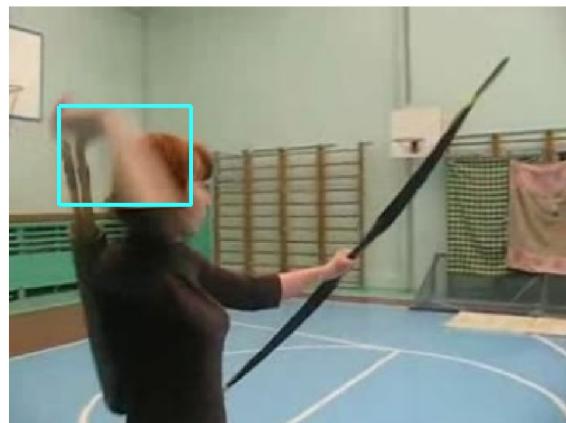
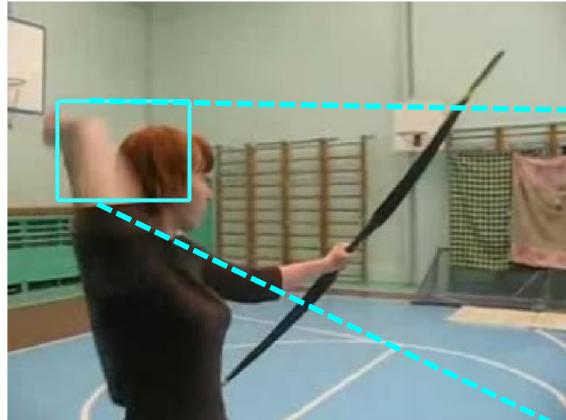
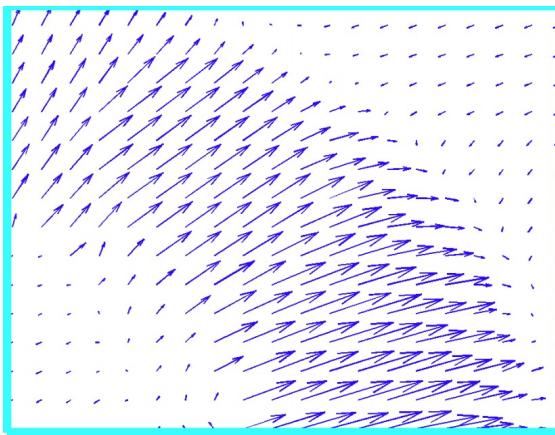


Image at frame $t+1$

Optical flow gives a displacement field F between images I_t and I_{t+1}



Tells where each pixel will move in the next frame:
 $F(x, y) = (dx, dy)$
 $I_{t+1}(x+dx, y+dy) = I_t(x, y)$

Measuring Motion: Optical Flow

Image at frame t

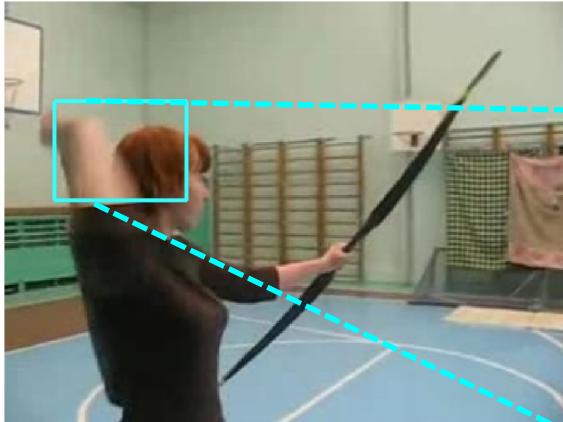
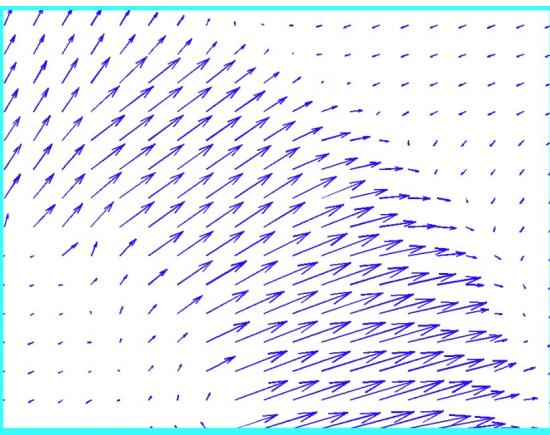


Image at frame $t+1$

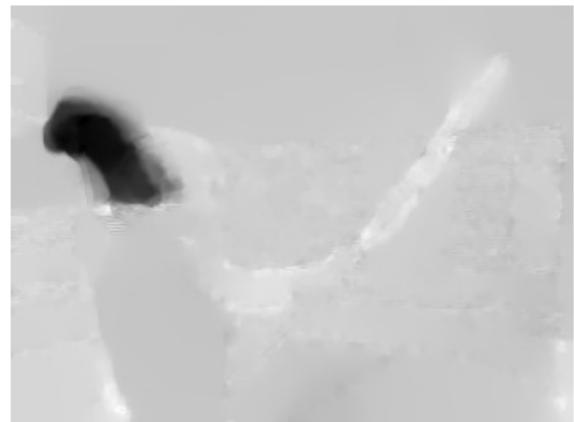
Optical flow gives a displacement field F between images I_t and I_{t+1}



Tells where each pixel will move in the next frame:
 $F(x, y) = (dx, dy)$
 $I_{t+1}(x+dx, y+dy) = I_t(x, y)$

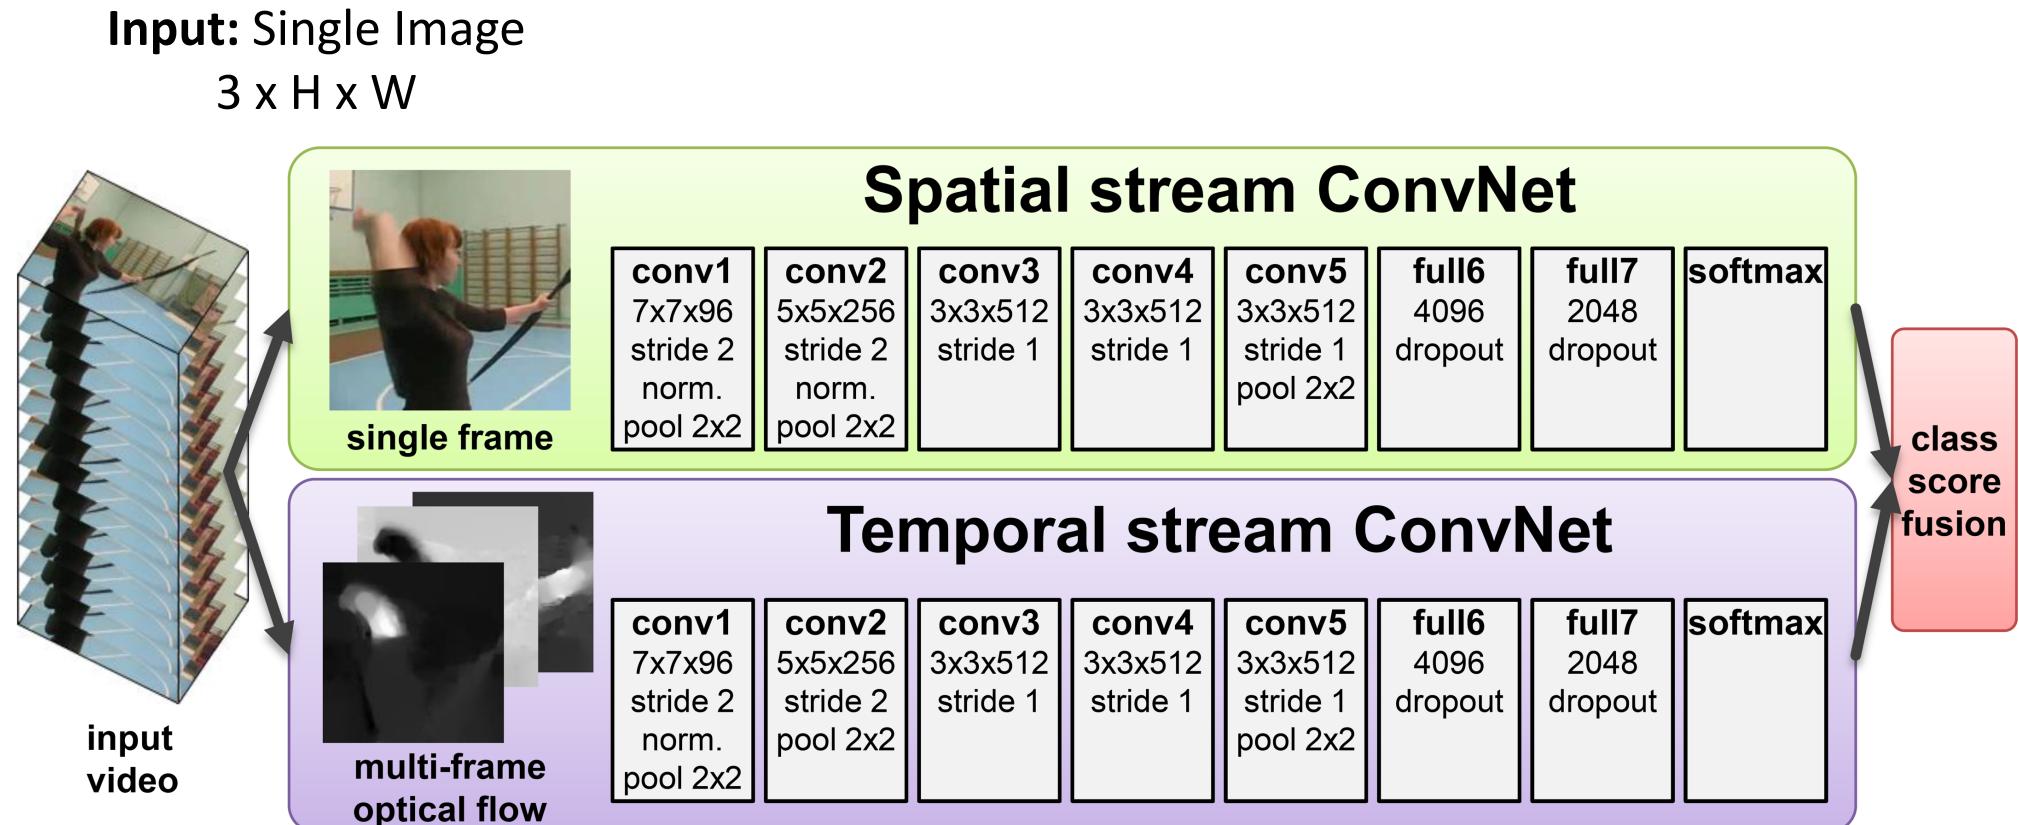
Optical Flow highlights
local motion

Horizontal flow dx



Vertical Flow dy

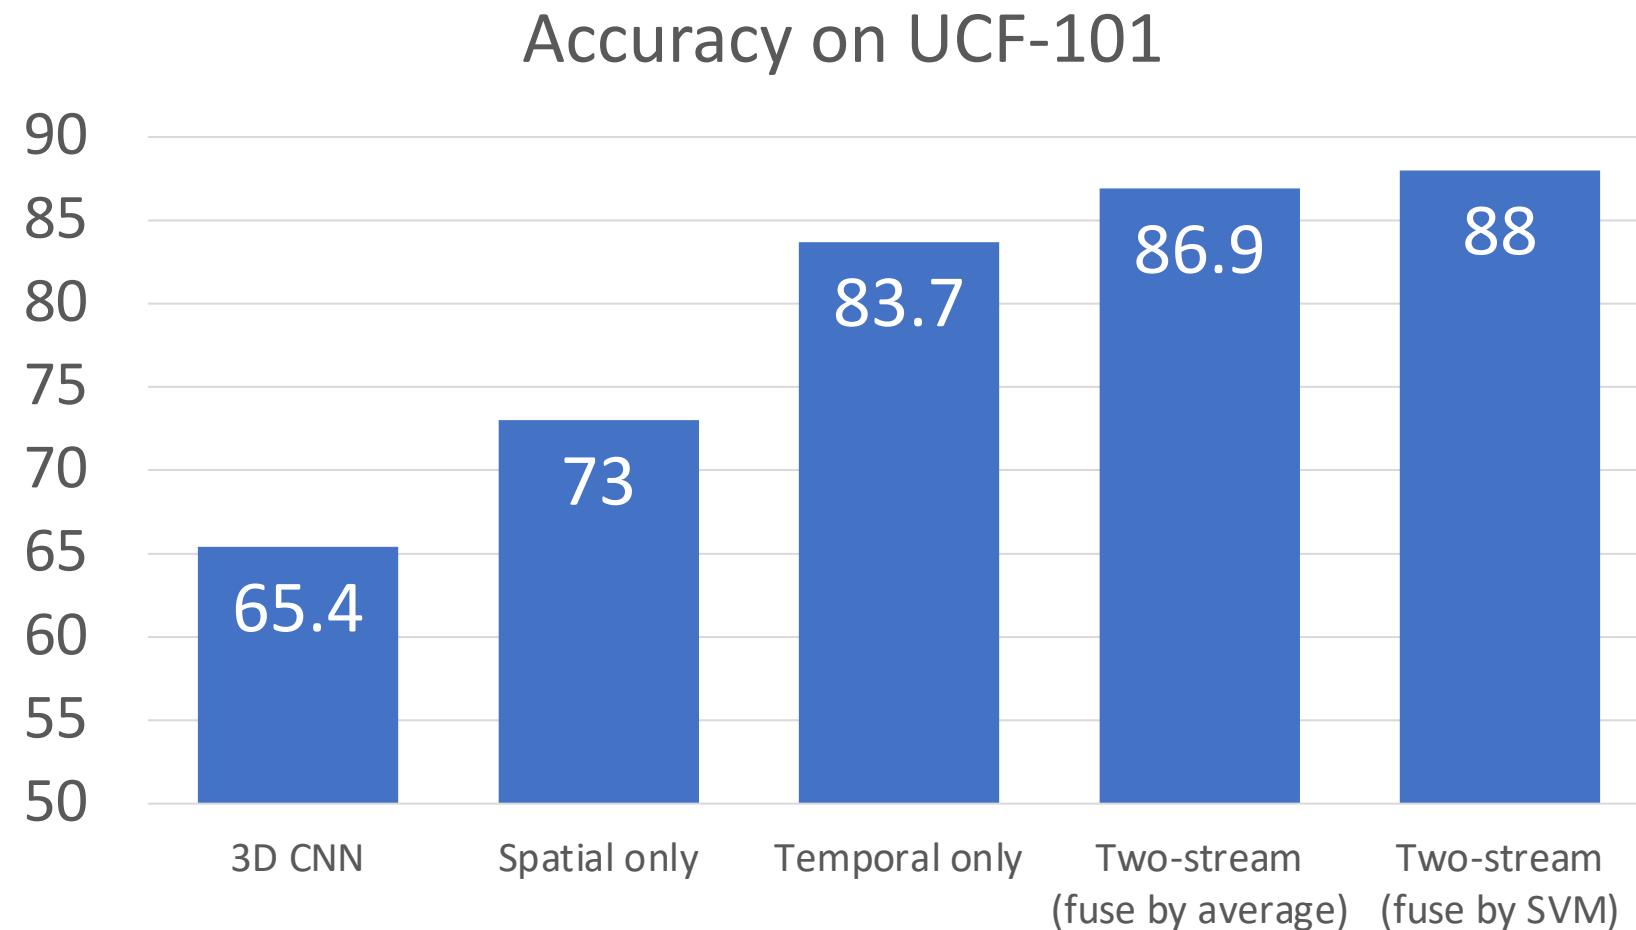
Separating Motion and Appearance: Two-Stream Networks



Input: Stack of optical flow:
 $[2*(T-1)] \times H \times W$

Early fusion: First 2D conv
processes all flow images

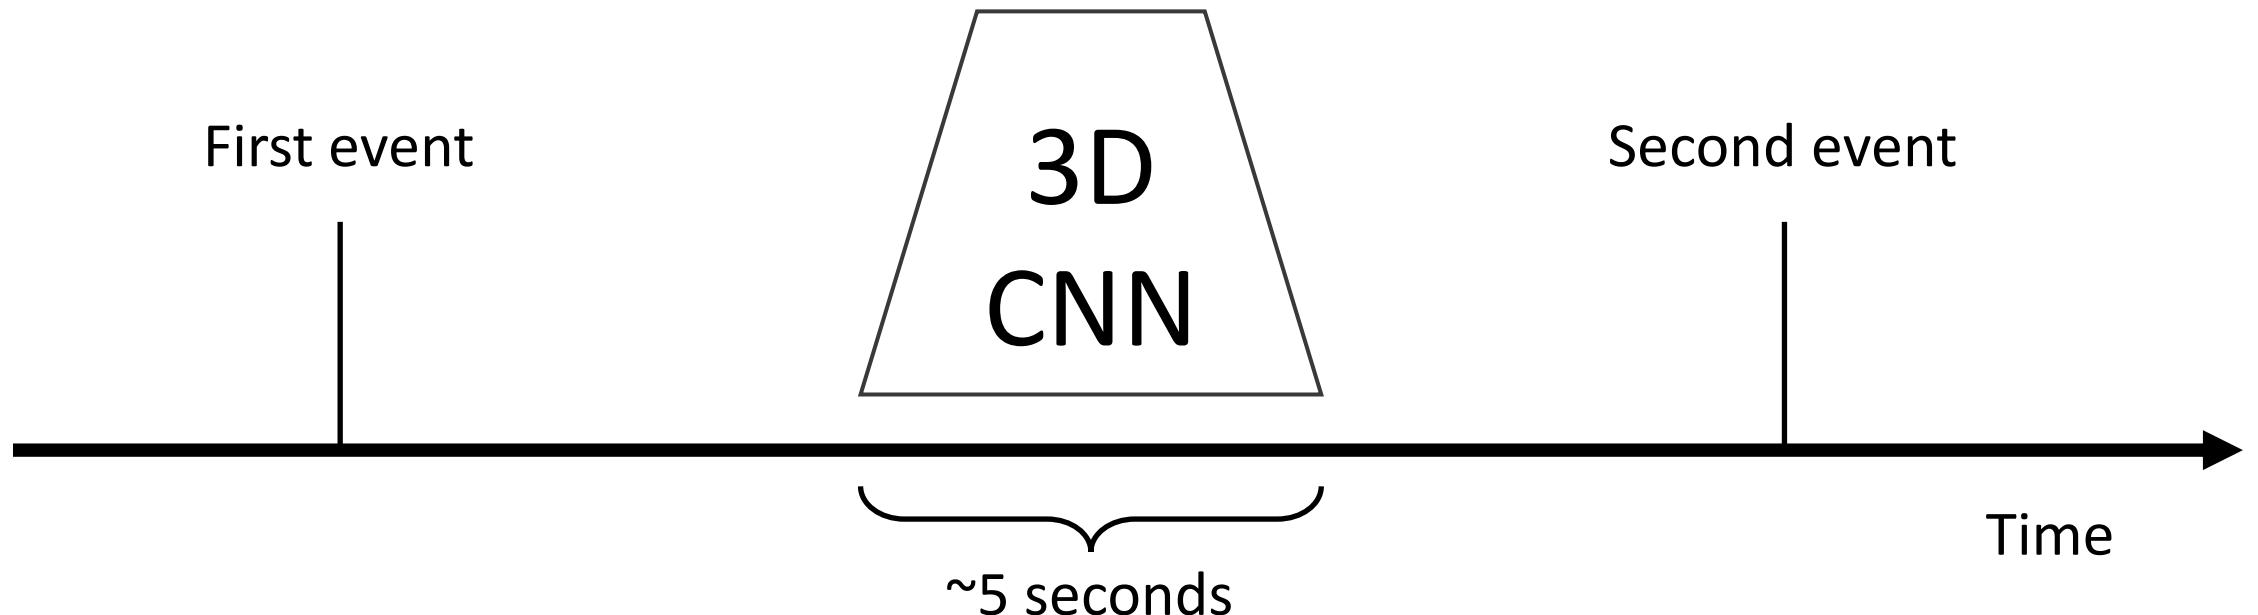
Separating Motion and Appearance: Two-Stream Networks



Simonyan and Zisserman, "Two-stream convolutional networks for action recognition in videos", NeurIPS 2014

Modeling long-term temporal structure

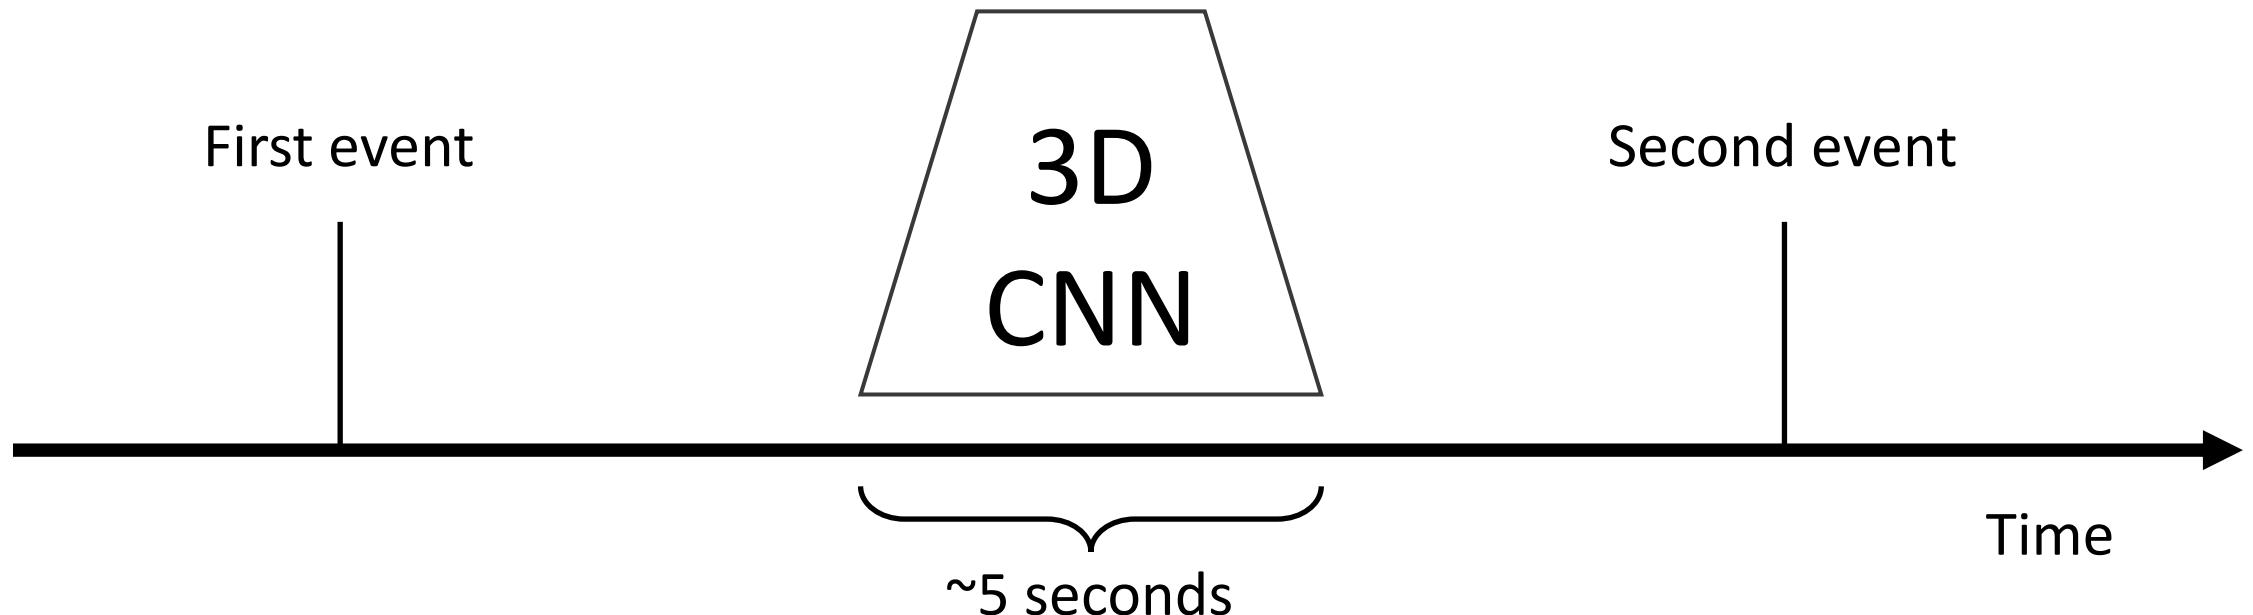
So far all our temporal CNNs only model local motion between frames in very short clips of ~2-5 seconds. What about long-term structure?



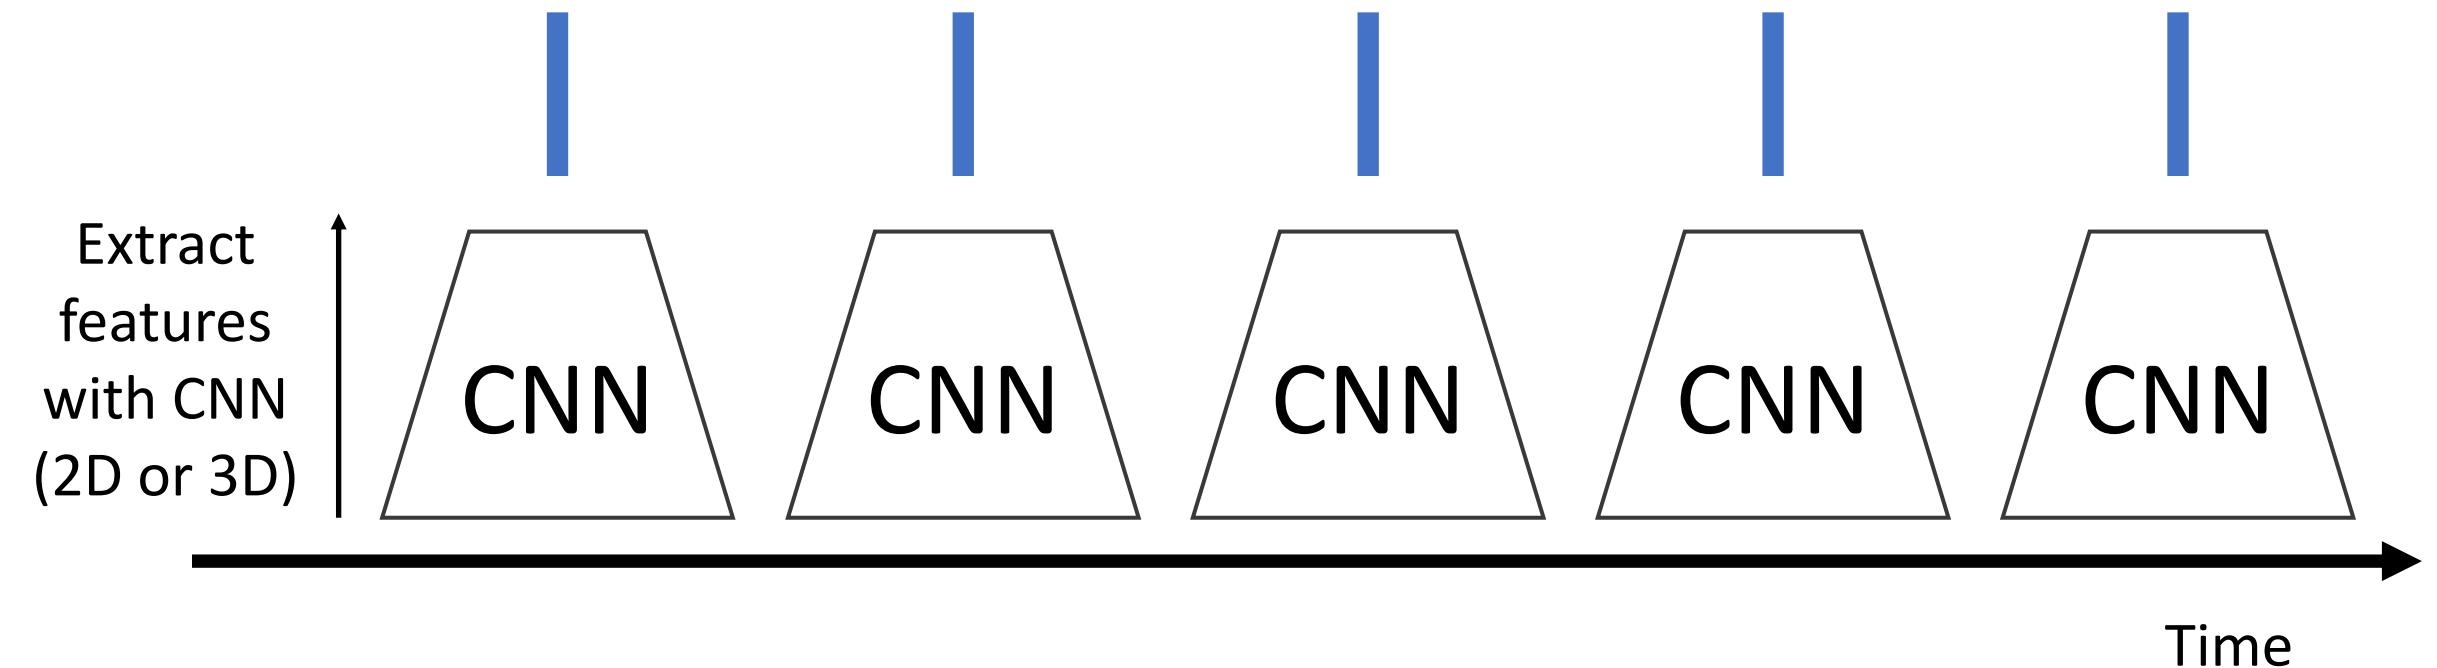
Modeling long-term temporal structure

So far all our temporal CNNs only model local motion between frames in very short clips of ~2-5 seconds. What about long-term structure?

We know how to handle sequences!
How about recurrent networks?

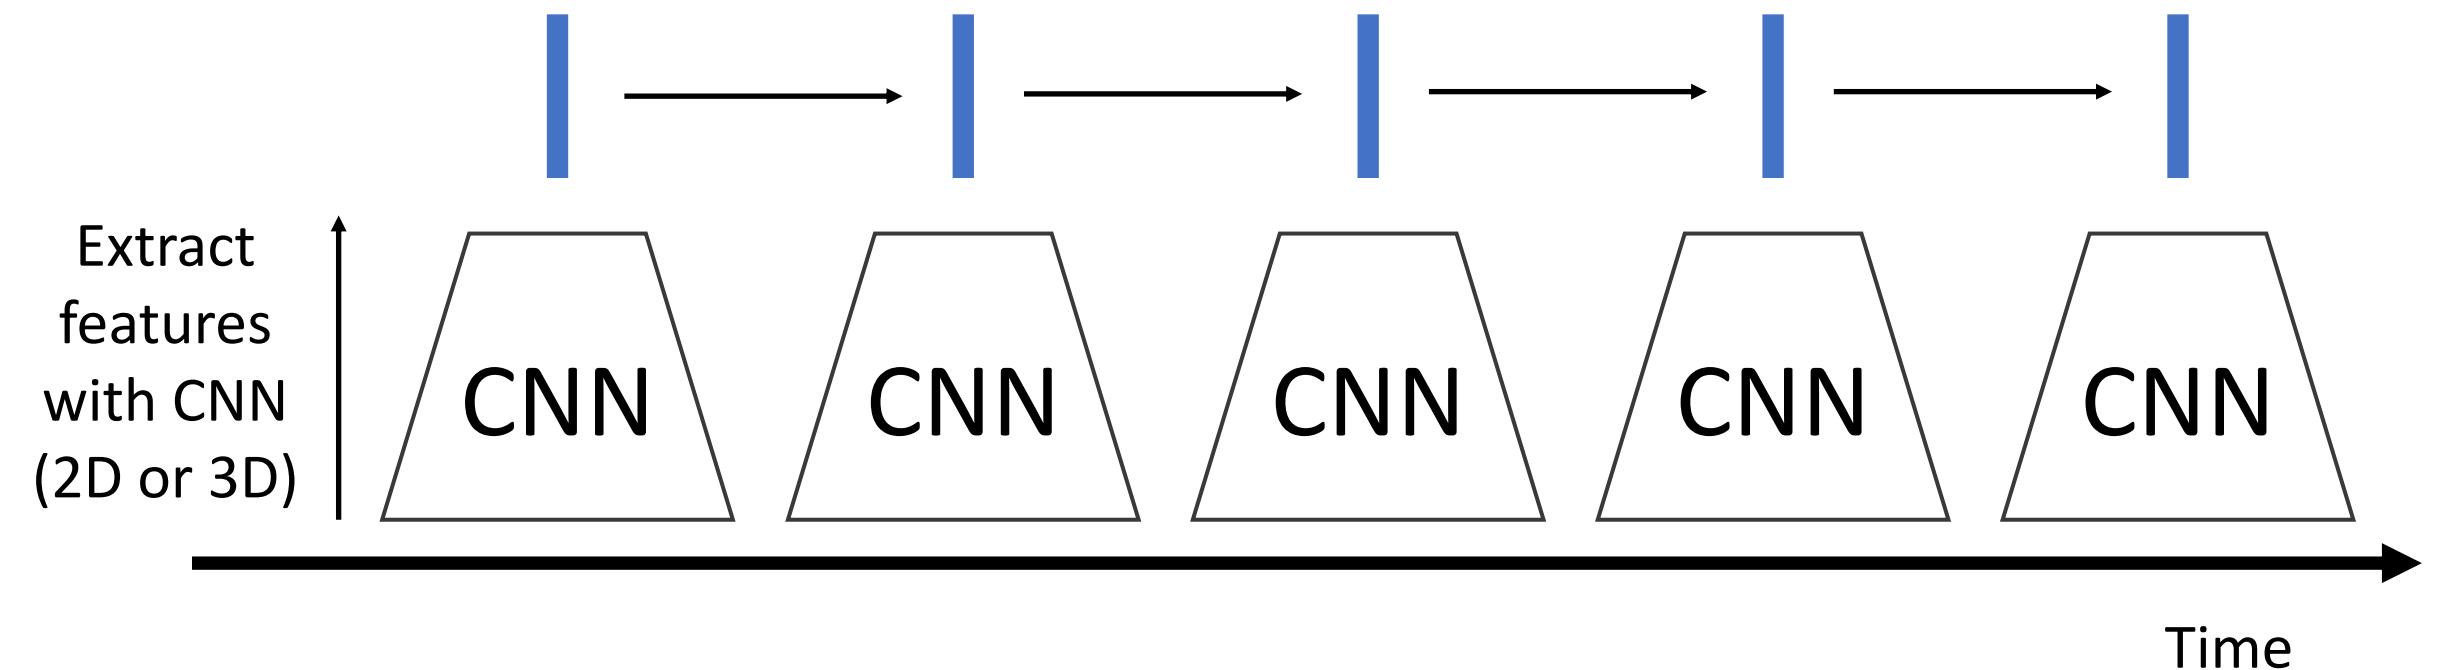


Modeling long-term temporal structure



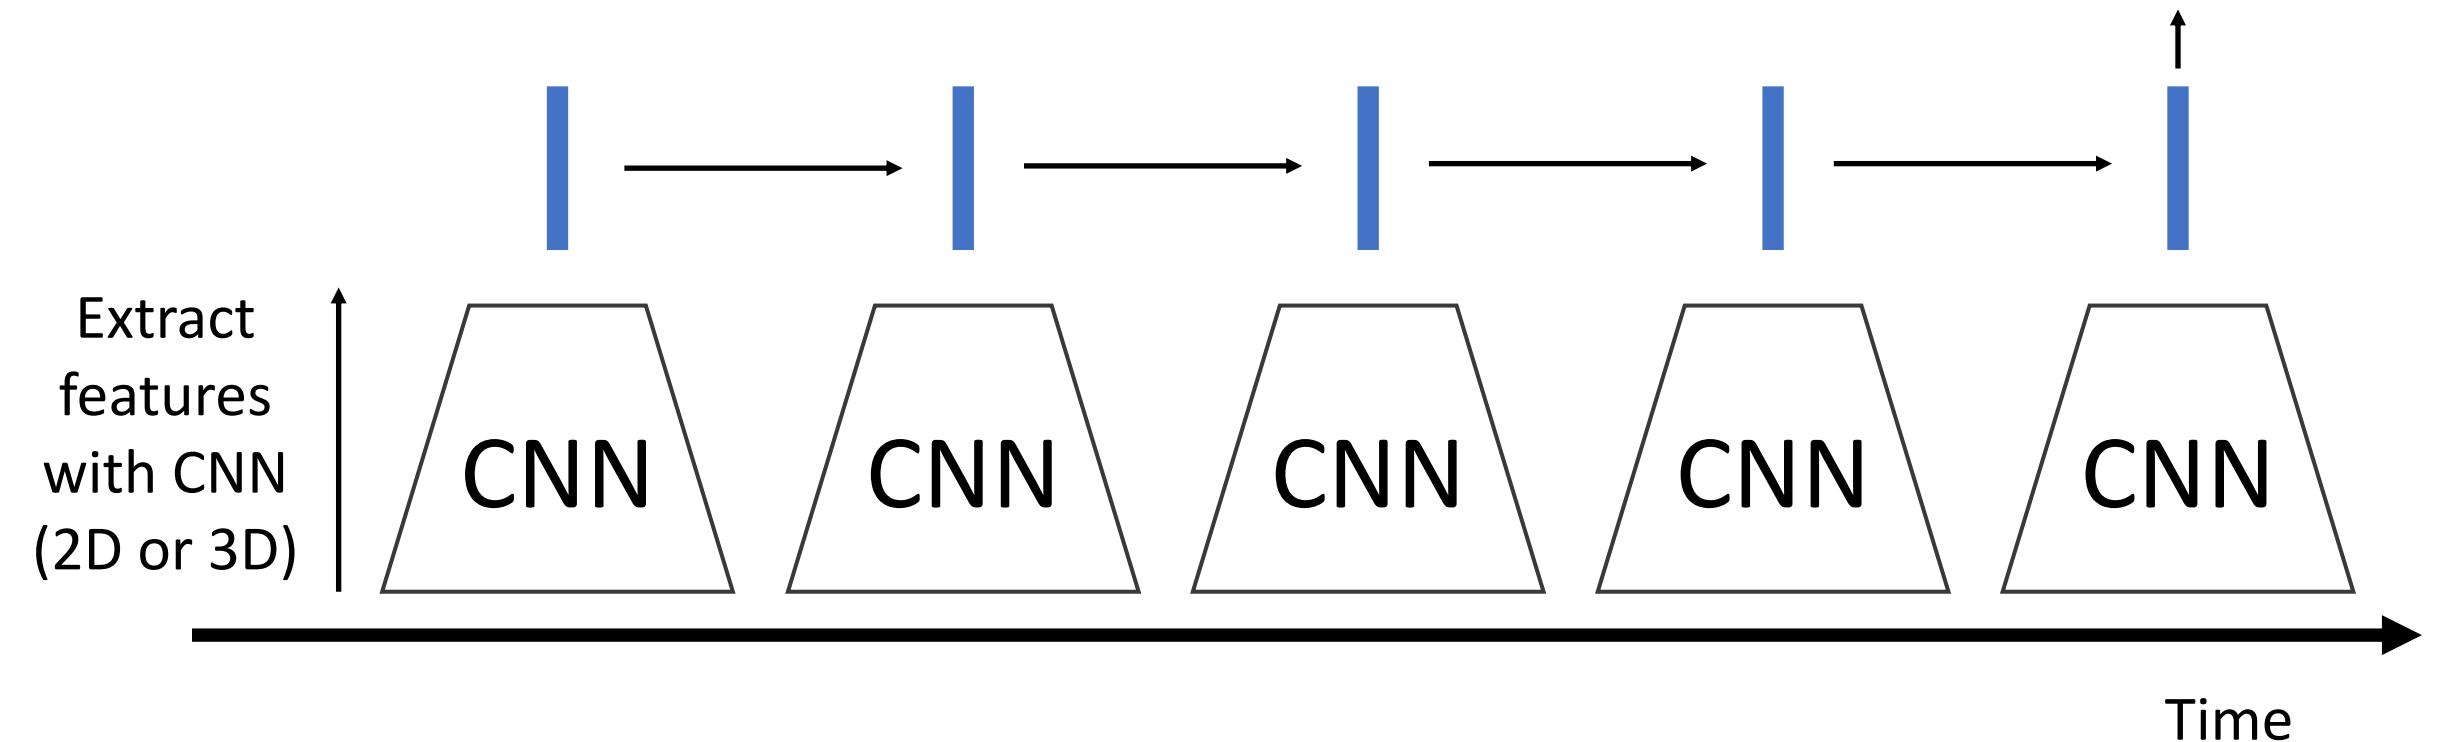
Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)



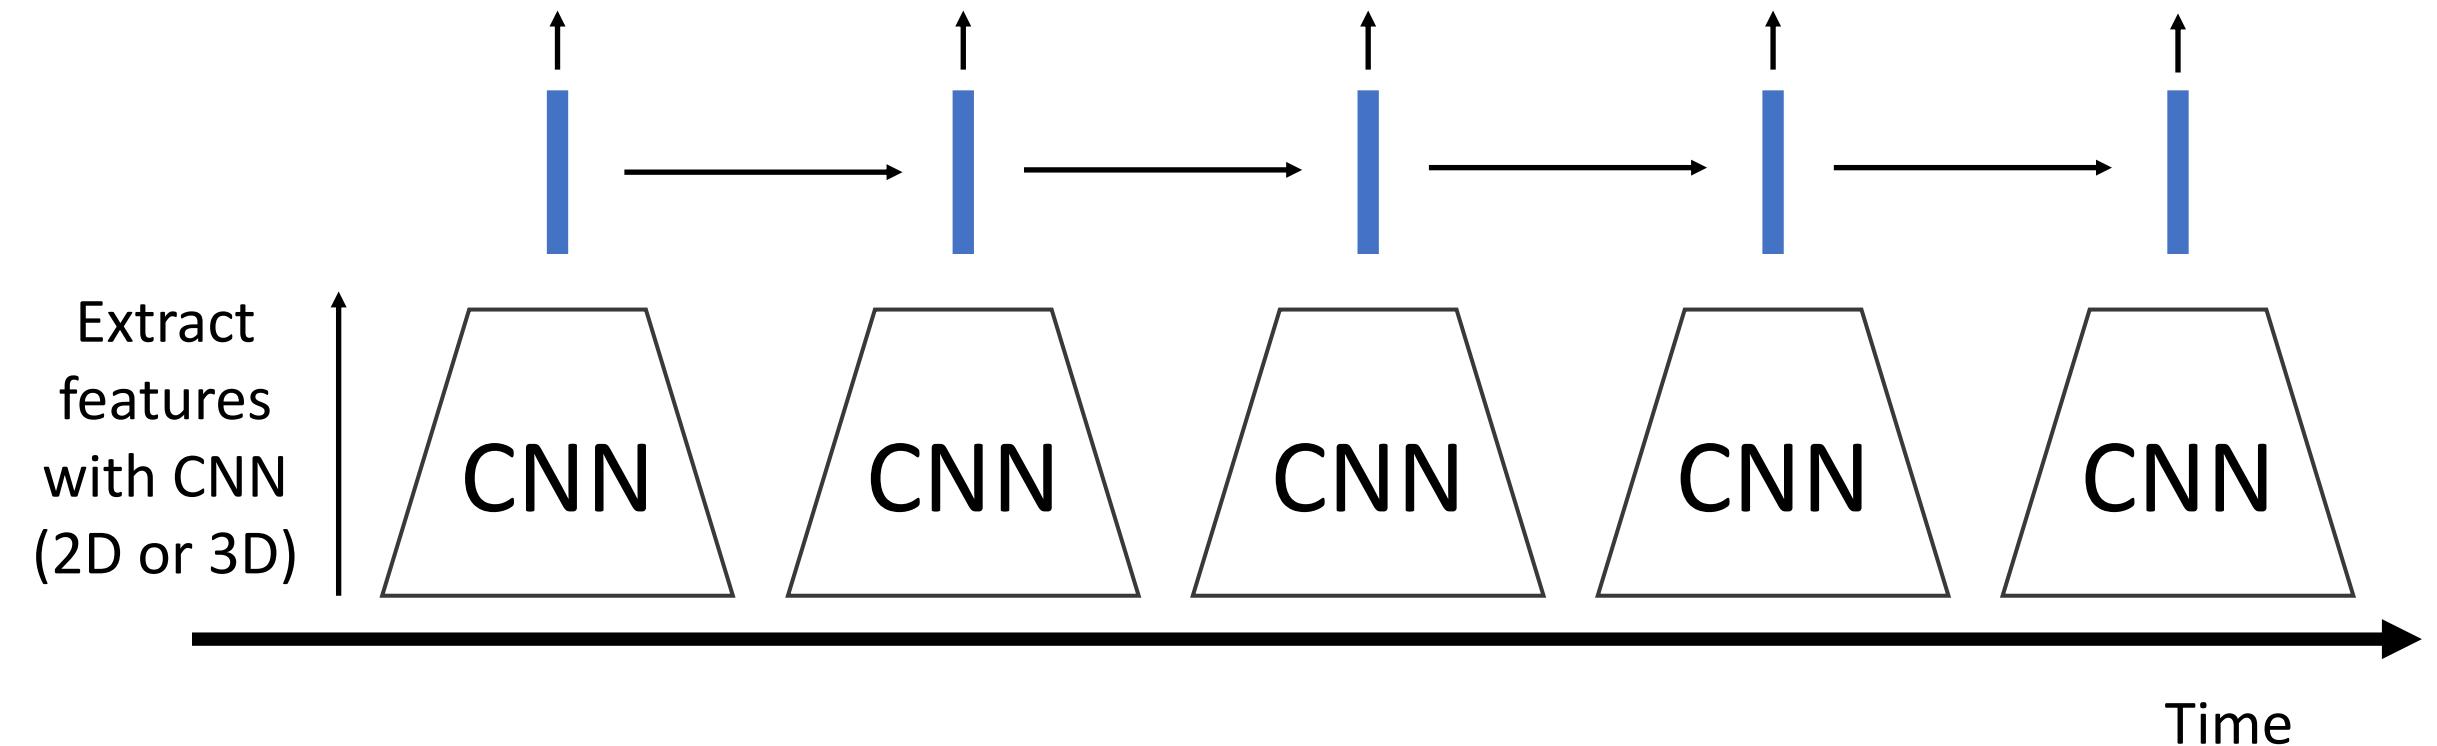
Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)
Many to one: One output at end of video



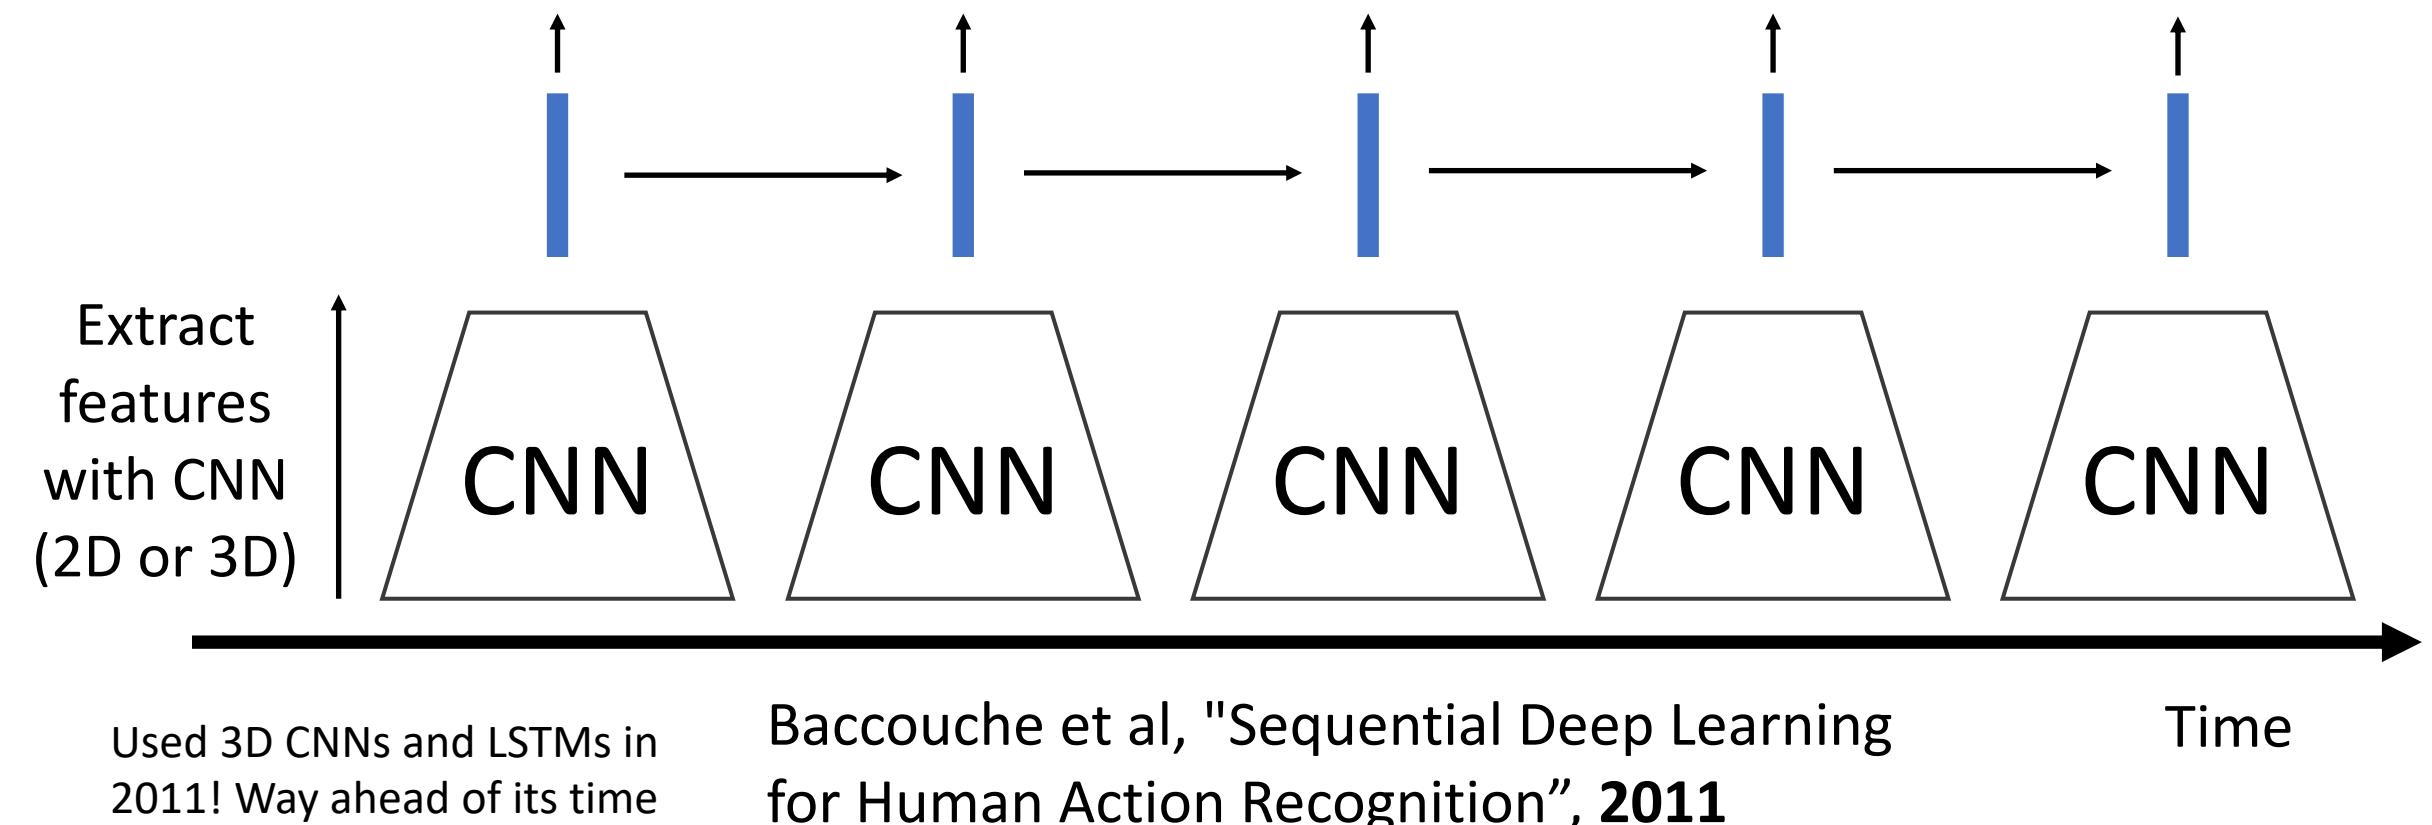
Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)
Many to many: one output per video frame



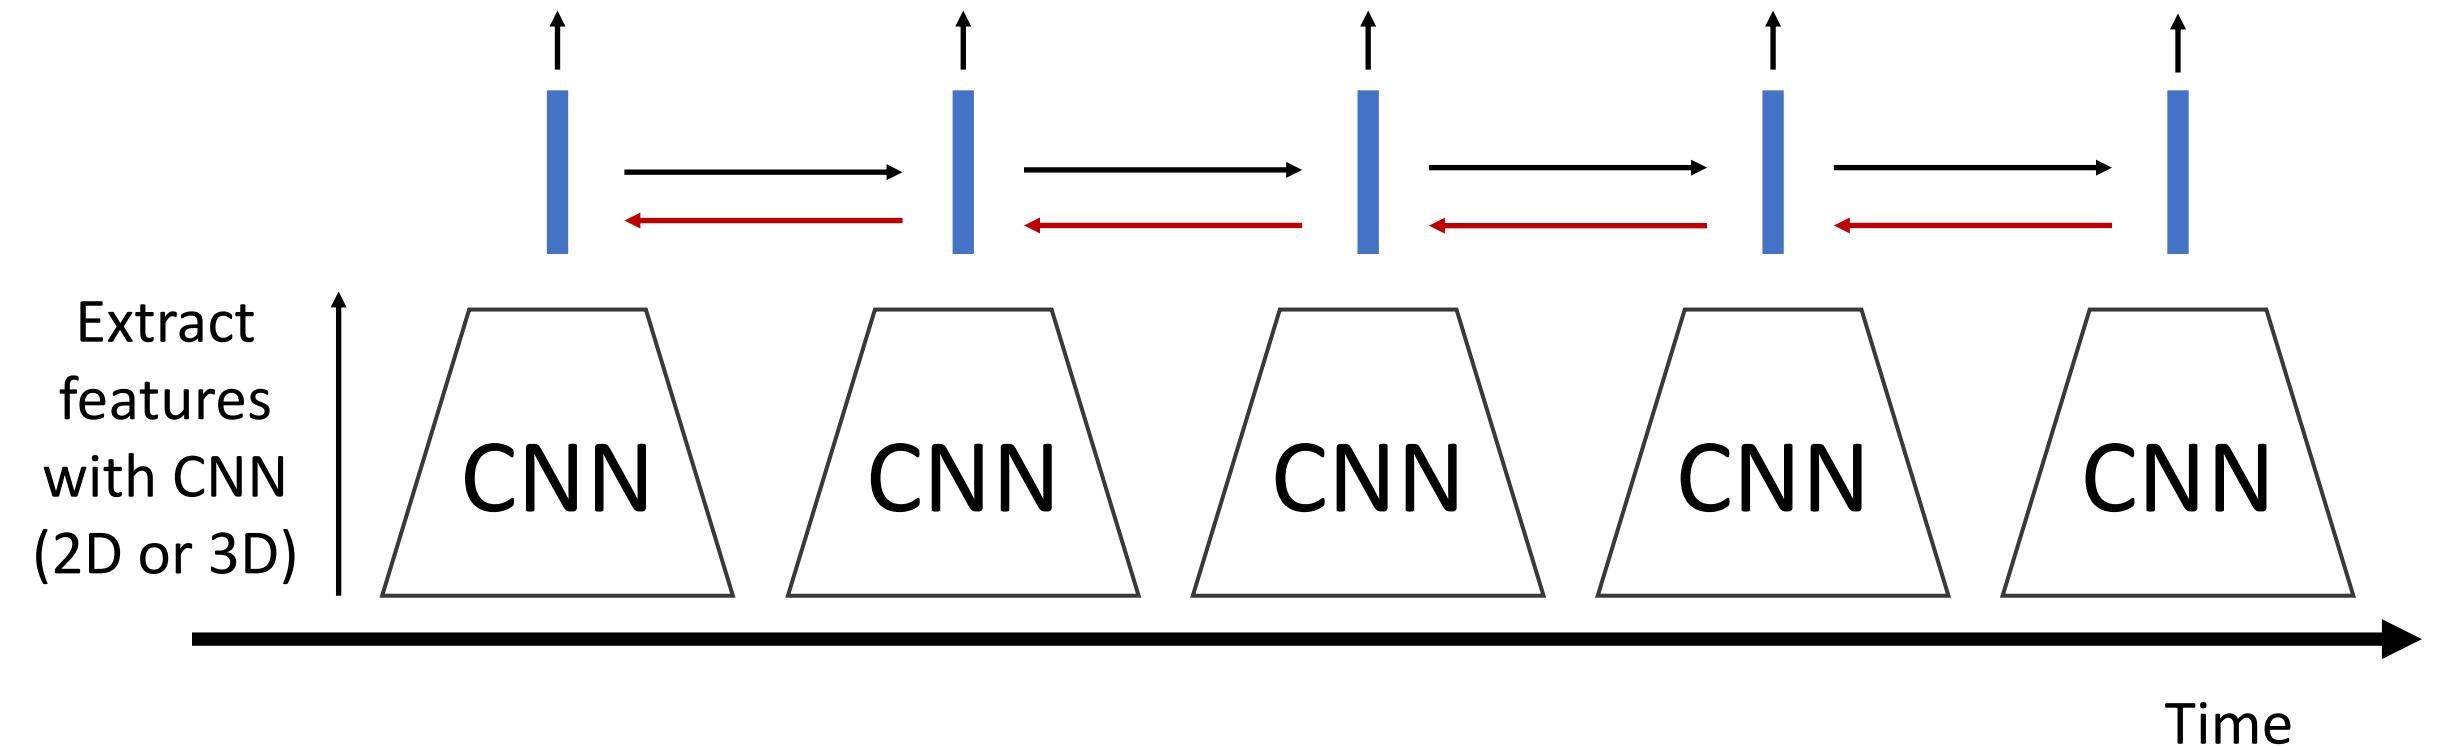
Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)
Many to many: one output per video frame



Modeling long-term temporal structure

Sometimes don't backprop to CNN to save memory;
pretrain and use it as a feature extractor



Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011

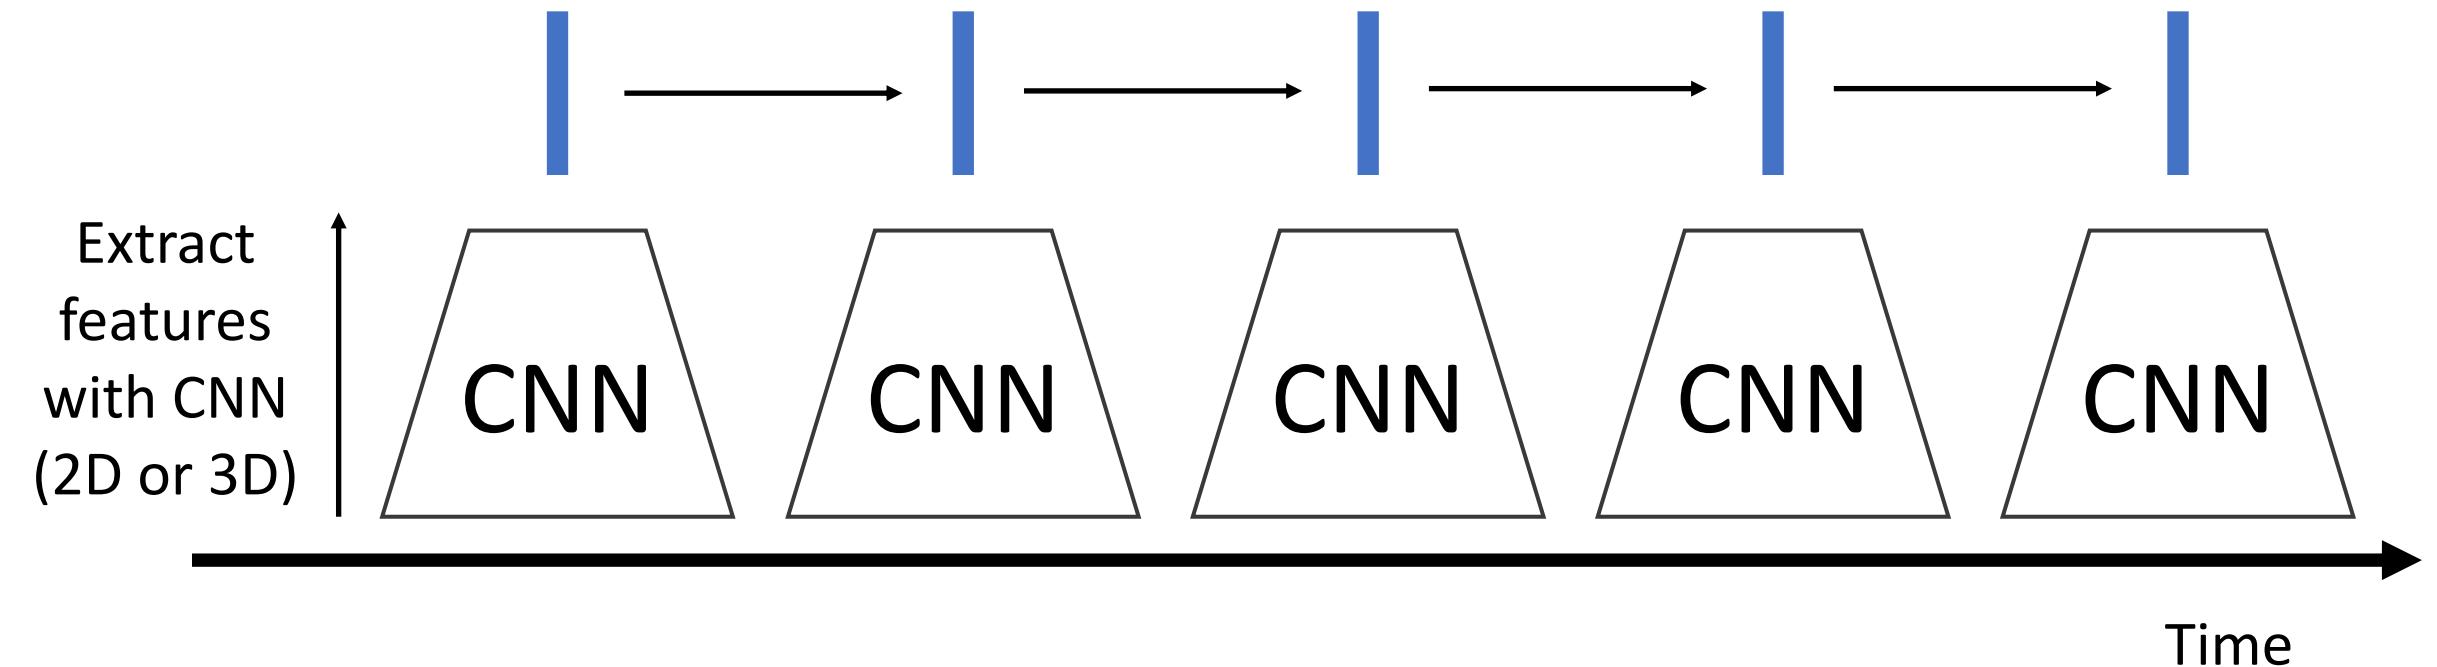
Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

Modeling long-term temporal structure

Inside CNN: Each value a function of a fixed temporal window (local temporal structure)

Inside RNN: Each vector is a function of all previous vectors (global temporal structure)

Can we merge both approaches?

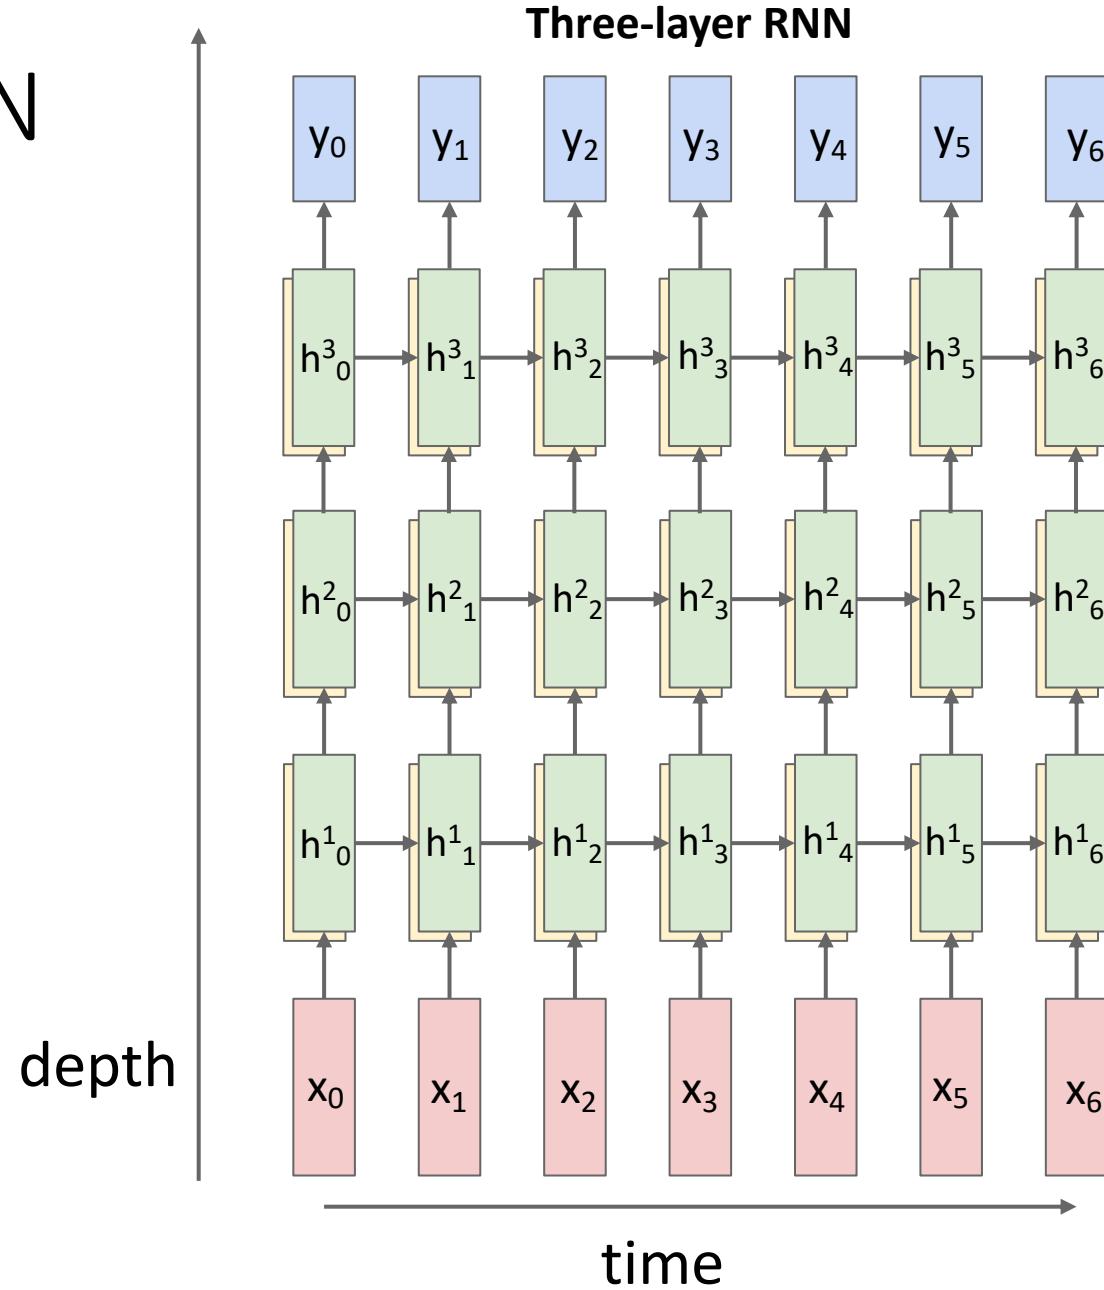


Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011

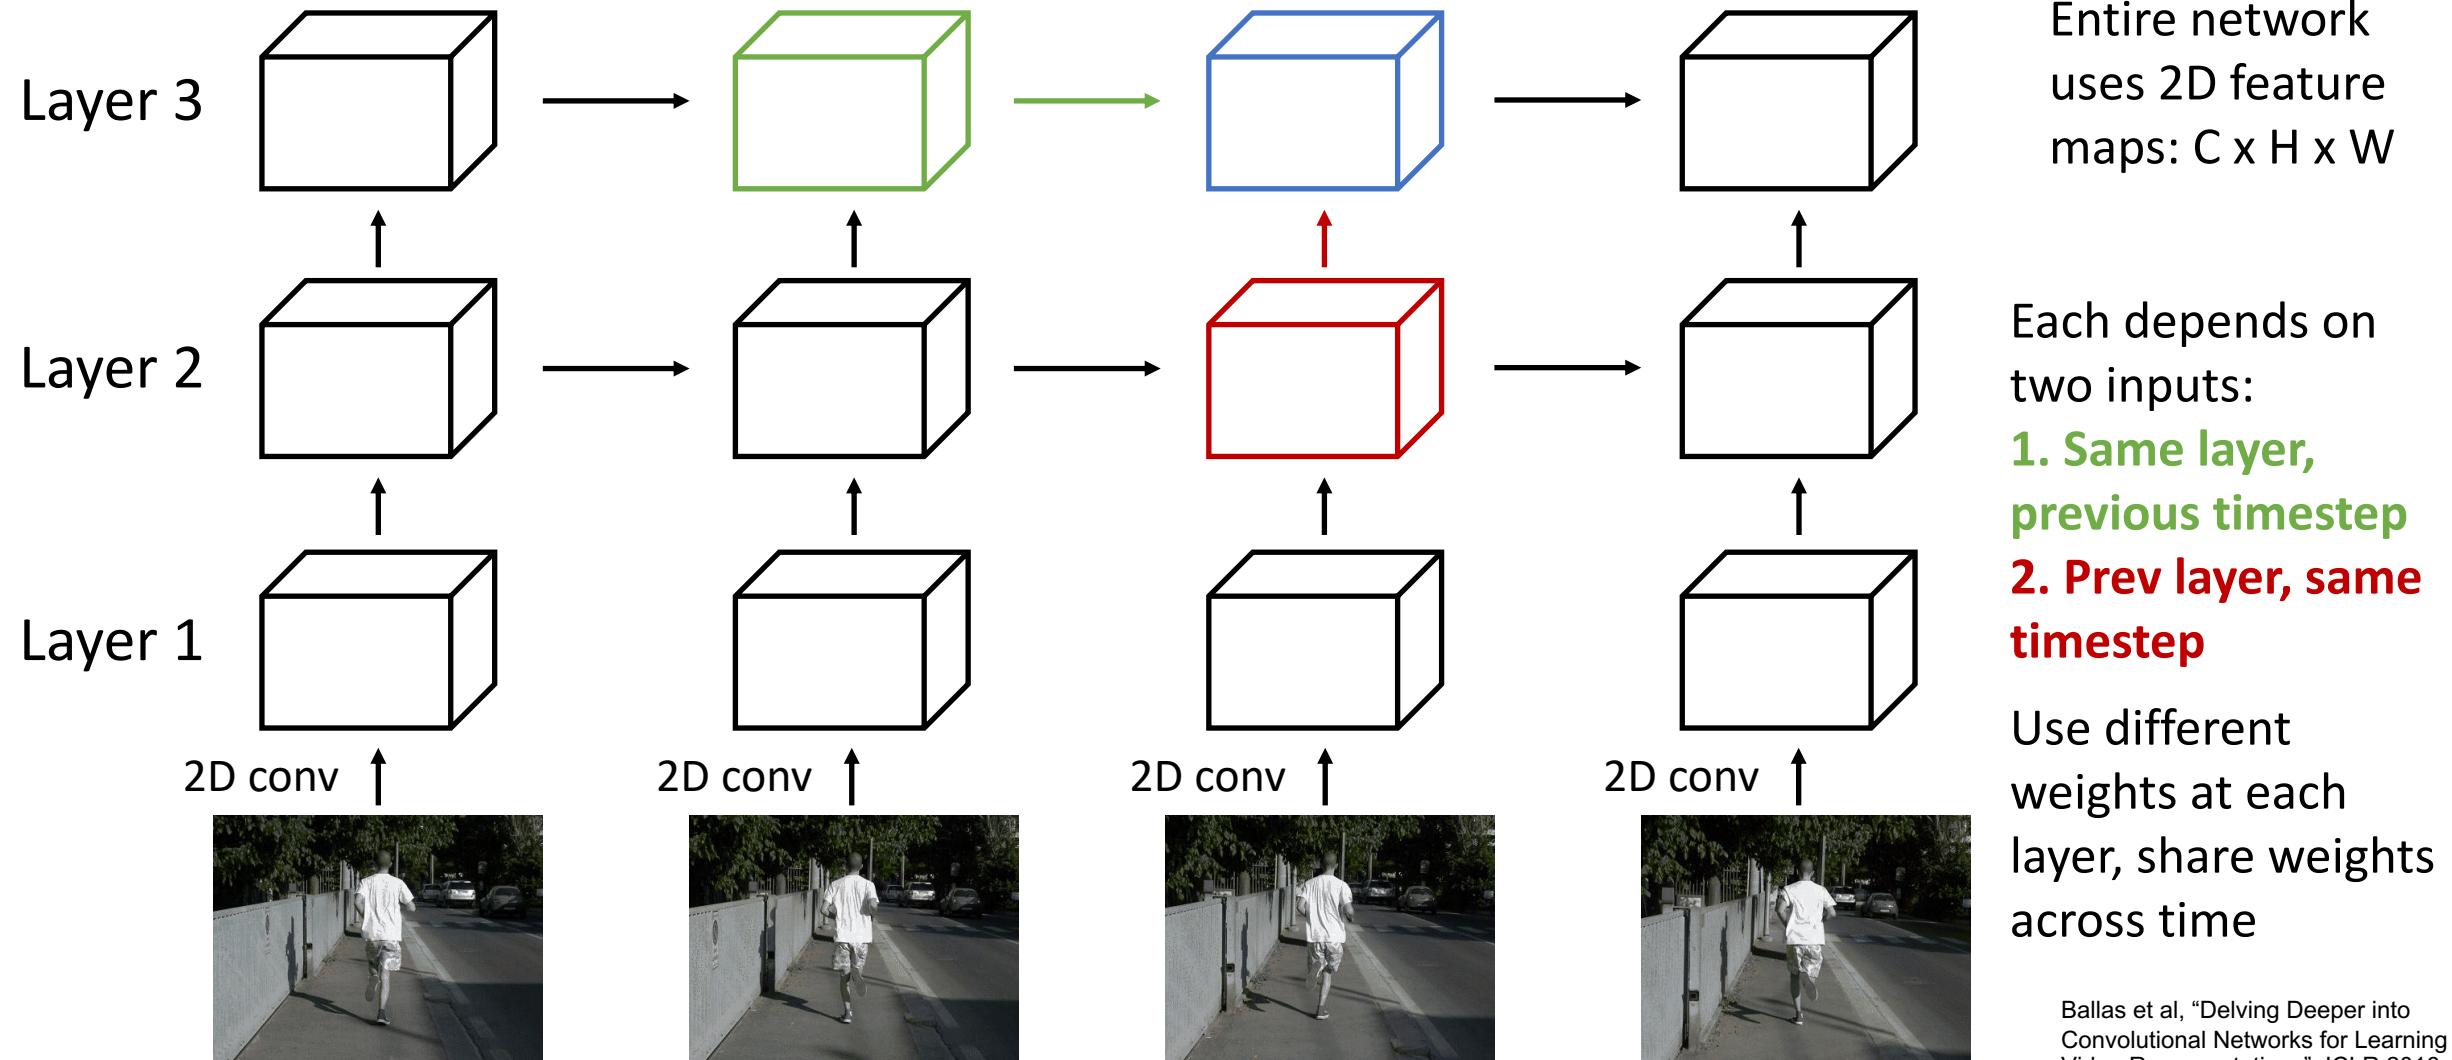
Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

Recall: Multi-layer RNN

We can use a similar structure to process videos!



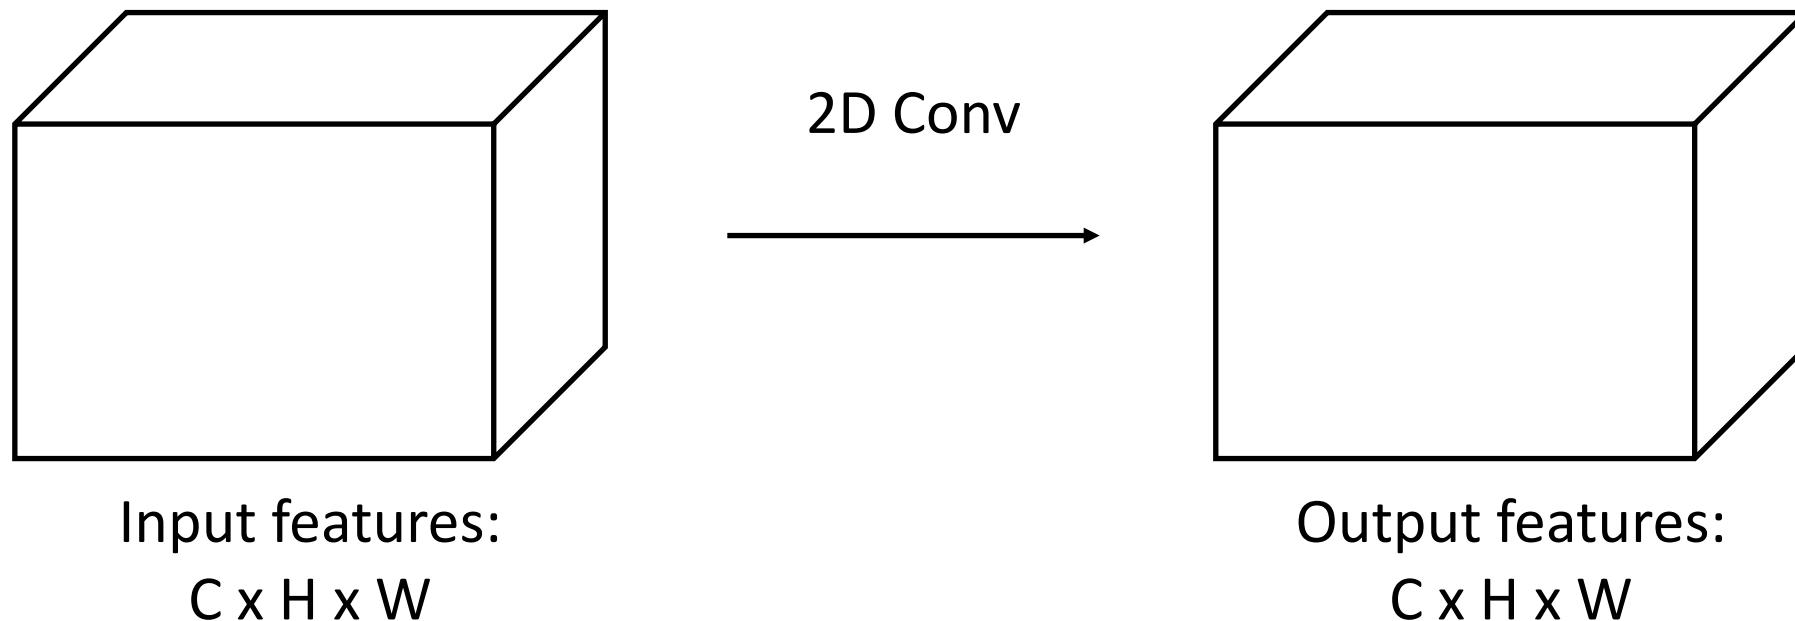
Recurrent Convolutional Network



Ballas et al, "Delving Deeper into Convolutional Networks for Learning Video Representations", ICLR 2016

Recurrent Convolutional Network

Normal 2D CNN:



Recurrent Convolutional Network

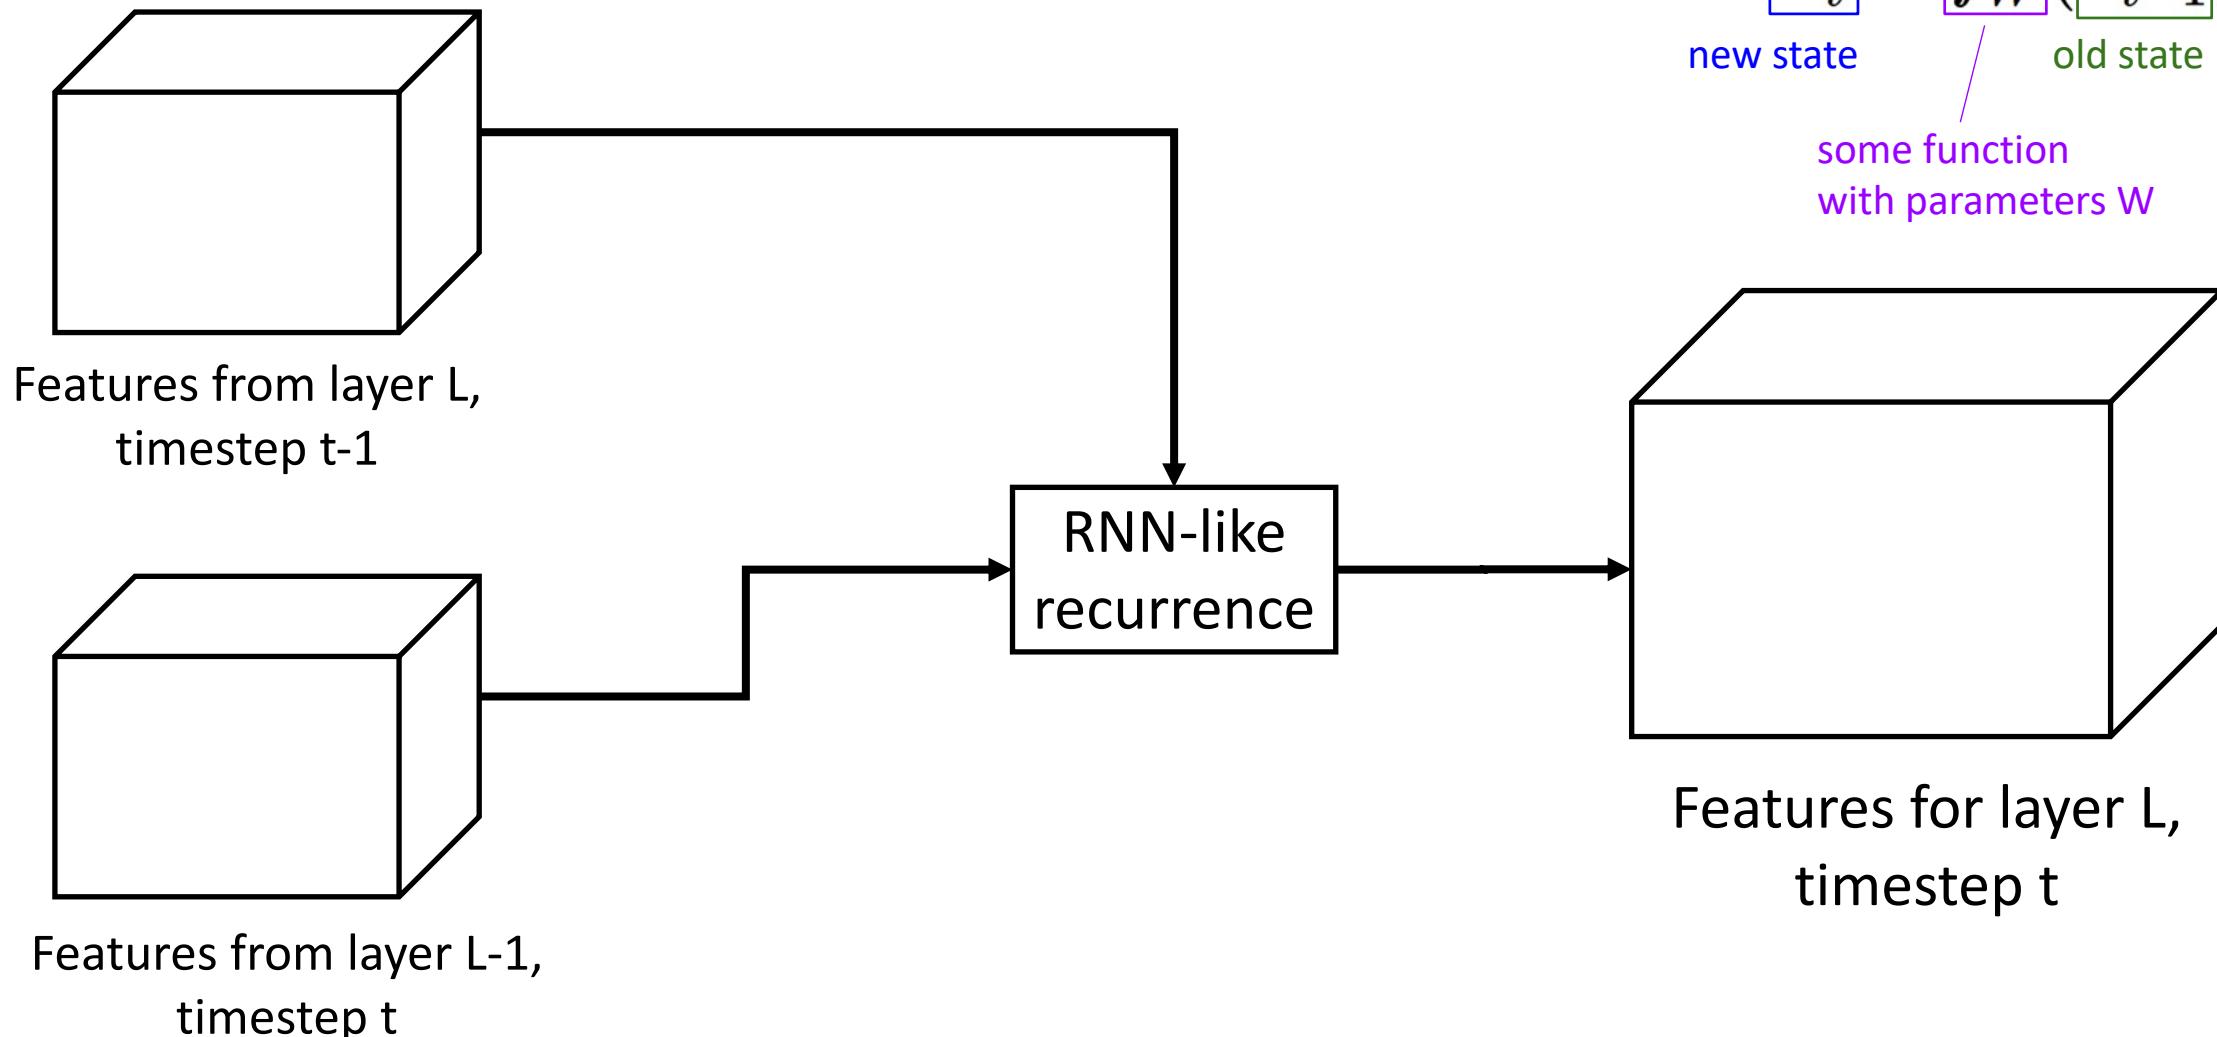
Recall: Recurrent Network

$$h_t = f_W(h_{t-1}, x_t)$$

new state

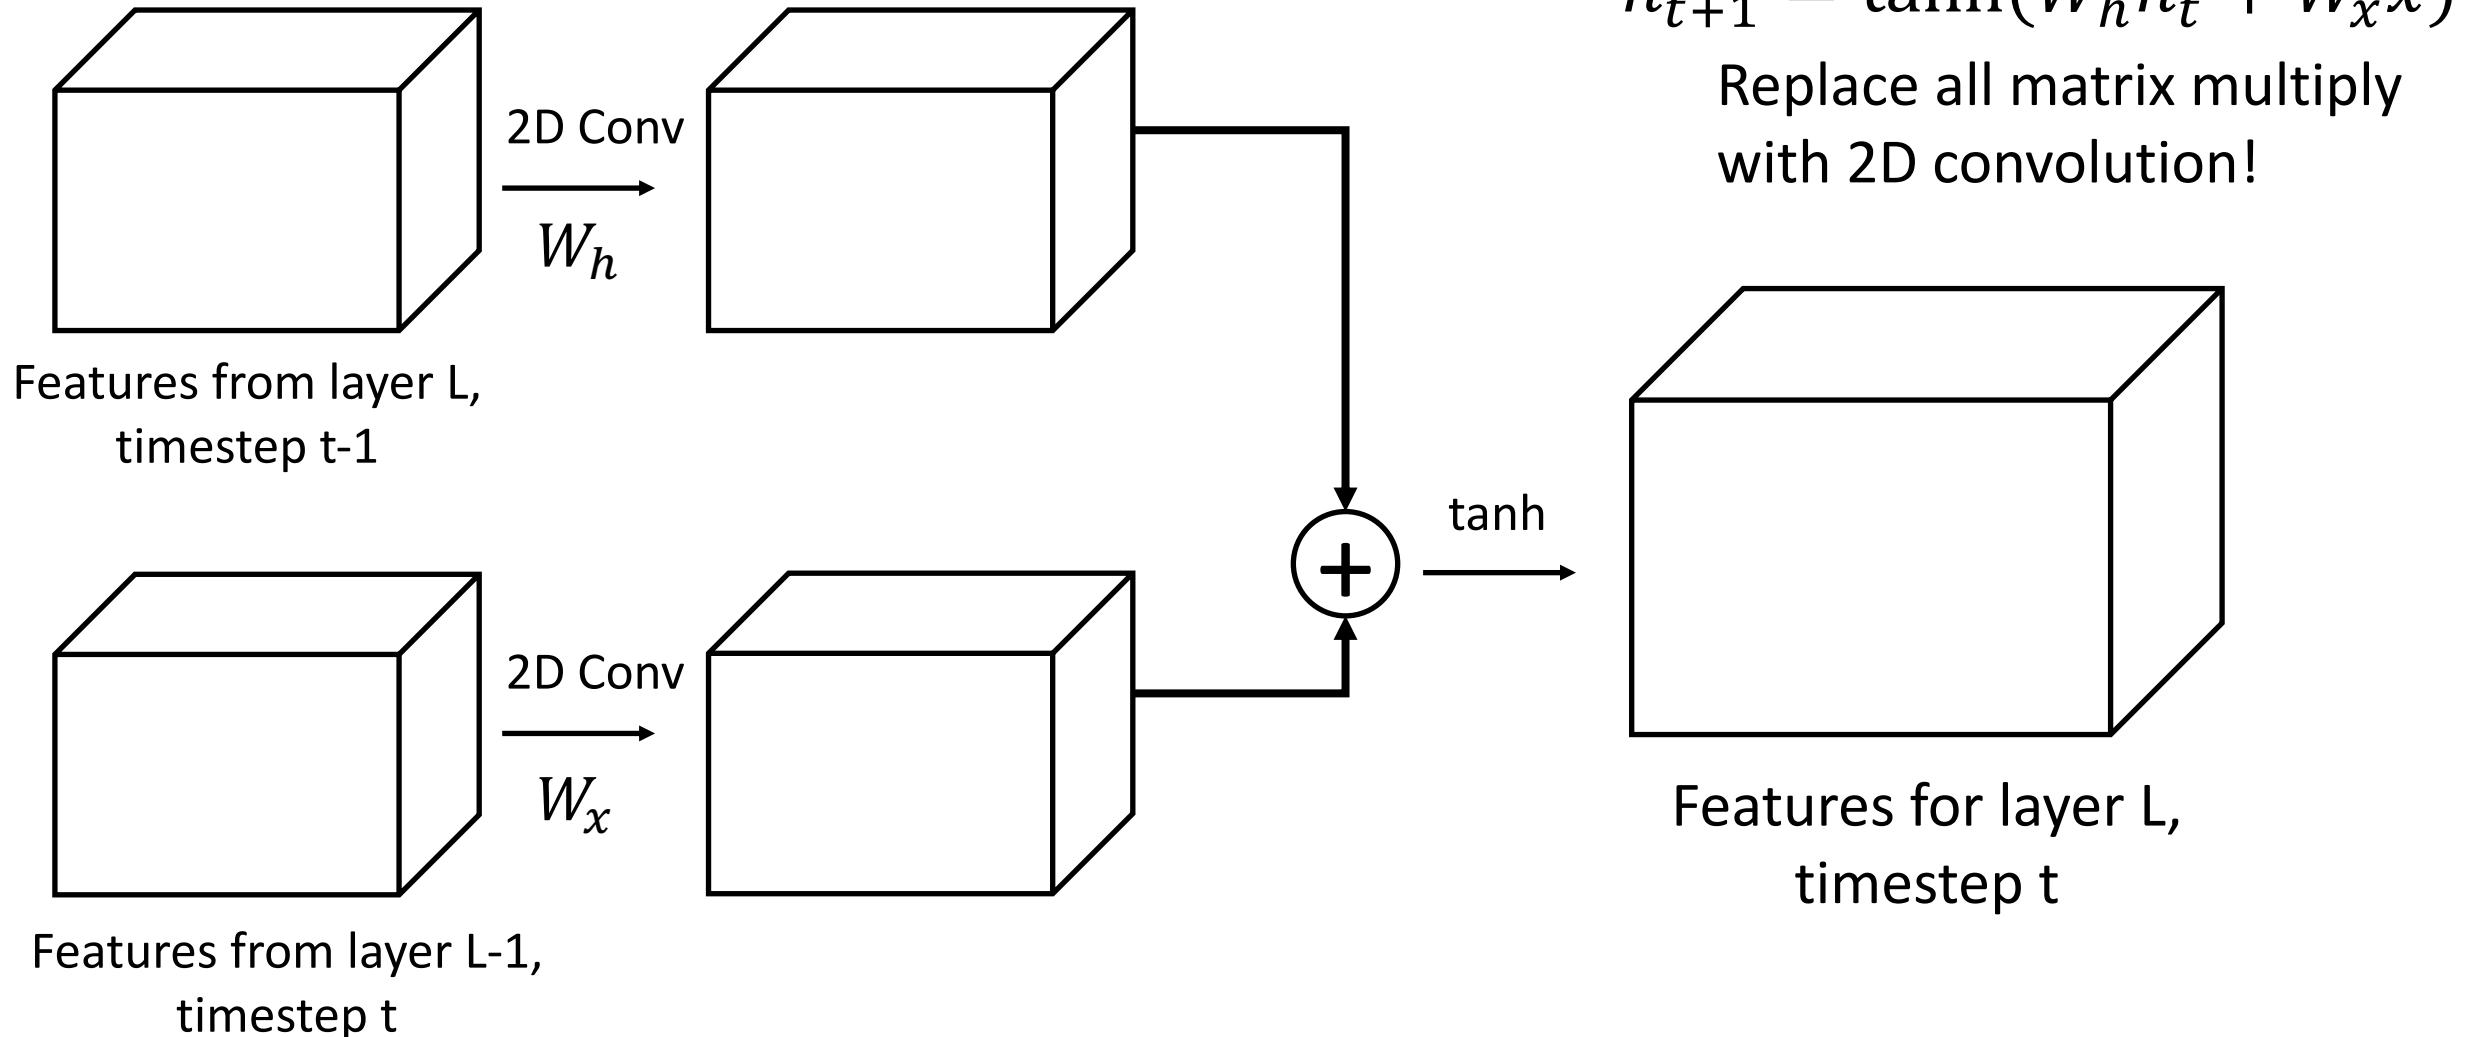
old state

some function
with parameters W



Ballas et al, "Delving Deeper into Convolutional Networks for Learning Video Representations", ICLR 2016

Recurrent Convolutional Network

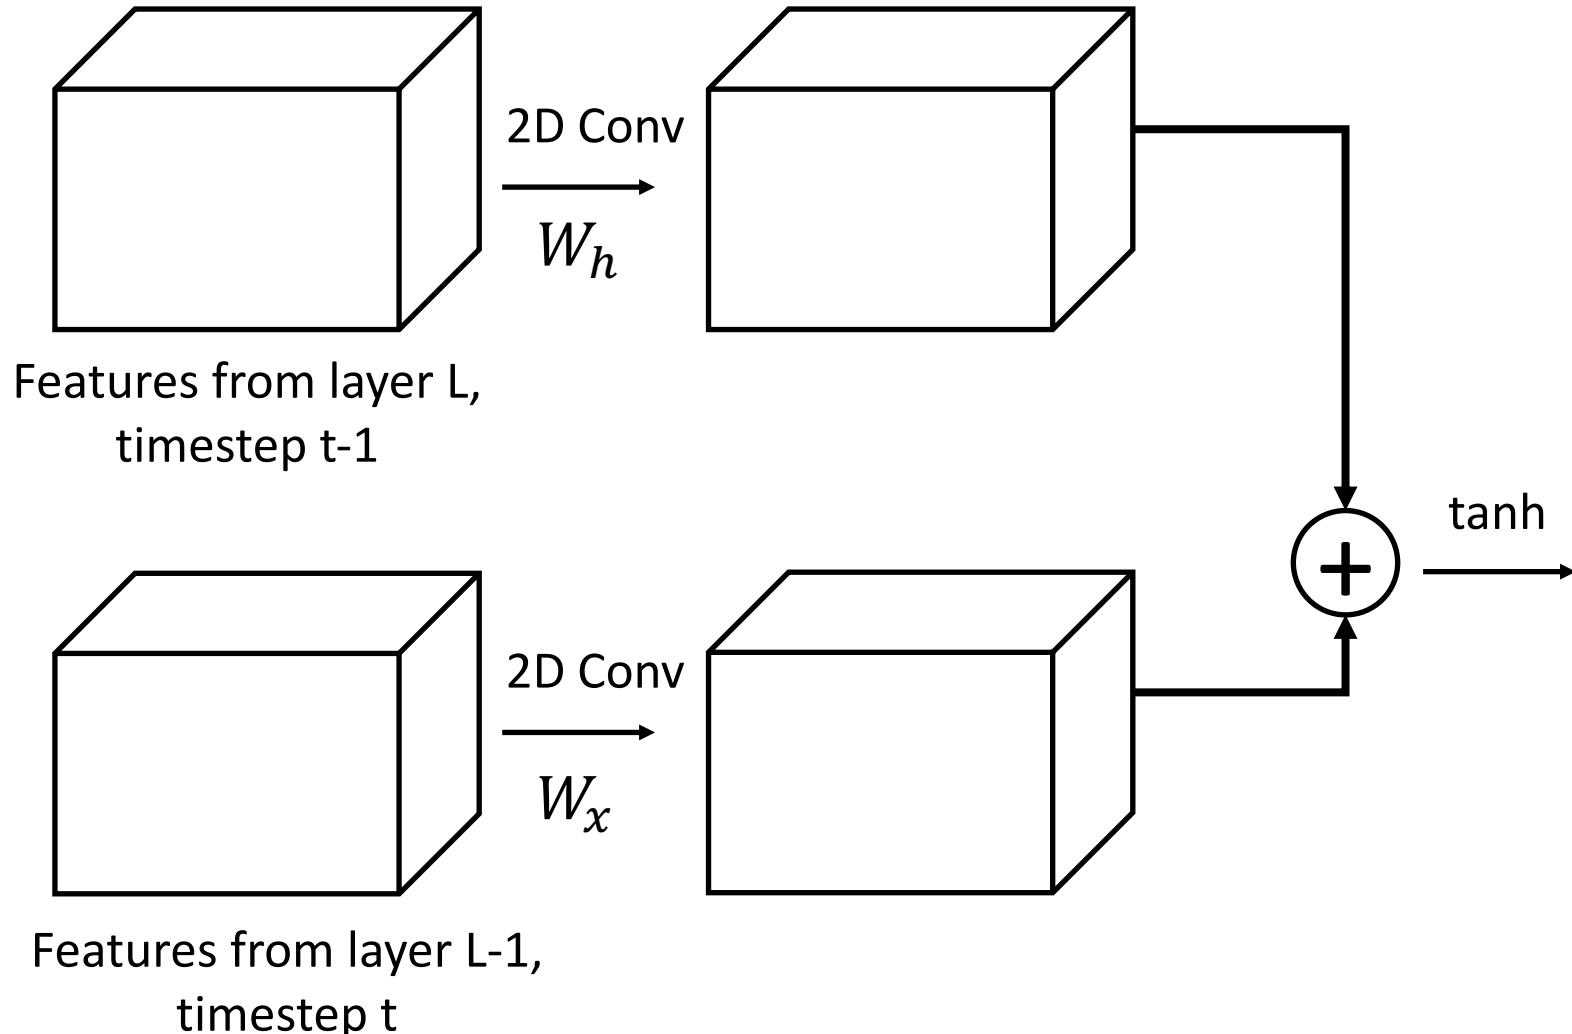


Recall: Vanilla RNN

$$h_{t+1} = \tanh(W_h h_t + W_x x)$$

Replace all matrix multiply
with 2D convolution!

Recurrent Convolutional Network



Recall: GRU

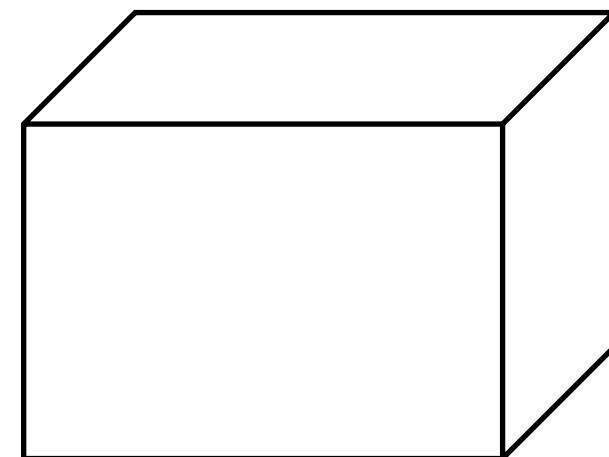
$$r_t = \sigma(W_{xr}x_t + W_{hr}h_{t-1} + b_r)$$

$$z_t = \sigma(W_{xz}x_t + W_{hz}h_{t-1} + b_z)$$

$$\tilde{h}_t = \tanh(W_{xh}x_t + W_{hh}(r_t \odot h_{t-1}) + b_h)$$

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \tilde{h}_t$$

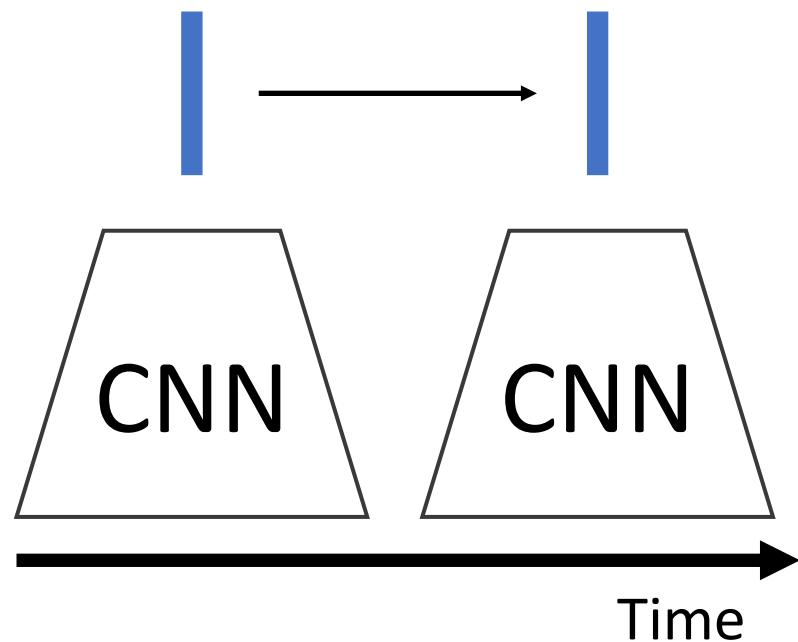
Can do similar transform for
other RNN variants (GRU, LSTM)



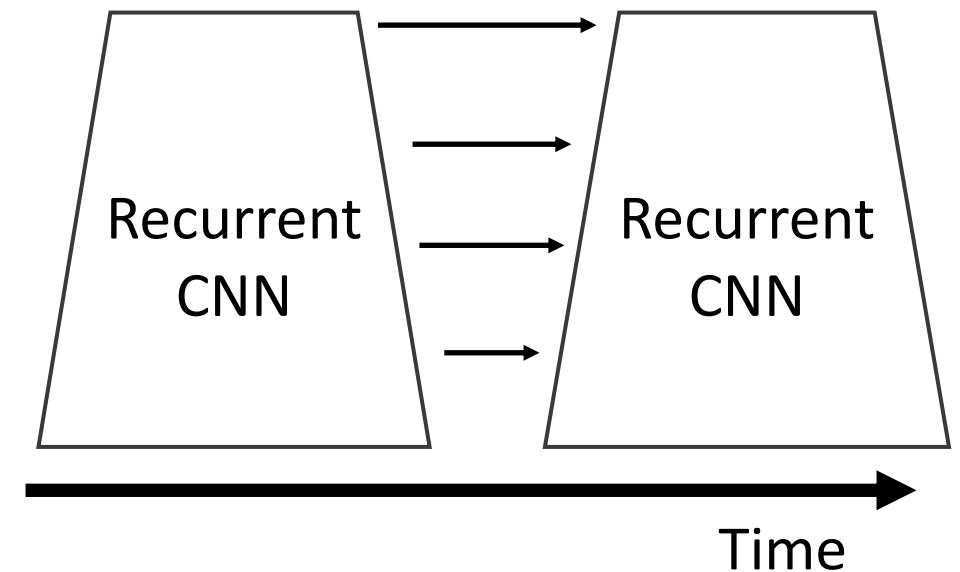
Modeling long-term temporal structure

RNN: Infinite
temporal extent
(fully-connected)

CNN: finite
temporal extent
(convolutional)



Recurrent CNN: Infinite
temporal extent
(convolutional)



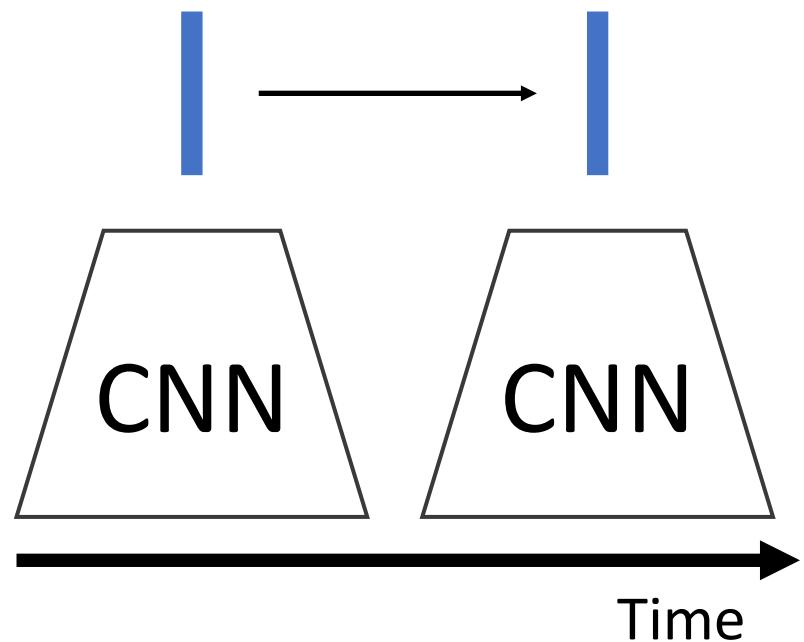
Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011
Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

Ballas et al, "Delving Deeper into Convolutional Networks for Learning Video Representations", ICLR 2016

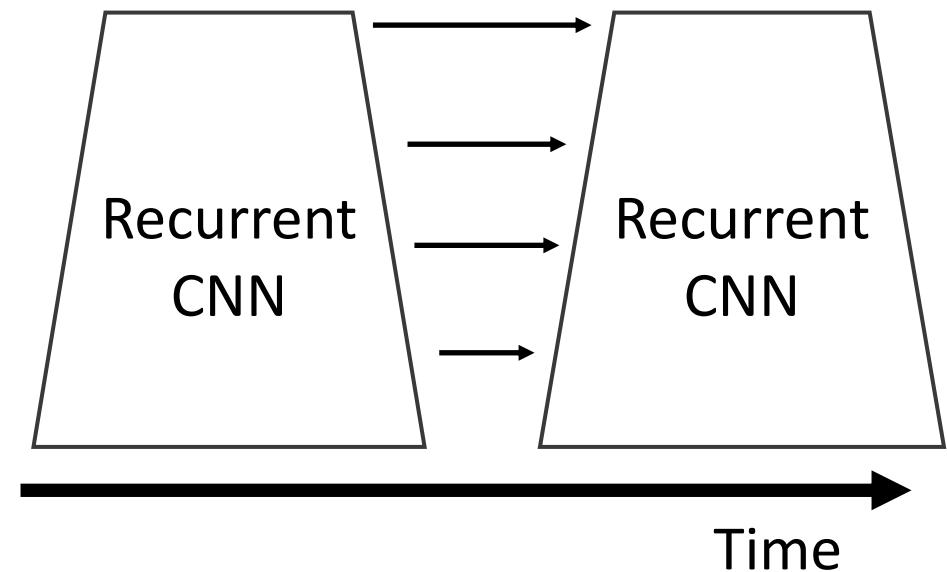
Modeling long-term temporal structure

Problem: RNNs are slow for long sequences (can't be parallelized)

RNN: Infinite temporal extent (fully-connected)



Recurrent CNN: Infinite temporal extent (convolutional)

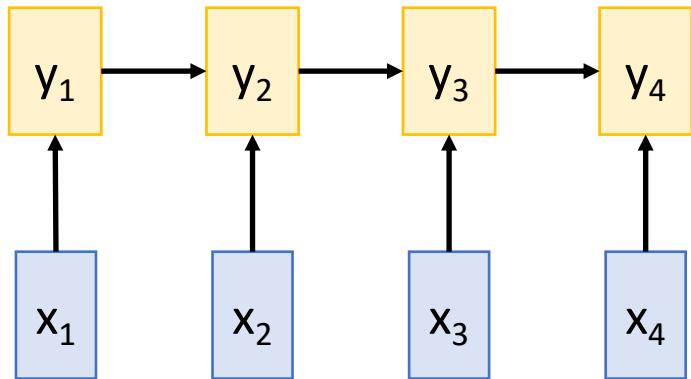


Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011
Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

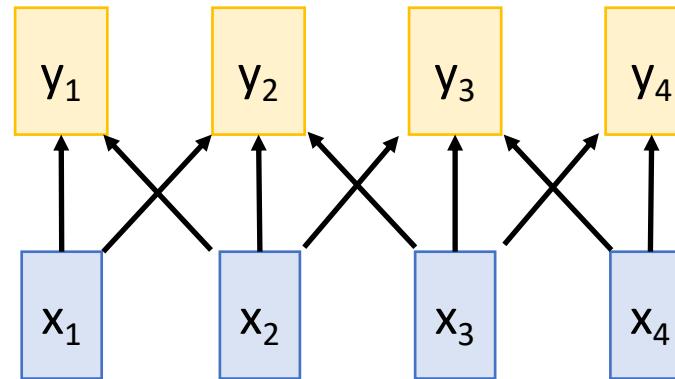
Ballas et al, "Delving Deeper into Convolutional Networks for Learning Video Representations", ICLR 2016

Recall: Different ways of processing sequences

Recurrent Neural Network



1D Convolution



Works on **Ordered Sequences**

(+) Good at long sequences: After one RNN layer, h_T "sees" the whole sequence

(-) Not parallelizable: need to compute hidden states sequentially

In video: CNN+RNN, or recurrent CNN

Works on **Multidimensional Grids**

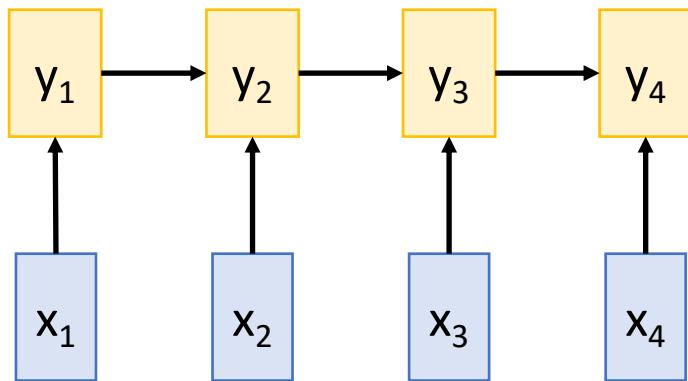
(-) Bad at long sequences: Need to stack many conv layers for outputs to "see" the whole sequence

(+) Highly parallel: Each output can be computed in parallel

In video: 3D convolution

Recall: Different ways of processing sequences

Recurrent Neural Network



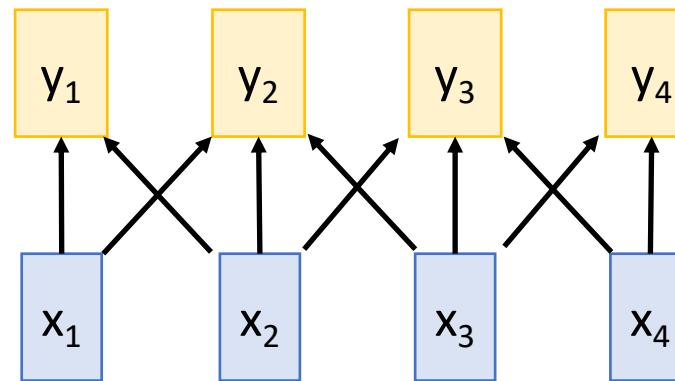
Works on **Ordered Sequences**

(+) Good at long sequences: After one RNN layer, h_T "sees" the whole sequence

(-) Not parallelizable: need to compute hidden states sequentially

In video: CNN+RNN, or recurrent CNN

1D Convolution



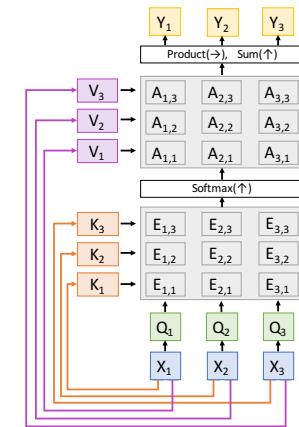
Works on **Multidimensional Grids**

(-) Bad at long sequences: Need to stack many conv layers for outputs to "see" the whole sequence

(+) Highly parallel: Each output can be computed in parallel

In video: 3D convolution

Self-Attention



Works on **Sets of Vectors**

(-) Good at long sequences: after one self-attention layer, each output "sees" all inputs!

(+) Highly parallel: Each output can be computed in parallel

(-) Very memory intensive

In video: ????

Recall: Self-Attention

Input: Set of vectors x_1, \dots, x_N

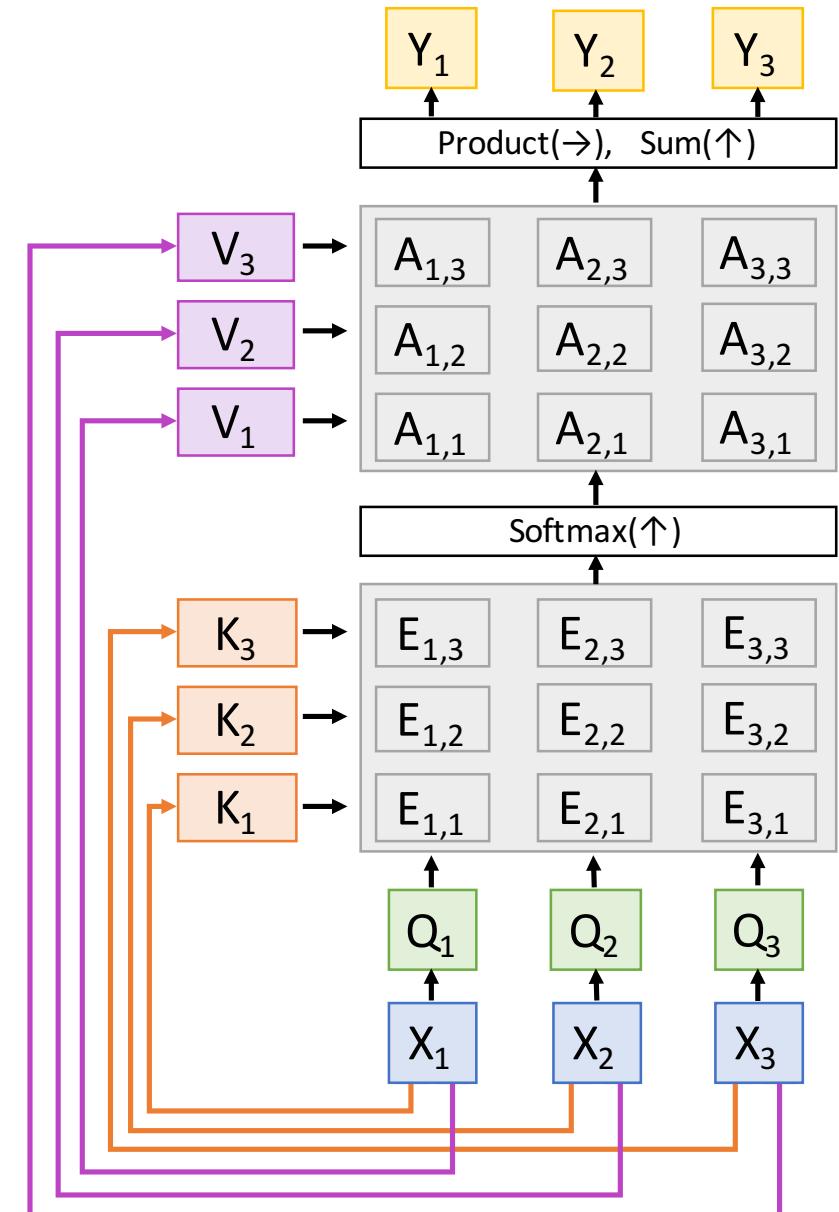
Keys, Queries, Values: Project each x to a key, query, and value using linear layer

Affinity matrix: Compare each pair of x , (using scaled dot-product between keys and values) and normalize using softmax

Output: Weighted sum of values, with weights given by affinity matrix

Features in 3D CNN: $C \times T \times H \times W$

Interpret as a set of THW vectors of dim C

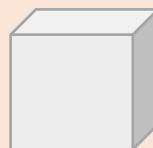


Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



3D
CNN



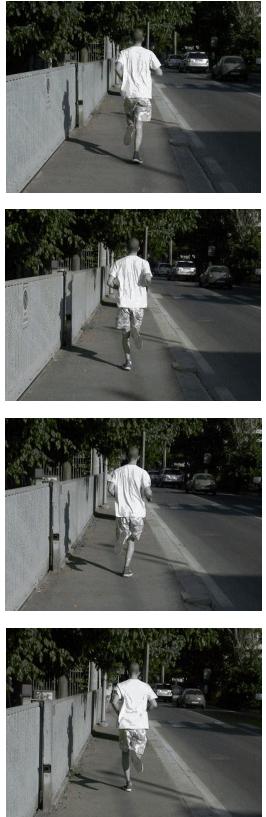
Features:
 $C \times T \times H \times W$

Nonlocal Block

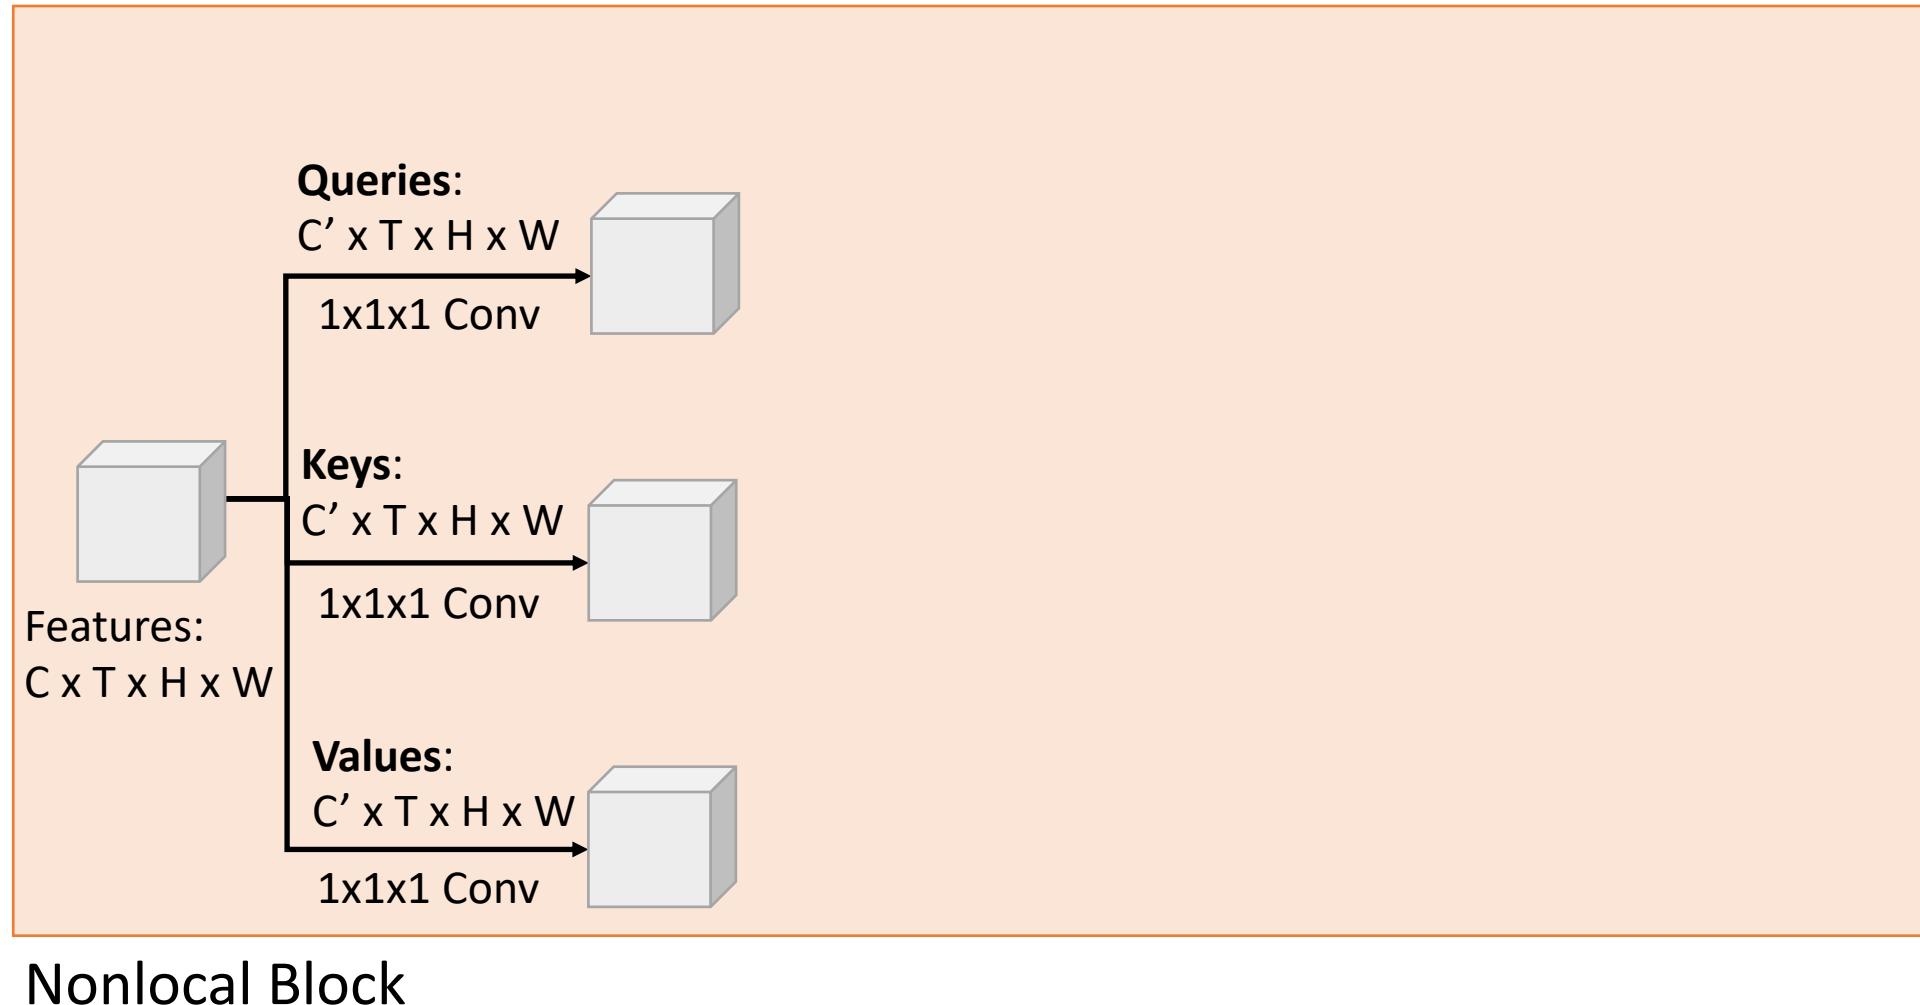
Wang et al, "Non-local neural networks", CVPR 2018

Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



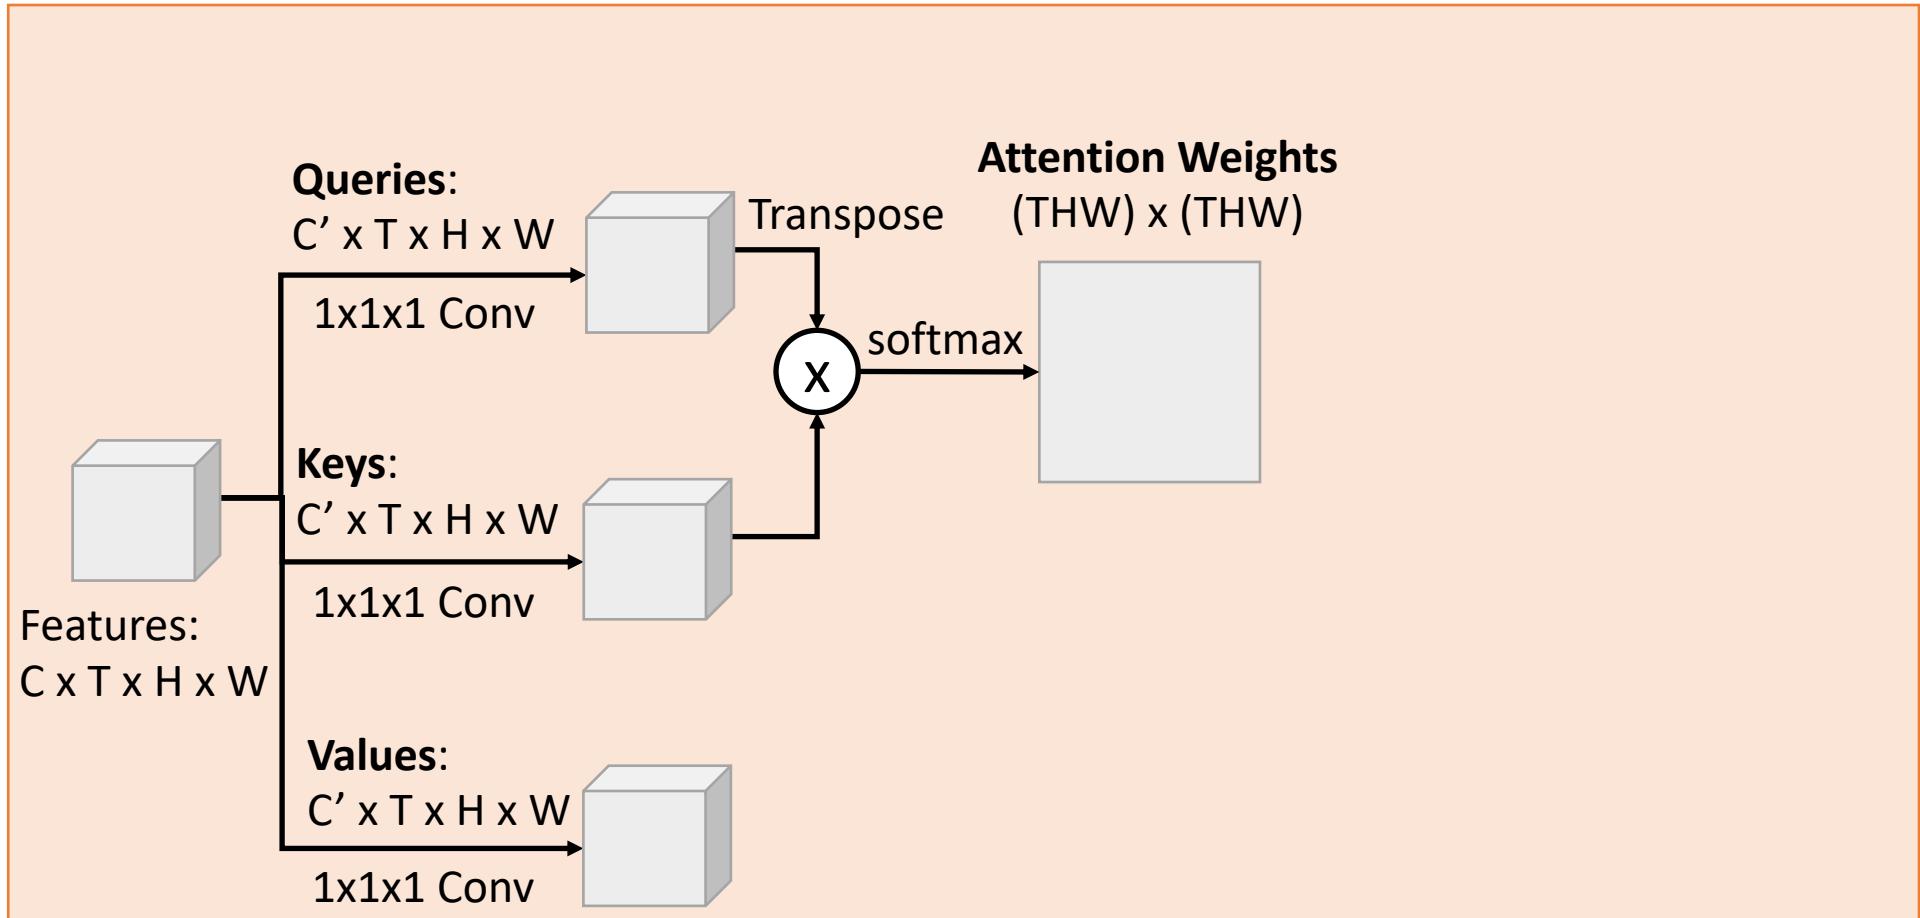
3D
CNN



Wang et al, "Non-local neural networks", CVPR 2018

Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

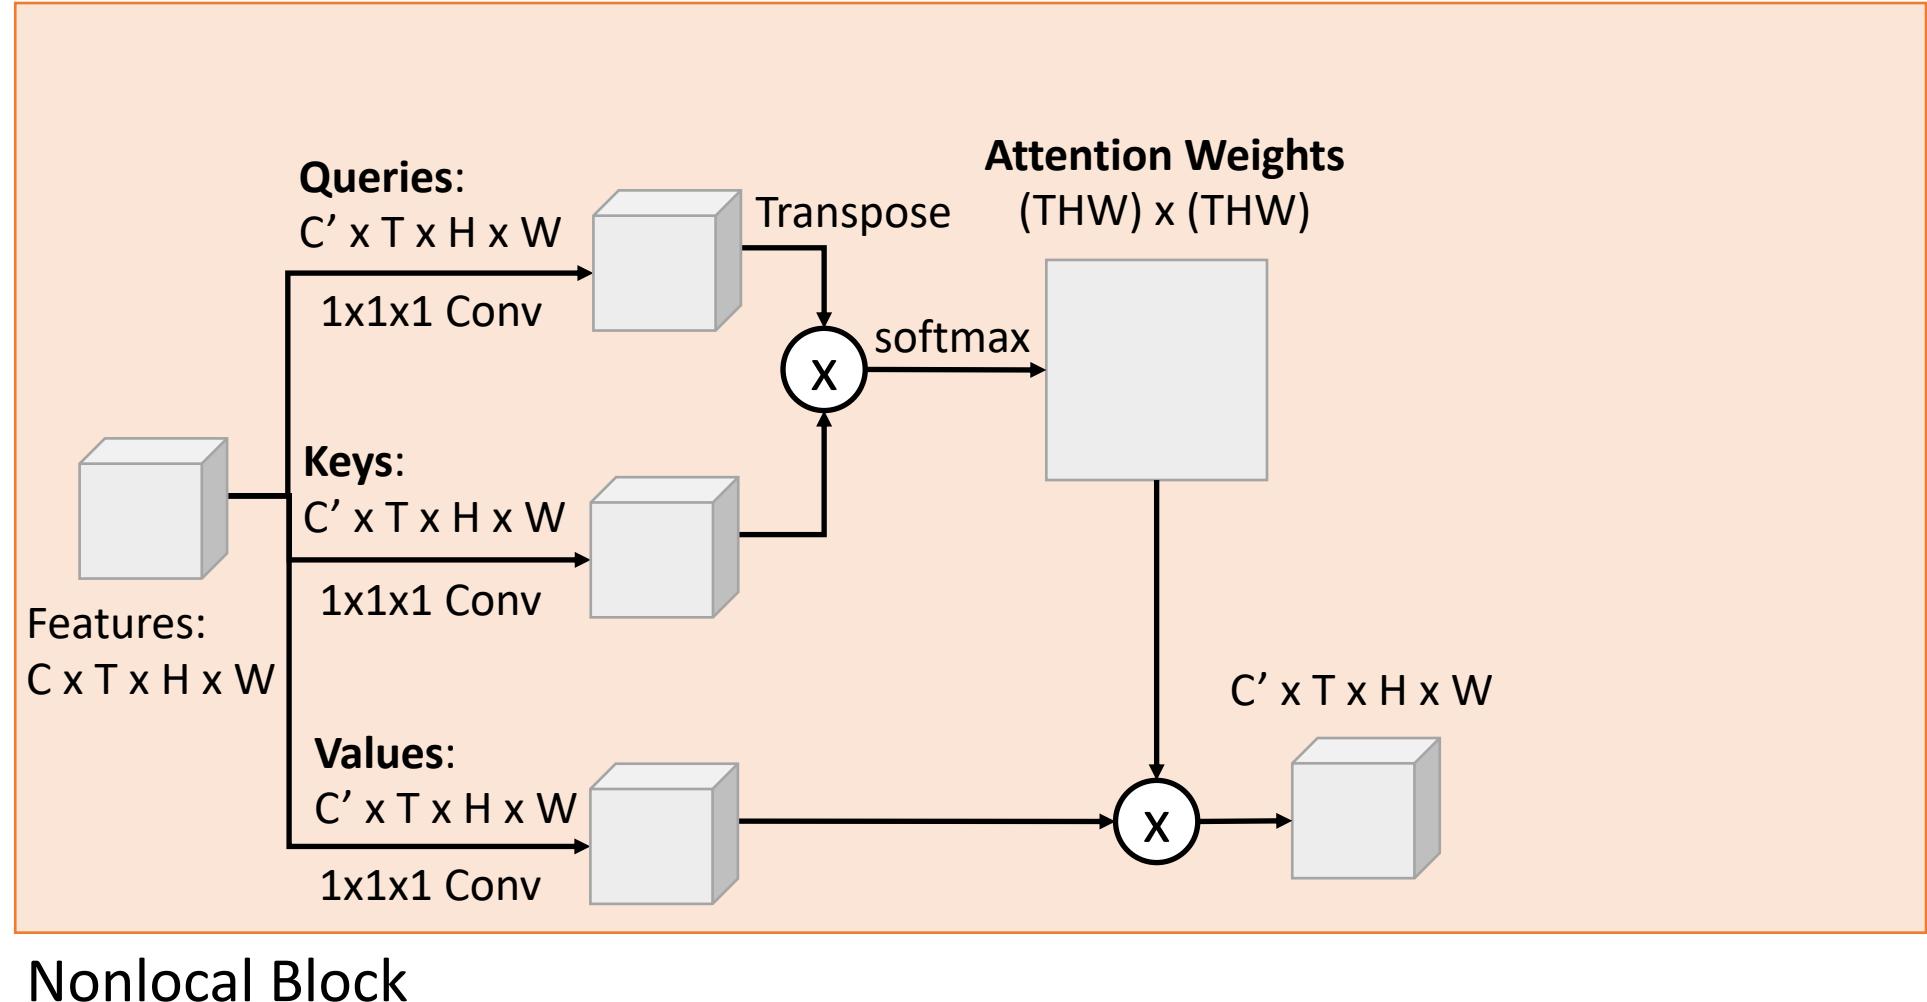


Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

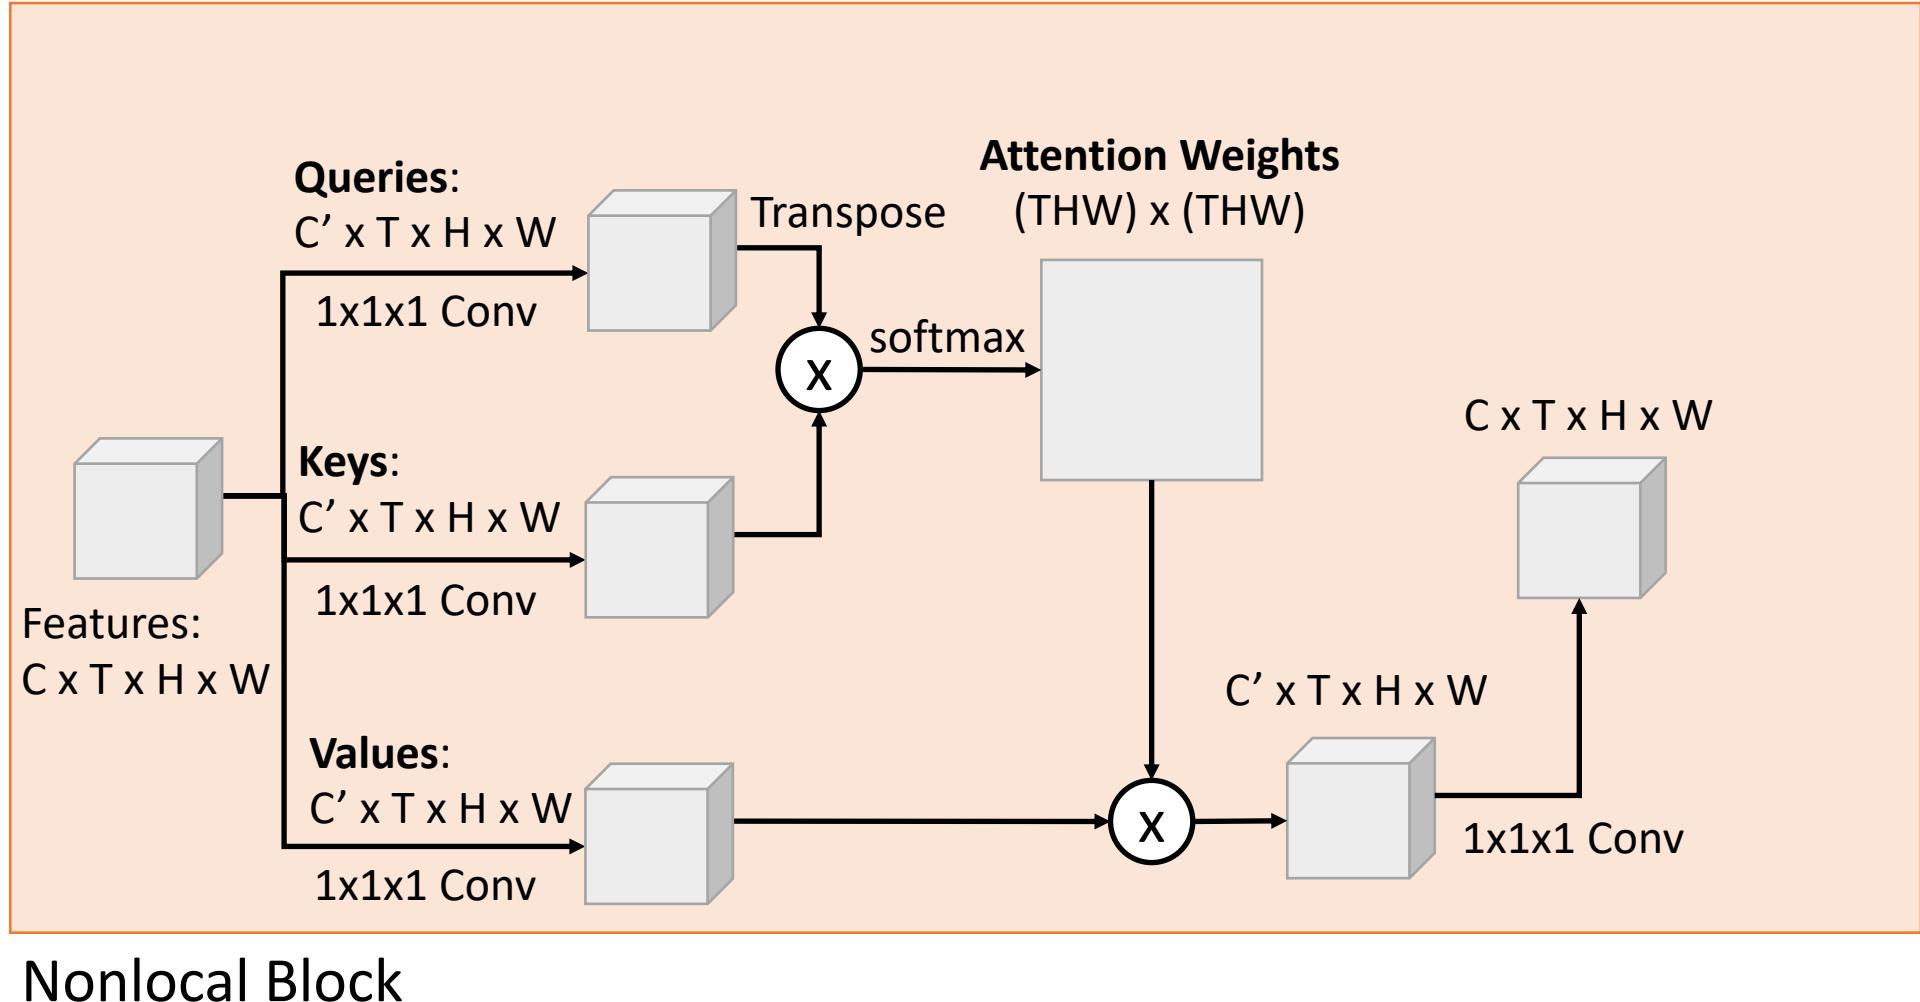


3D
CNN



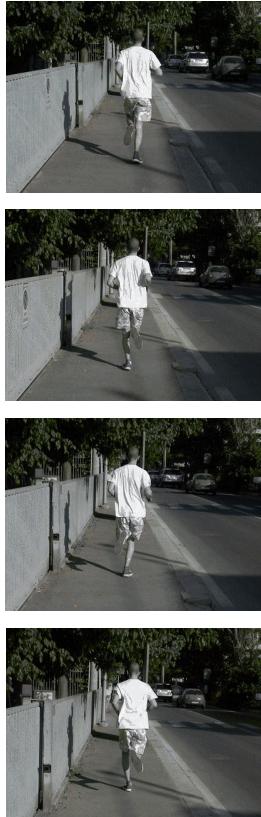
Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

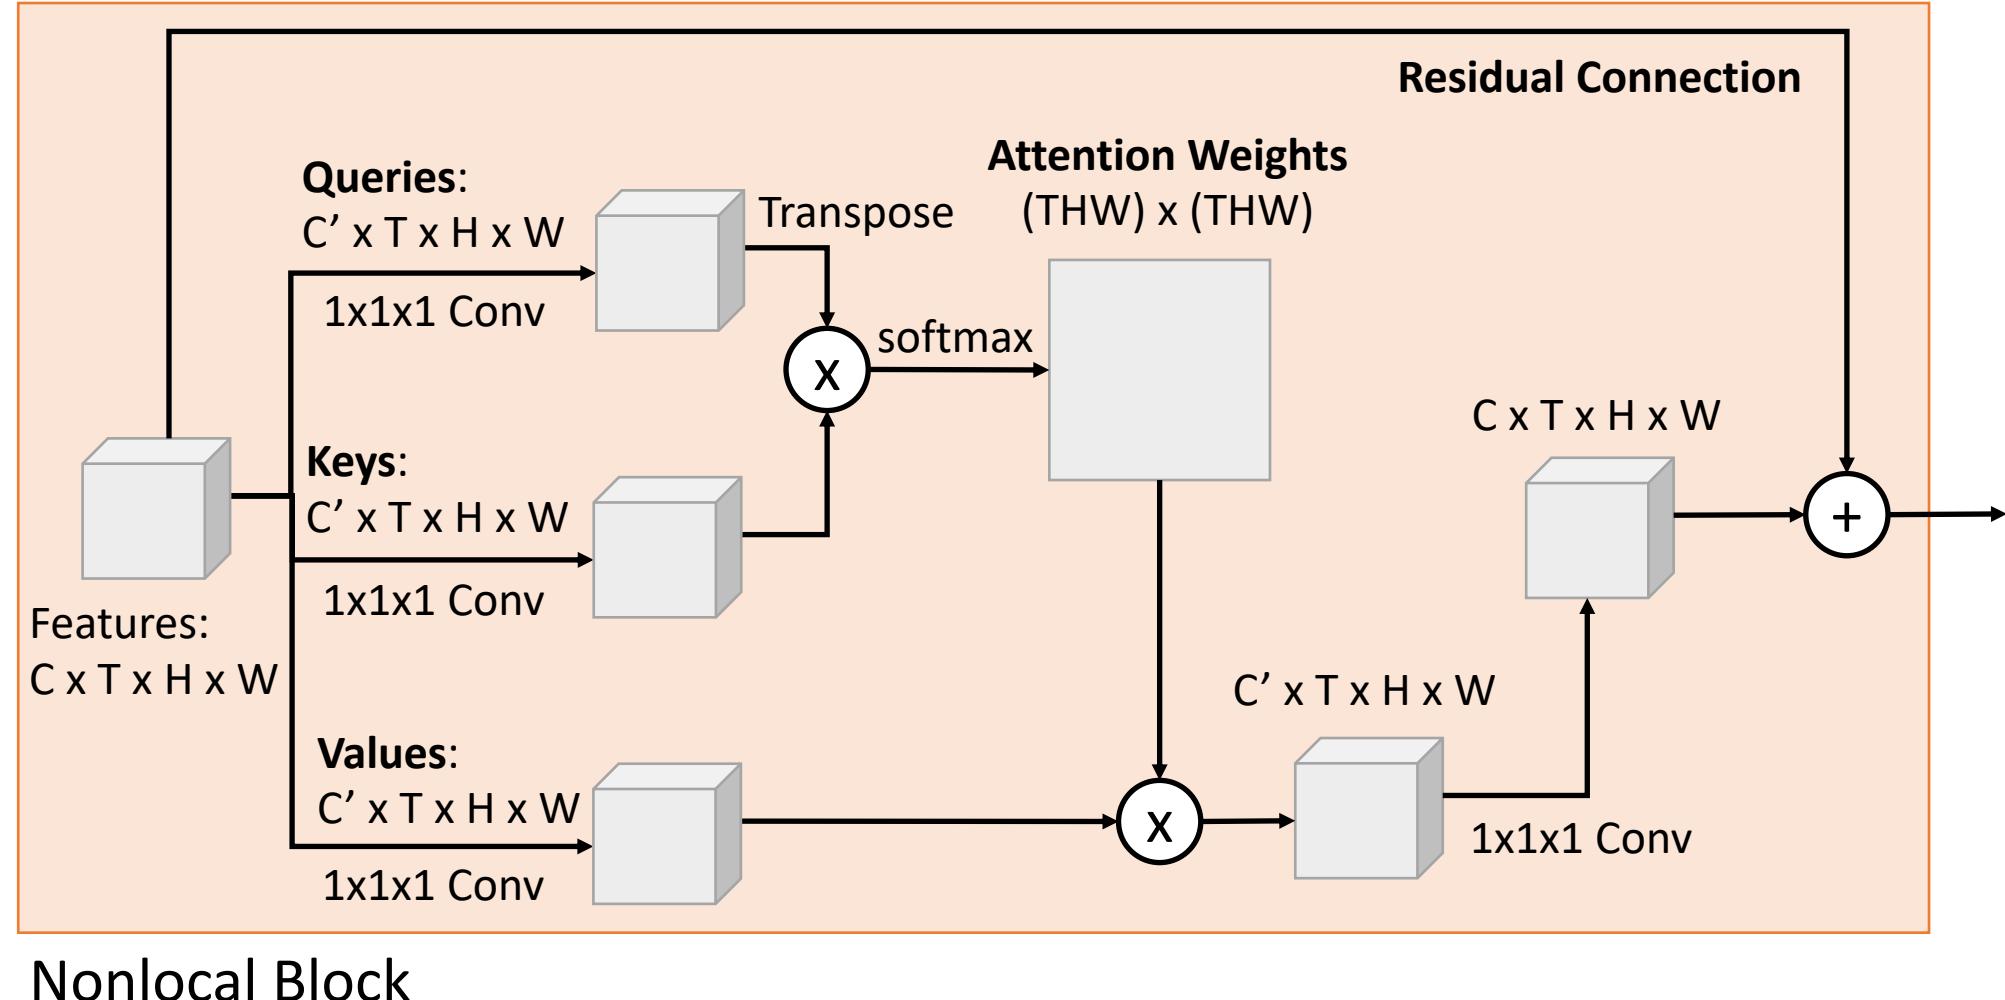


Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



3D
CNN



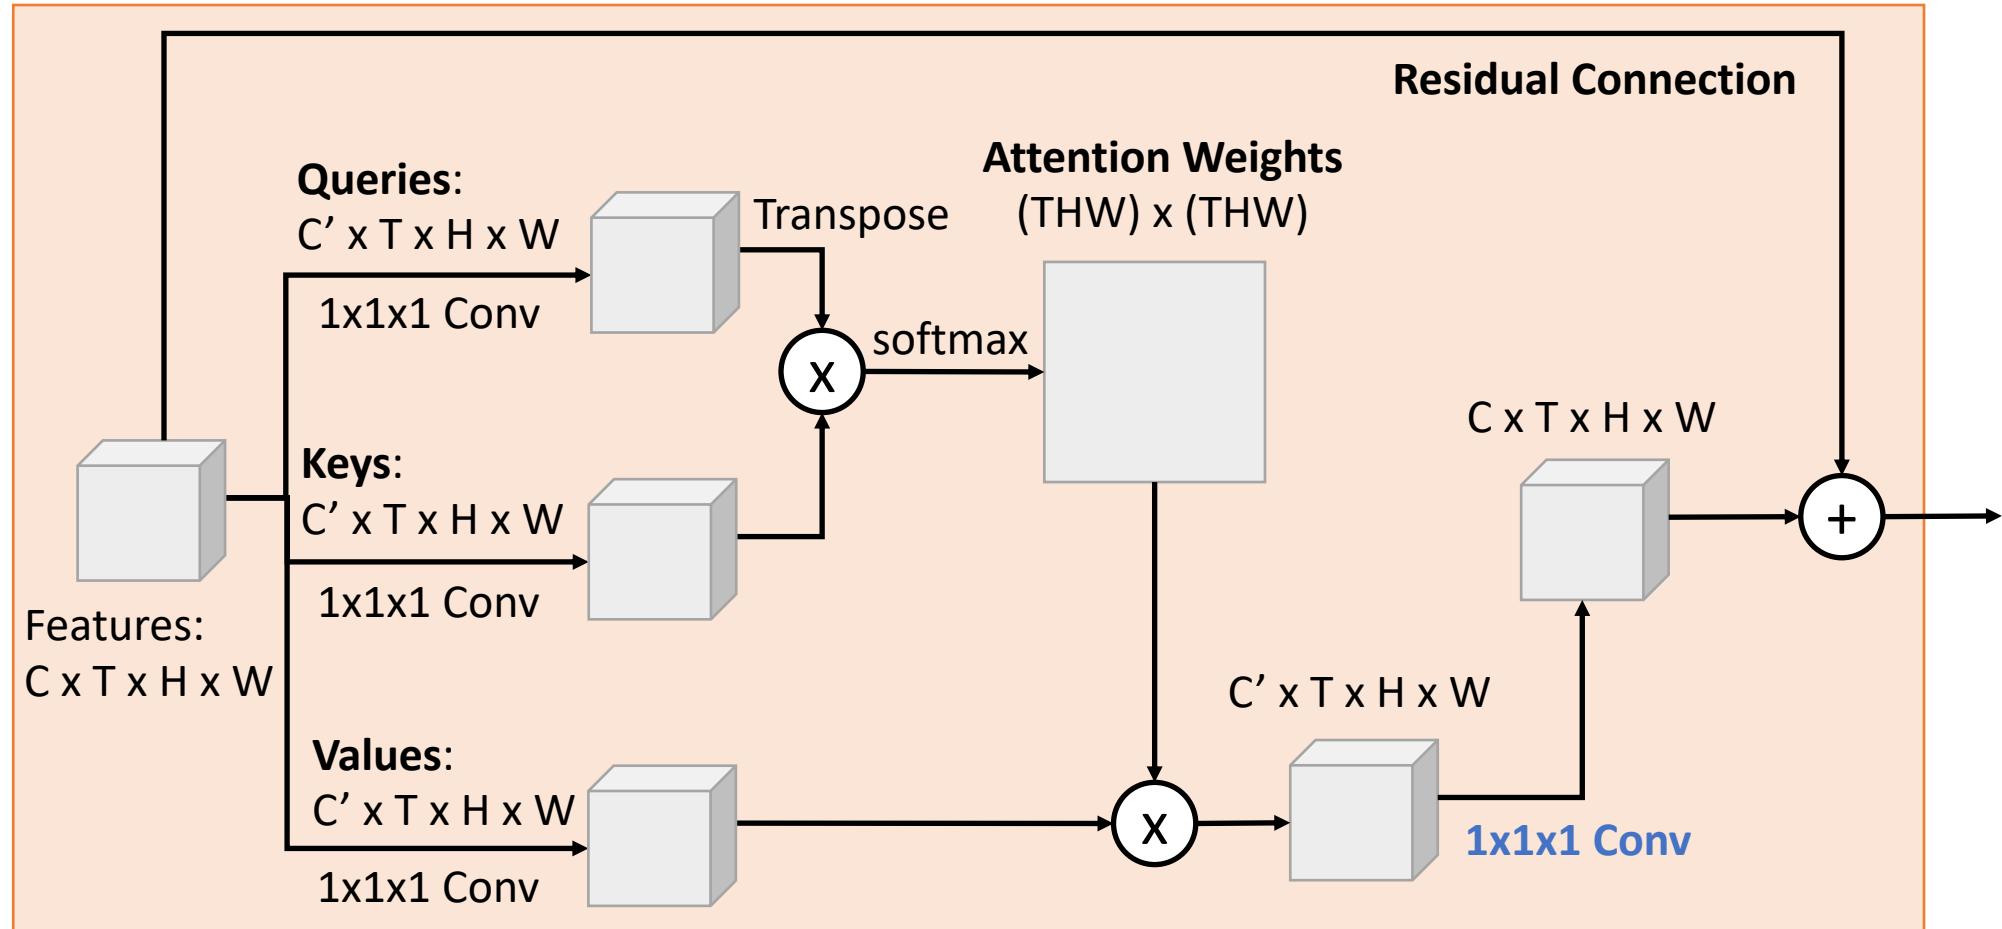
Wang et al, "Non-local neural networks", CVPR 2018

Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



3D
CNN



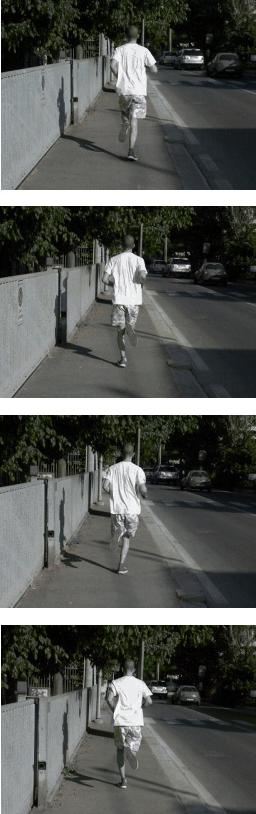
Nonlocal Block

Trick: Initialize **last conv** to 0, then entire block computes identity. Can insert into existing 3D CNNs

In practice, actually insert BatchNorm layer after final conv, and initialize scale parameter of BN layer to 0 rather than setting conv weight to 0

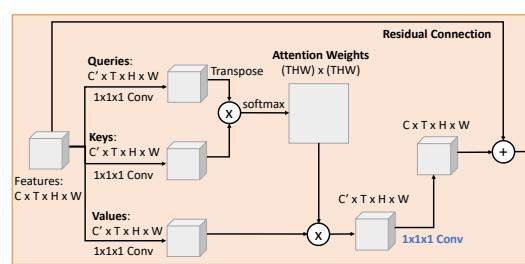
Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

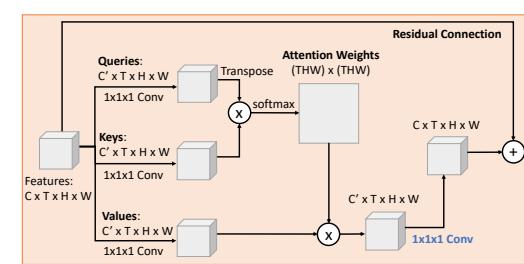
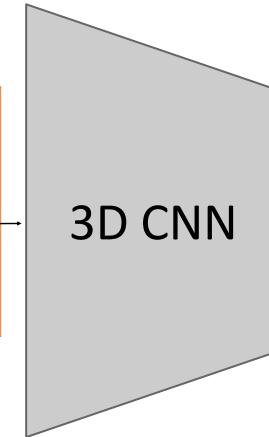


3D CNN

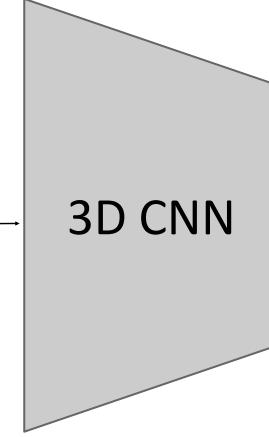
We can add nonlocal blocks into existing 3D CNN architectures.
But what is the best 3D CNN architecture?



Nonlocal Block



Nonlocal Block



Running

Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each 2D $K_h \times K_w$ conv/pool
layer with a 3D $K_t \times K_h \times K_w$ version

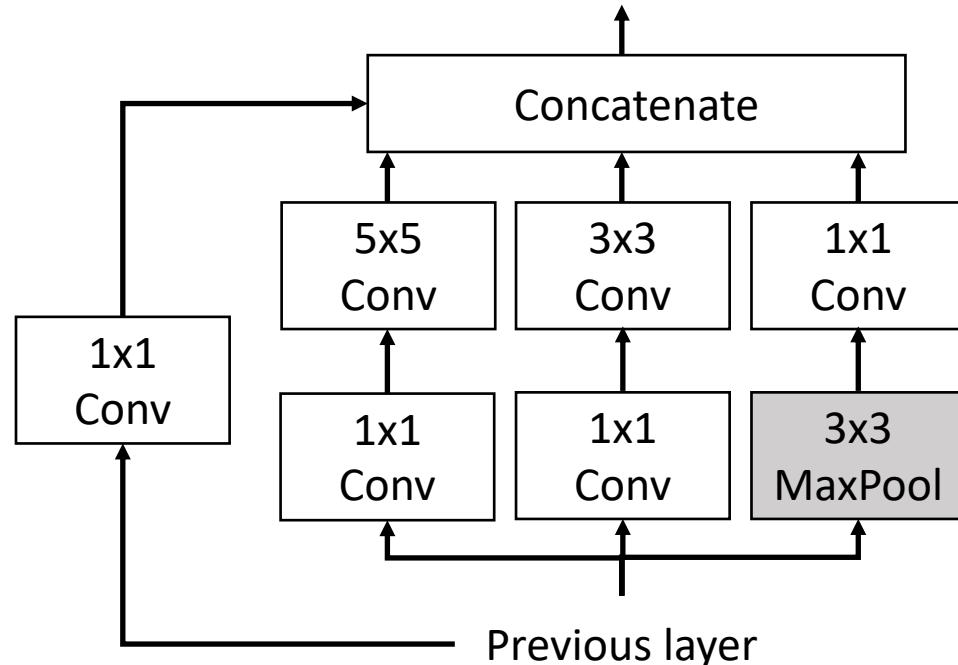
Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each $2D\ K_h \times K_w$ conv/pool layer with a $3D\ K_t \times K_h \times K_w$ version

Inception Block: Original



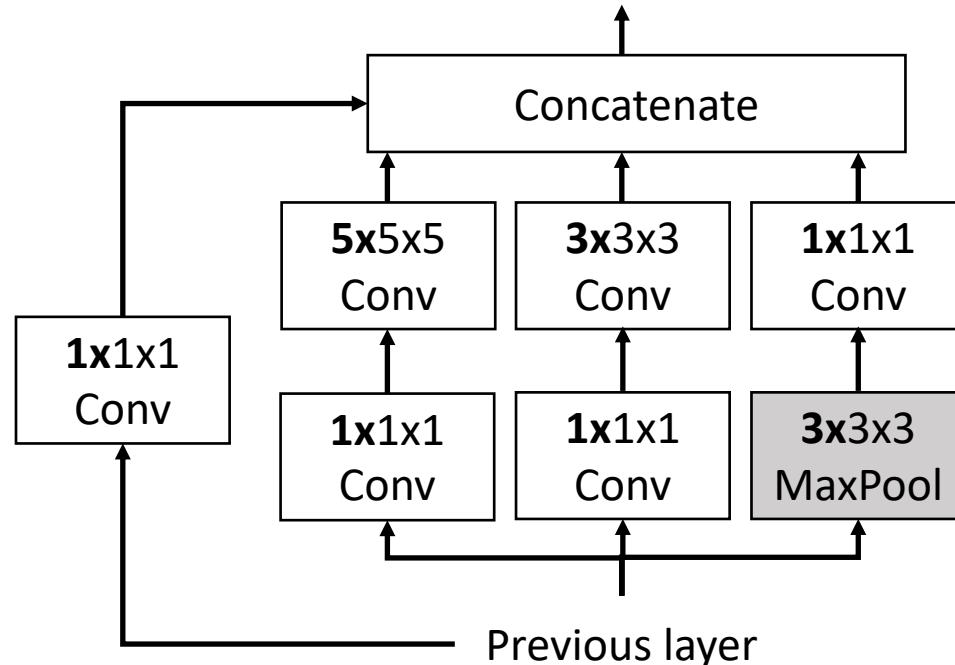
Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each $2D K_h \times K_w$ conv/pool layer with a $3D K_t \times K_h \times K_w$ version

Inception Block: Inflated



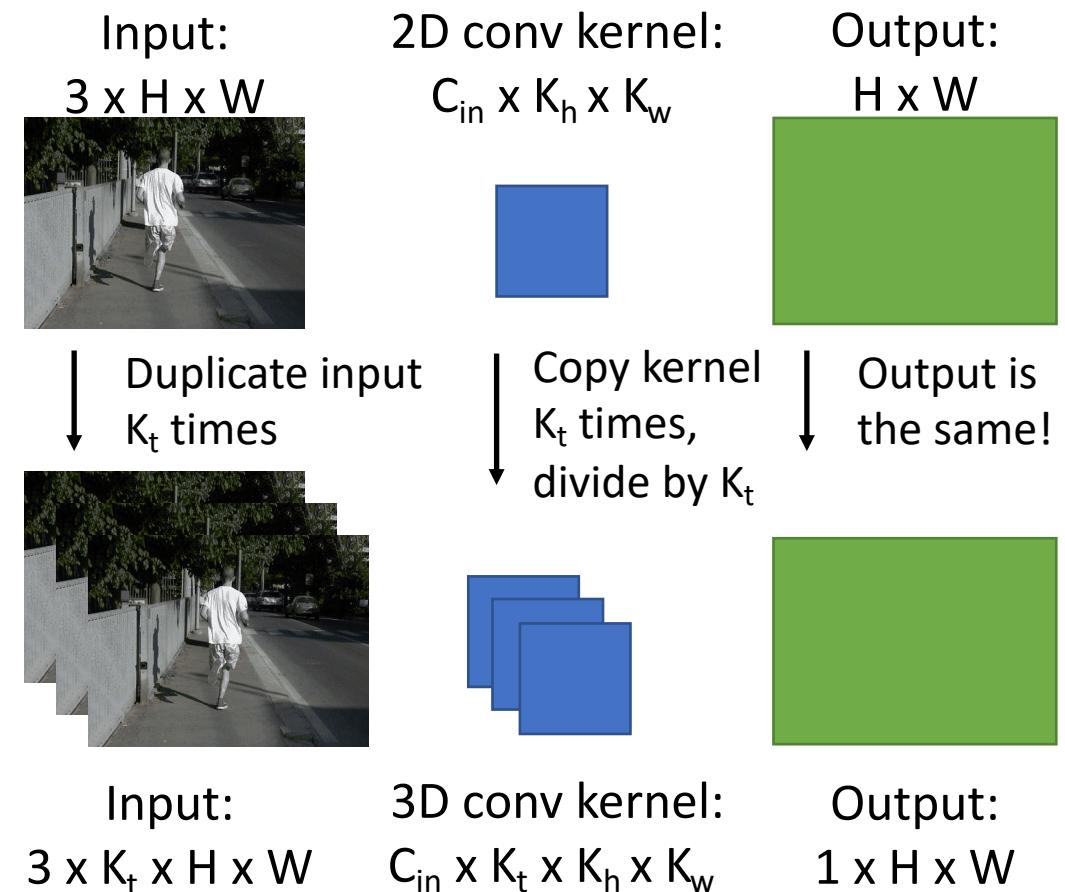
Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each 2D $K_h \times K_w$ conv/pool layer with a 3D $K_t \times K_h \times K_w$ version

Can use weights of 2D conv to initialize 3D conv: copy K_t times in space and divide by K_t
This gives the same result as 2D conv given “constant” video input



Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.
Can we reuse image architectures for video?

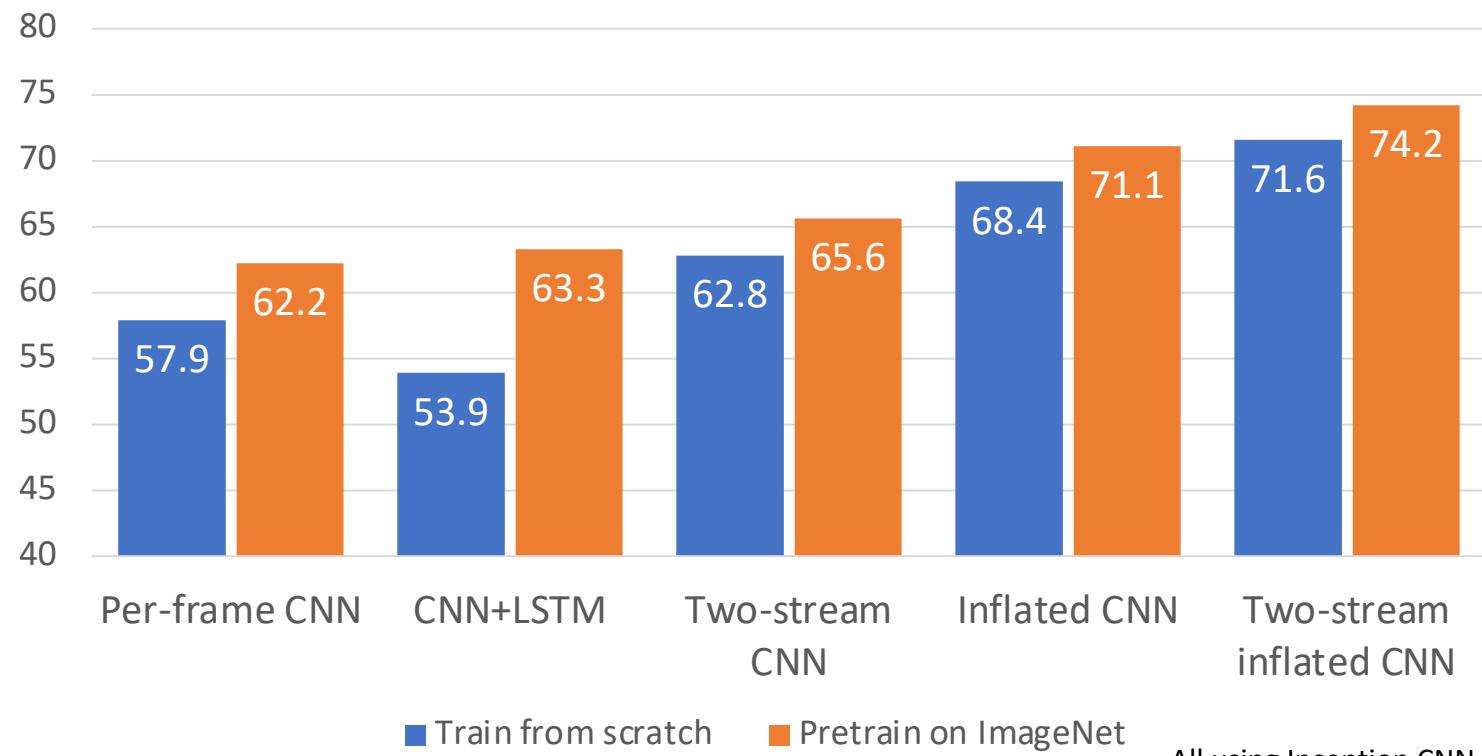
Idea: take a 2D CNN architecture.

Replace each $2D K_h \times K_w$ conv/pool layer with a $3D K_t \times K_h \times K_w$ version

Can use weights of 2D conv to initialize 3D conv: copy K_t times in space and divide by K_t

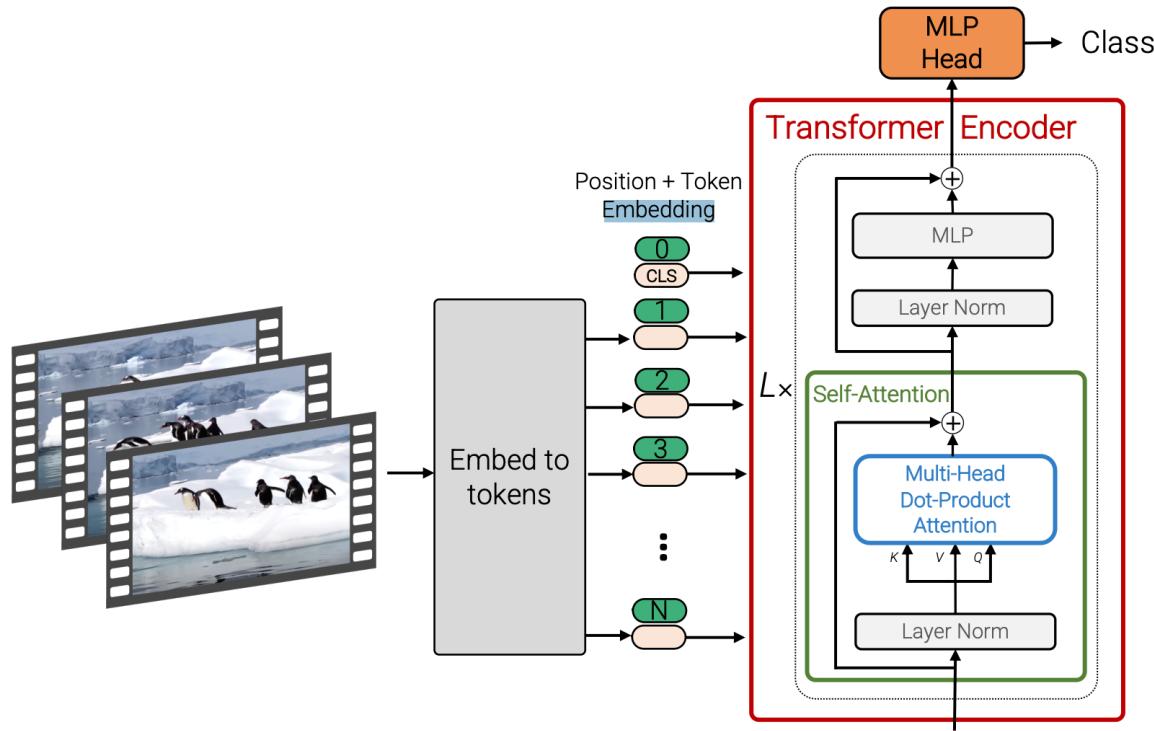
This gives the same result as 2D conv given “constant” video input

Top-1 Accuracy on Kinetics-400

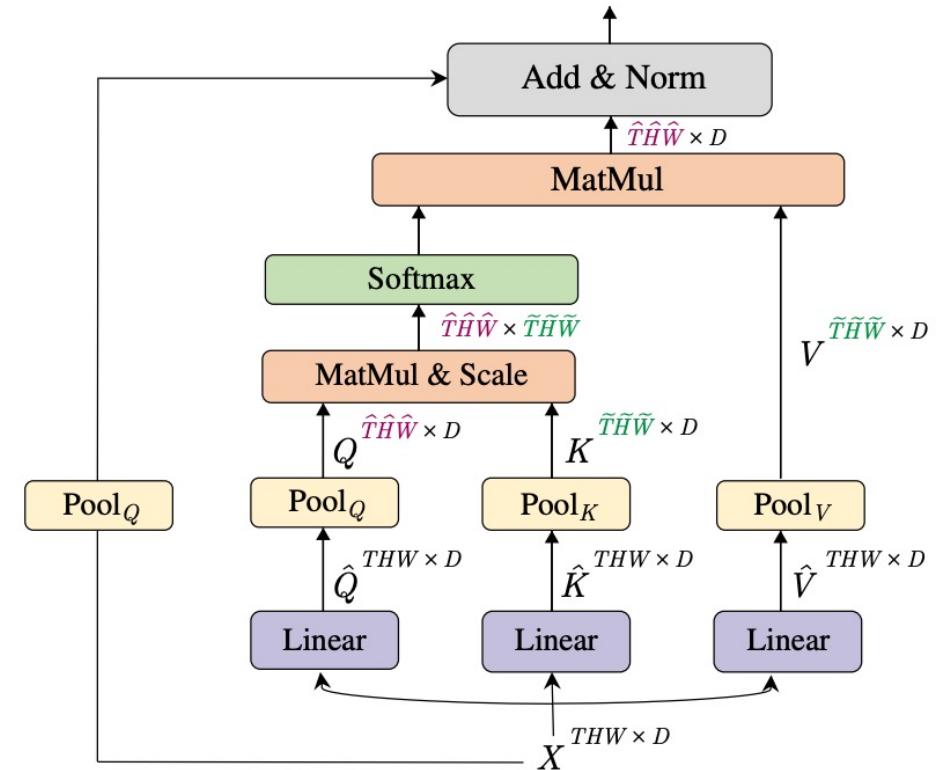


Vision Transformers for Video

Factorized attention: Attend over space / time



Pooling module: Reduce number of tokens



Bertasius et al, "Is Space-Time Attention All You Need for Video Understanding?", ICML 2021

Arnab et al, "ViViT: A Video Vision Transformer", ICCV 2021

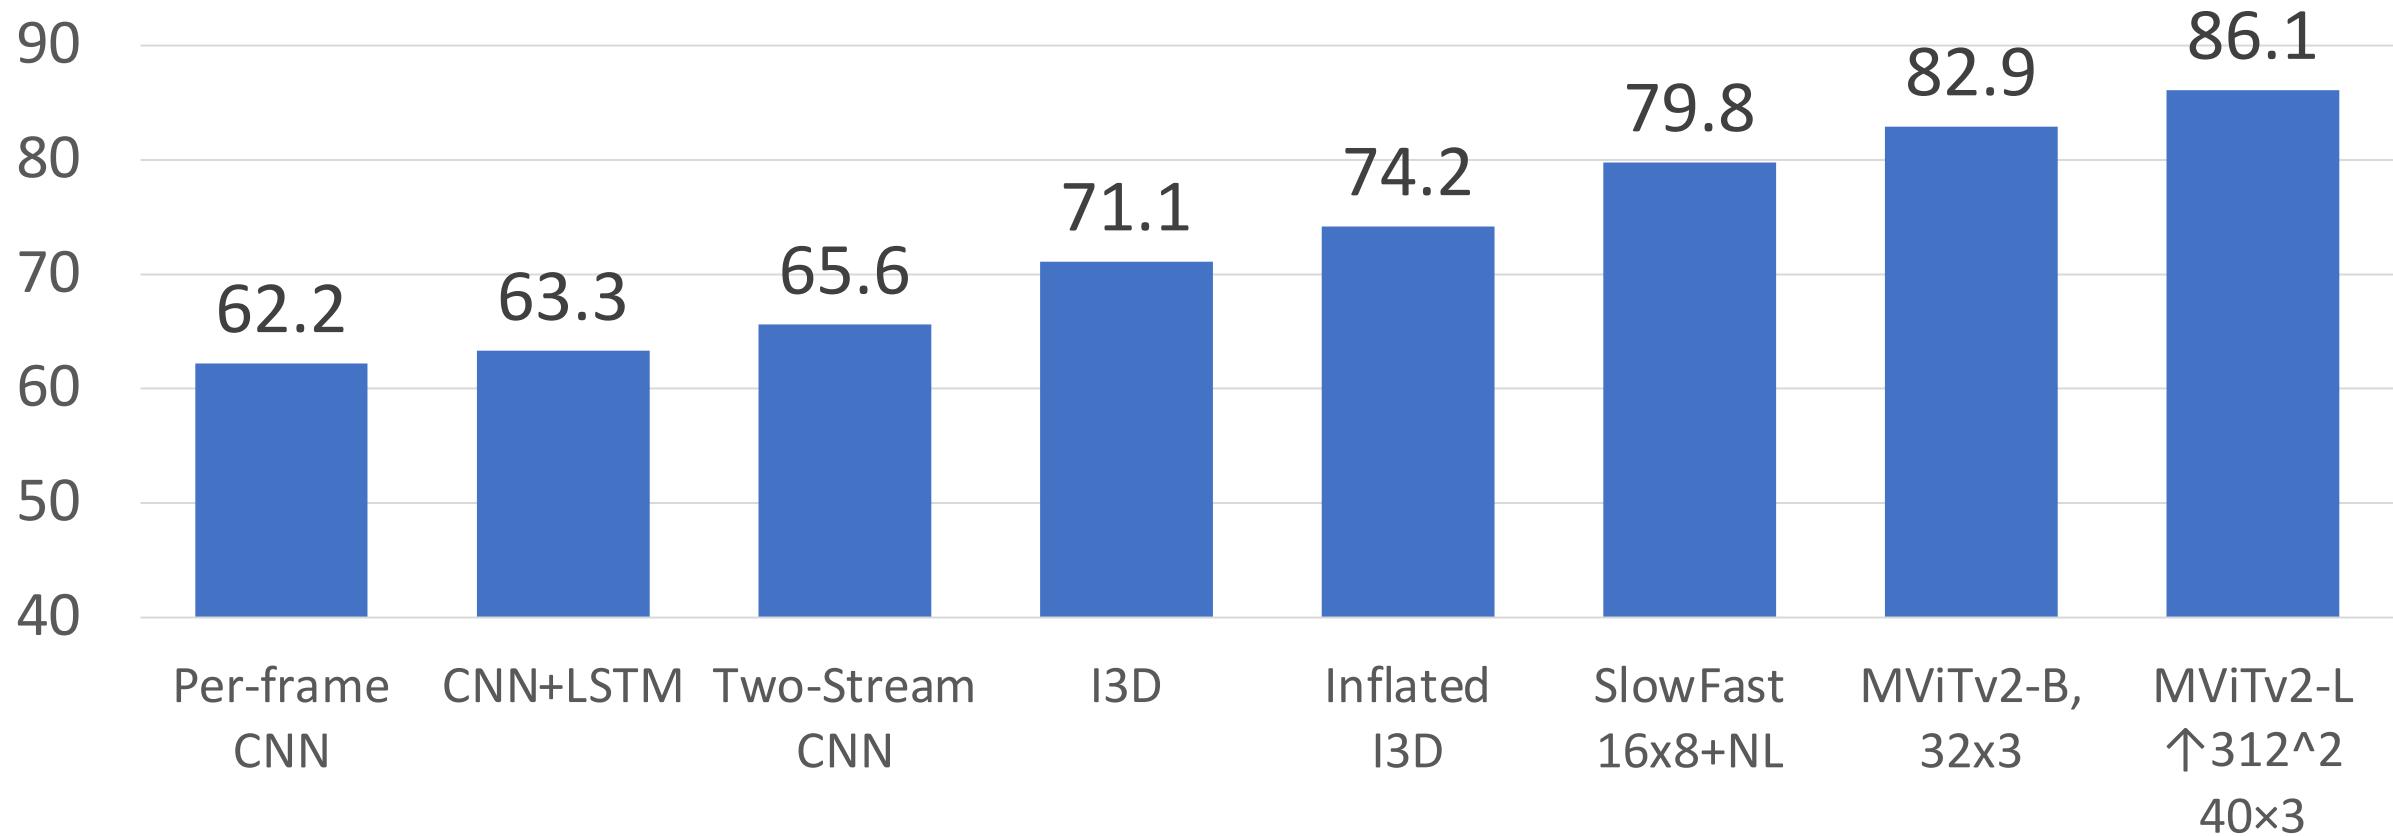
Neimark et al, "Video Transformer Network", ICCV 2021

Fan et al, "Multiscale Vision Transformers", ICCV 2021

Li et al, "MViTv2: Improved Multiscale Vision Transformers for Classification and Detection", CVPR 2022

Vision Transformers for Video

Top-1 Accuracy on Kinetics-400



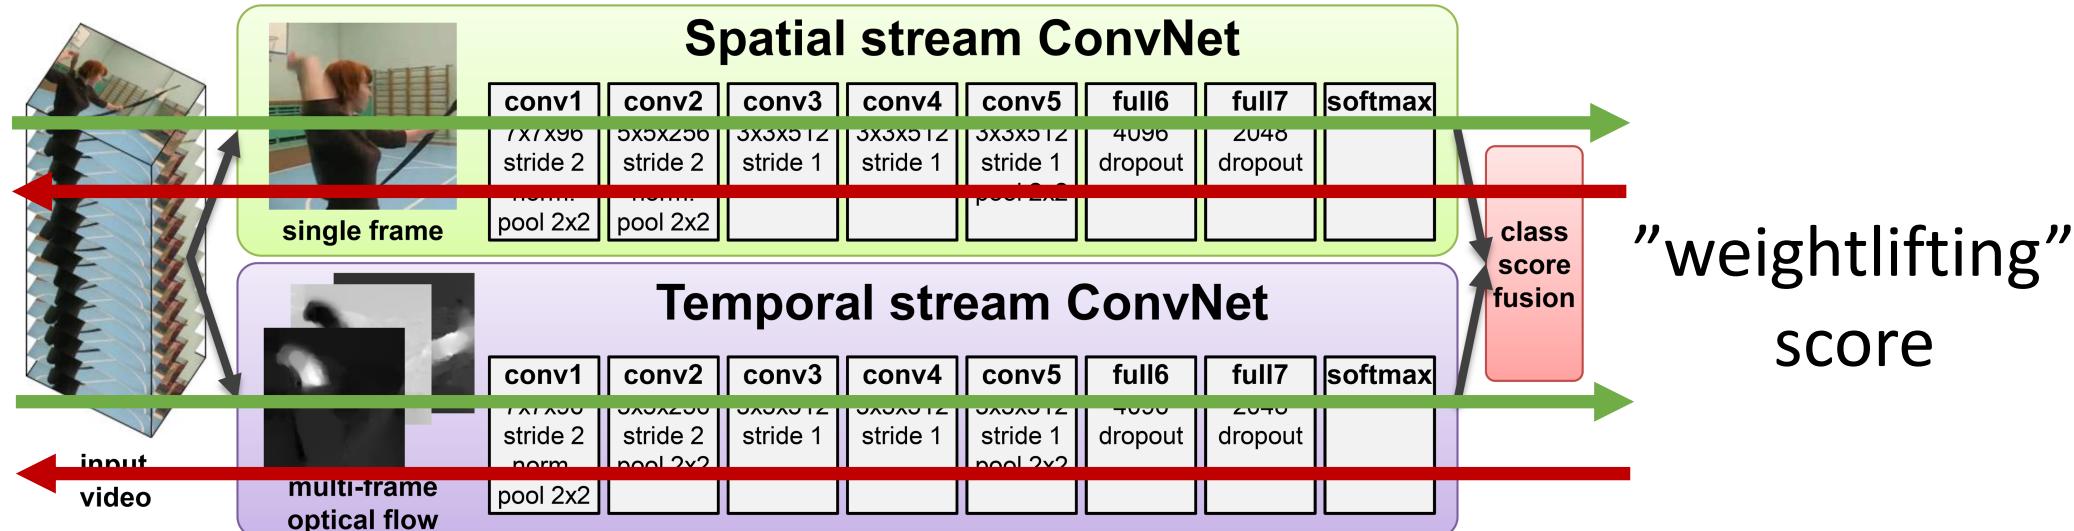
Li et al, "MViTv2: Improved Multiscale Vision Transformers for Classification and Detection", CVPR 2022

Visualizing Video Models

Image



Forward: Compute class score



"weightlifting"
score

Flow

Backward: Compute gradient

Add a term to encourage spatially smooth flow; tune penalty to pick out “slow” vs “fast” motion

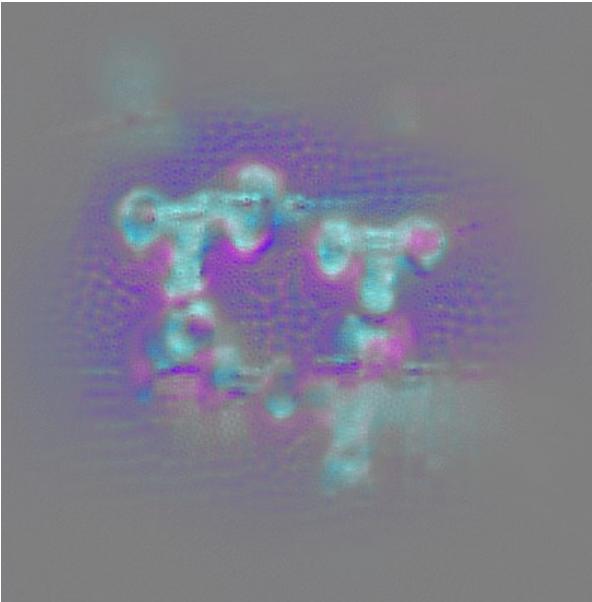
Figure credit: Simonyan and Zisserman, “Two-stream convolutional networks for action recognition in videos”, NeurIPS 2014
Feichtenhofer et al, “What have we learned from deep representations for action recognition?”, CVPR 2018
Feichtenhofer et al, “Deep insights into convolutional networks for video recognition?”, IJCV 2019.

Can you guess the action?

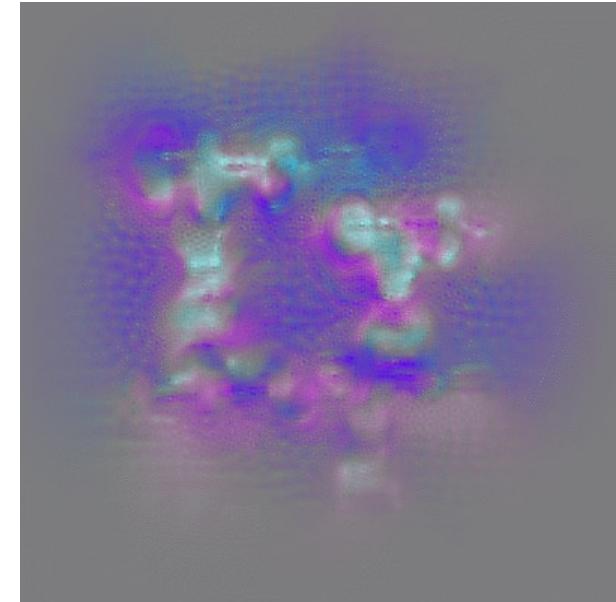
Appearance



“Slow” motion



“Fast” motion

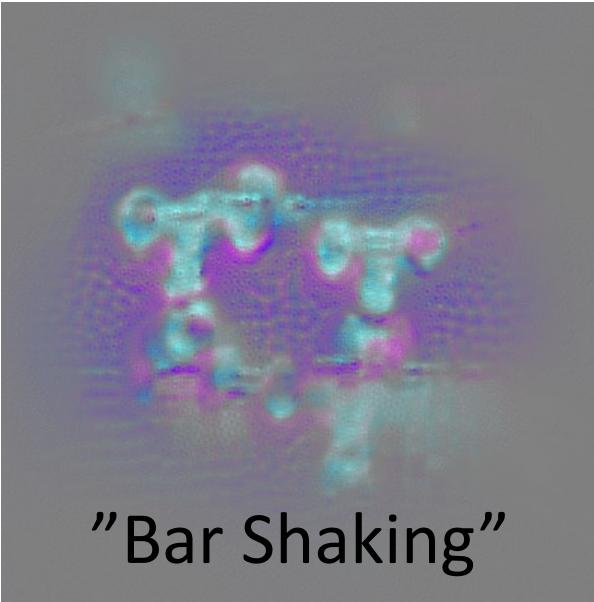


Can you guess the action? Weightlifting

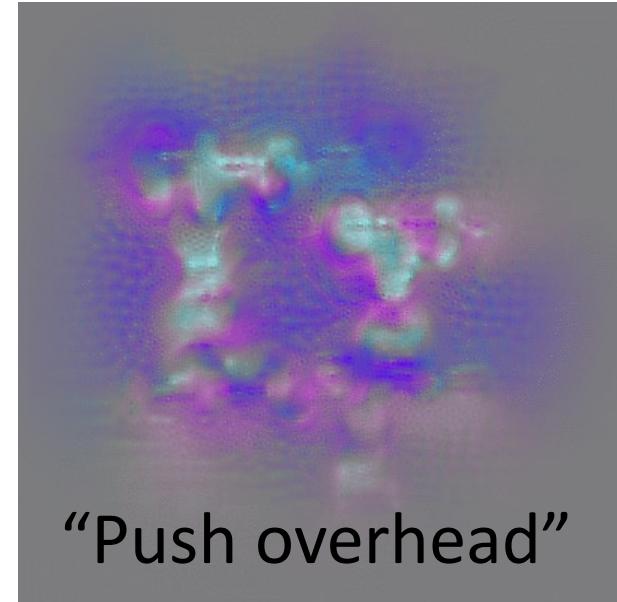
Appearance



“Slow” motion



“Fast” motion



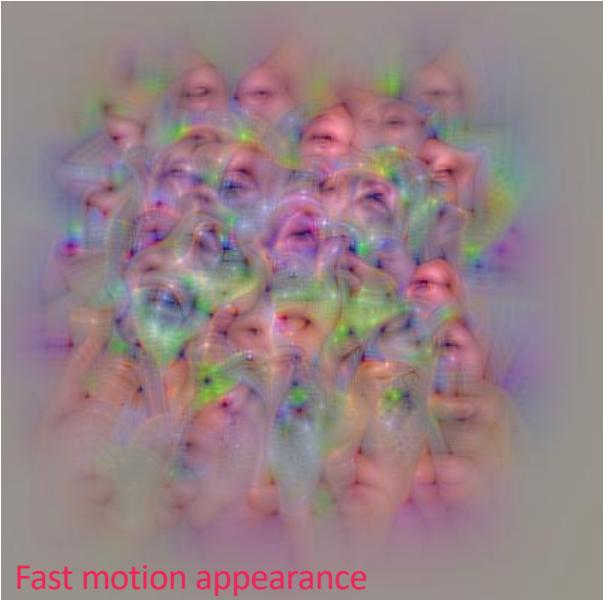
“Bar Shaking”

“Push overhead”

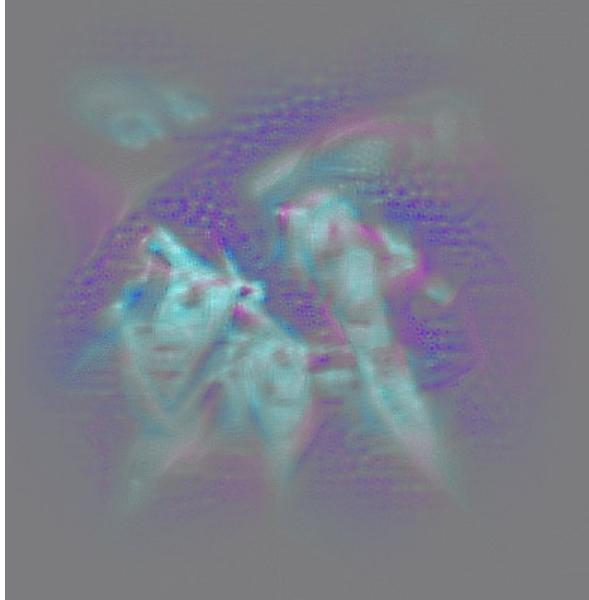


Can you guess the action?

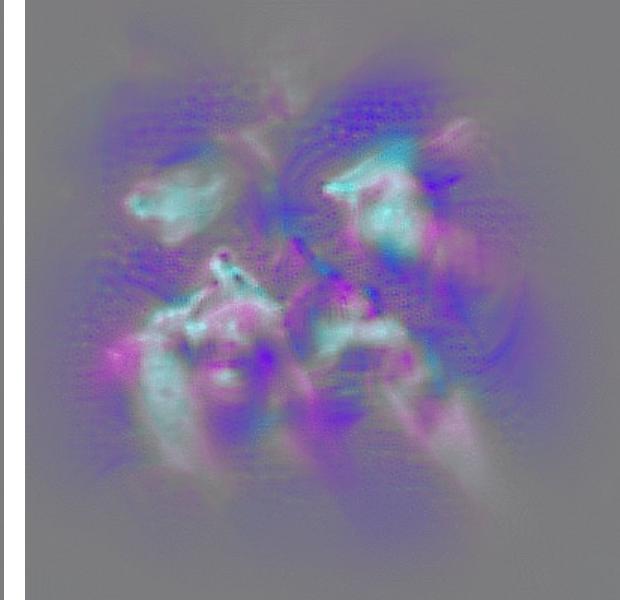
Appearance



“Slow” motion

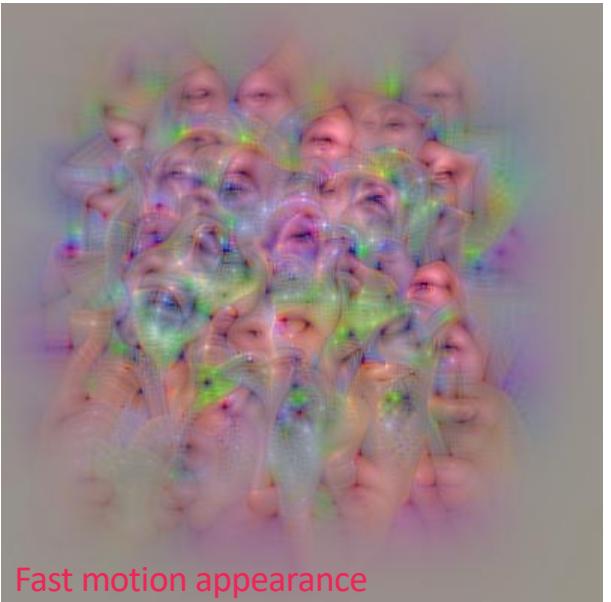


“Fast” motion

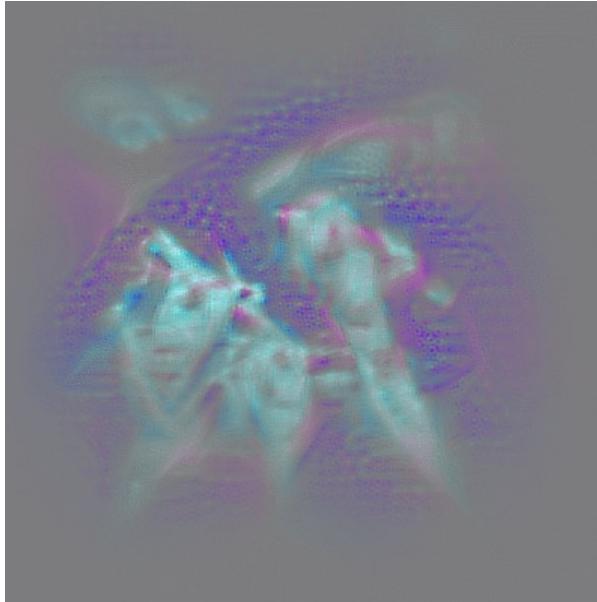


Can you guess the action? Apply Eye Makeup

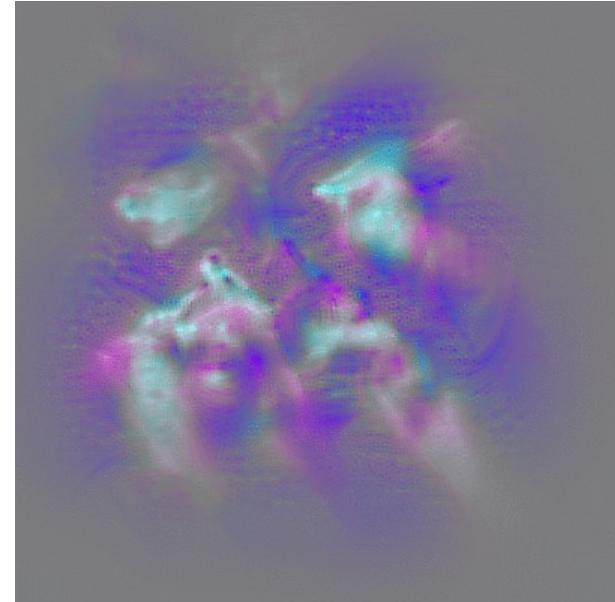
Appearance



“Slow” motion



“Fast” motion



So far: Classify short clips



Videos: Recognize **actions**



Swimming
Running
Jumping
Eating
Standing

Temporal Action Localization

Given a long untrimmed video sequence, identify frames corresponding to different actions

Running



Jumping

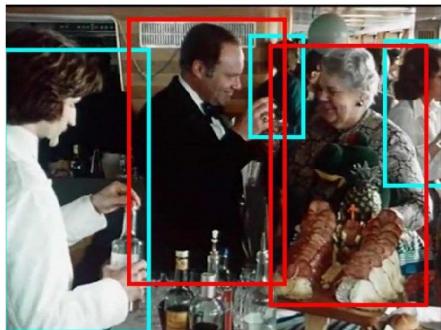


Can use architecture similar to Faster R-CNN:
first generate **temporal proposals** then **classify**

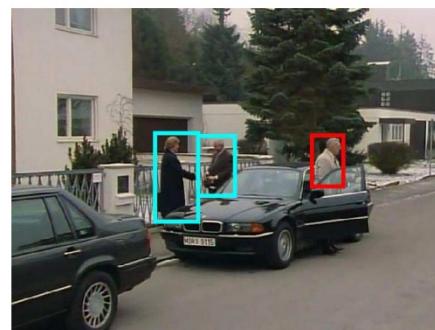
Chao et al, "Rethinking the Faster R-CNN Architecture for Temporal Action Localization", CVPR 2018

Spatio-Temporal Detection

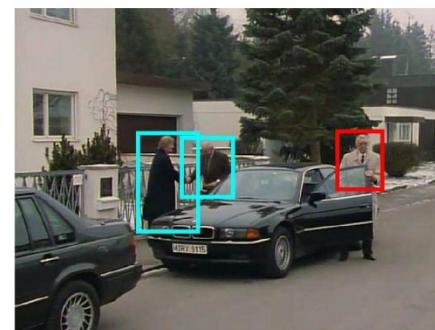
Given a long untrimmed video, detect all the people in space and time and classify the activities they are performing
Some examples from AVA Dataset:



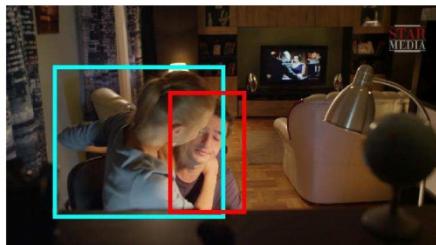
click glass → drink



open → close



grab (a person) → hug



look at phone → answer phone



Gu et al, "AVA: A Video Dataset of Spatio-temporally Localized Atomic Visual Actions", CVPR 2018

Ego4D: New Large-Scale Video Dataset

3670 hours of **egocentric** video (head-mounted cameras)

Long videos: 1-10 hours each

Diverse: data collected by 14 teams spread across 9 countries; 931 camera wearers (not just grad students!)

Natural-language narrations (3.85M sentences)

Support for 5 different tasks:

- Episodic Memory
- Hands and Objects
- Audio-Video Diarization
- Social Interactions
- Forecasting



Grauman et al, "Ego4D: Around the World in 3,000 Hours of Egocentric Video", CVPR 2022

Recap: Video Models

Many video models:

Single-frame CNN (Try this first!)

Late fusion

Early fusion

3D CNN / C3D

Two-stream networks

CNN + RNN

Convolutional RNN

Spatio-temporal self-attention

Next Time: Conclusion